


# Pentium I/O:


I/O, I/O and off to work we go

Kin K. Leung

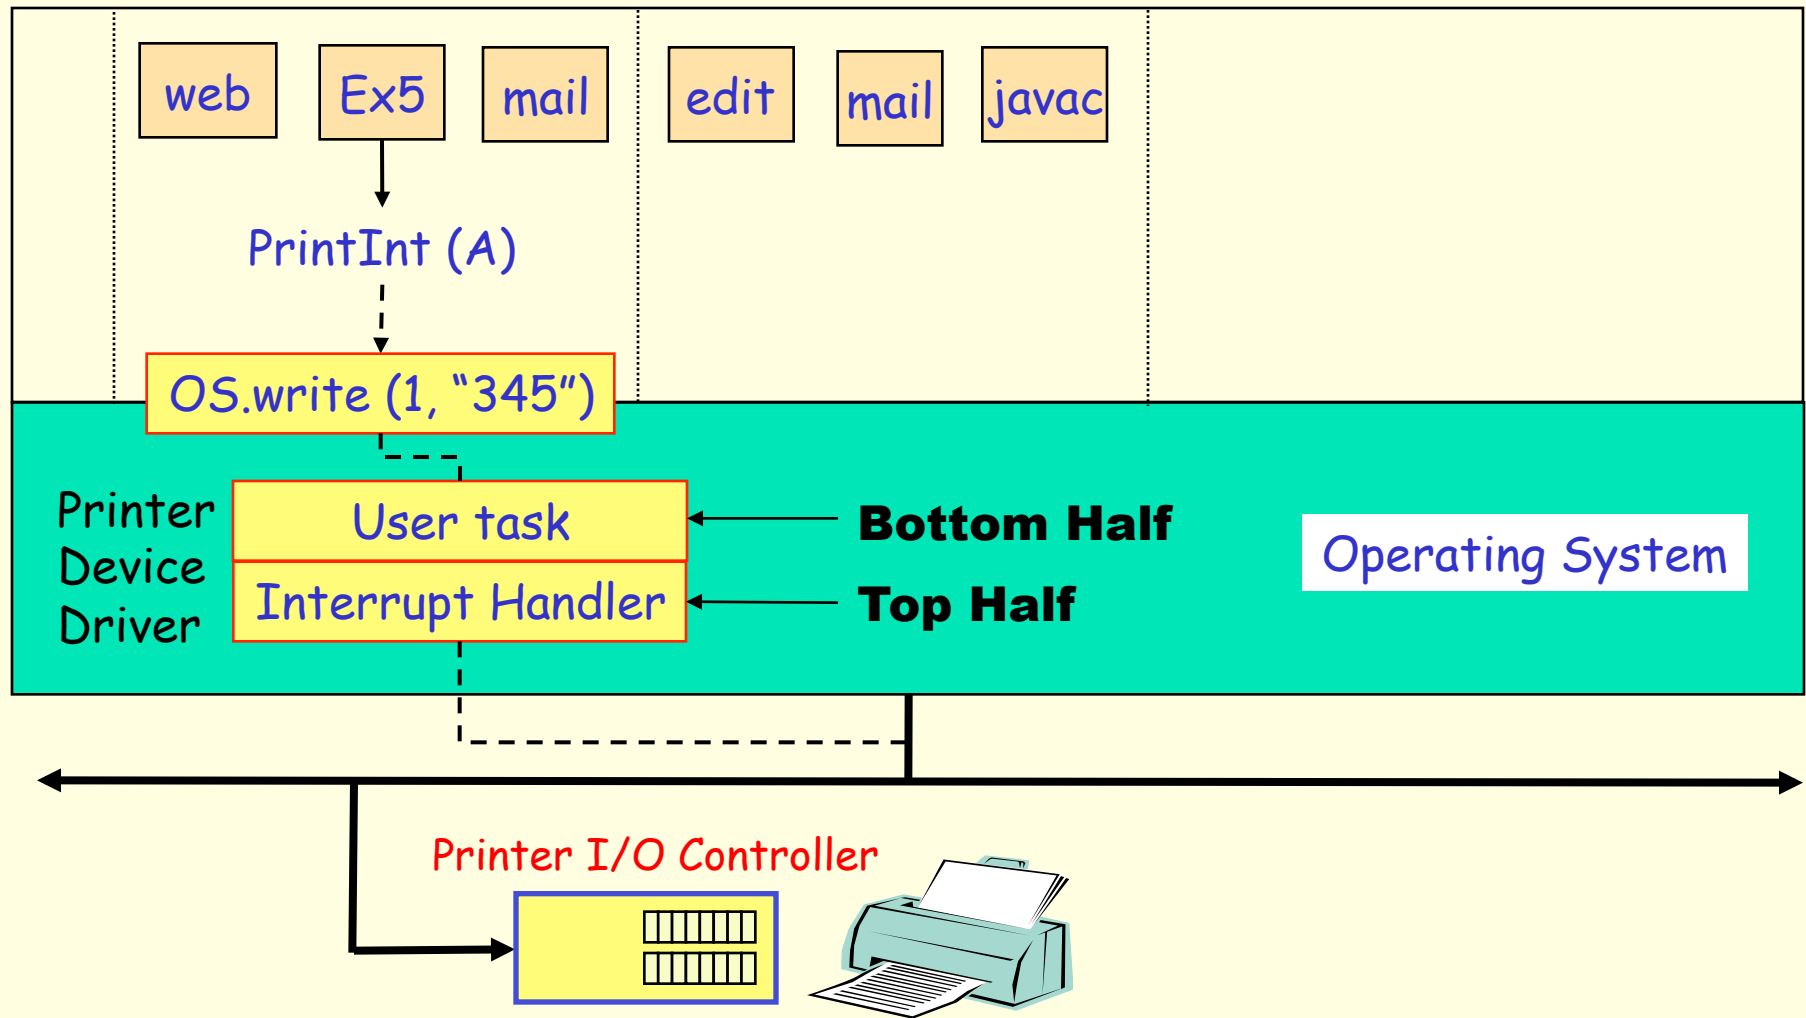
[kin.leung@imperial.ac.uk](mailto:kin.leung@imperial.ac.uk)



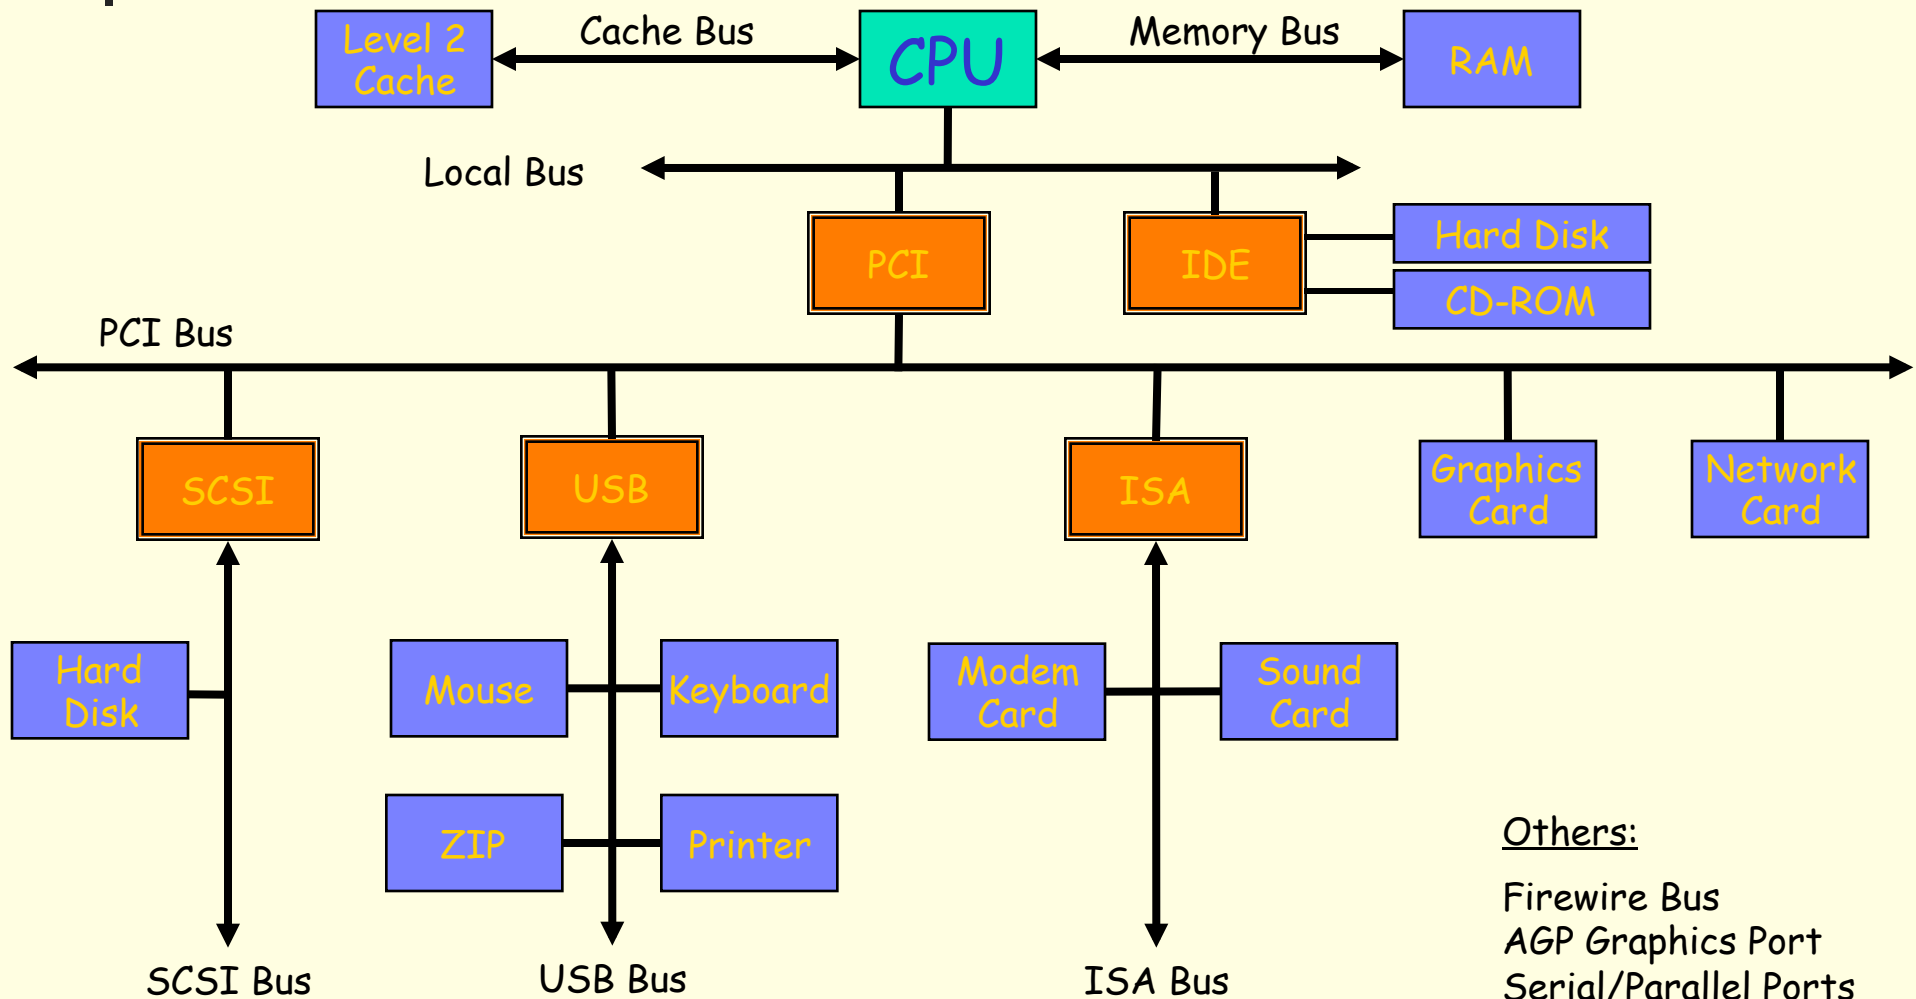
Heavily based on materials from Naranker Dulay



# Example: Introduction



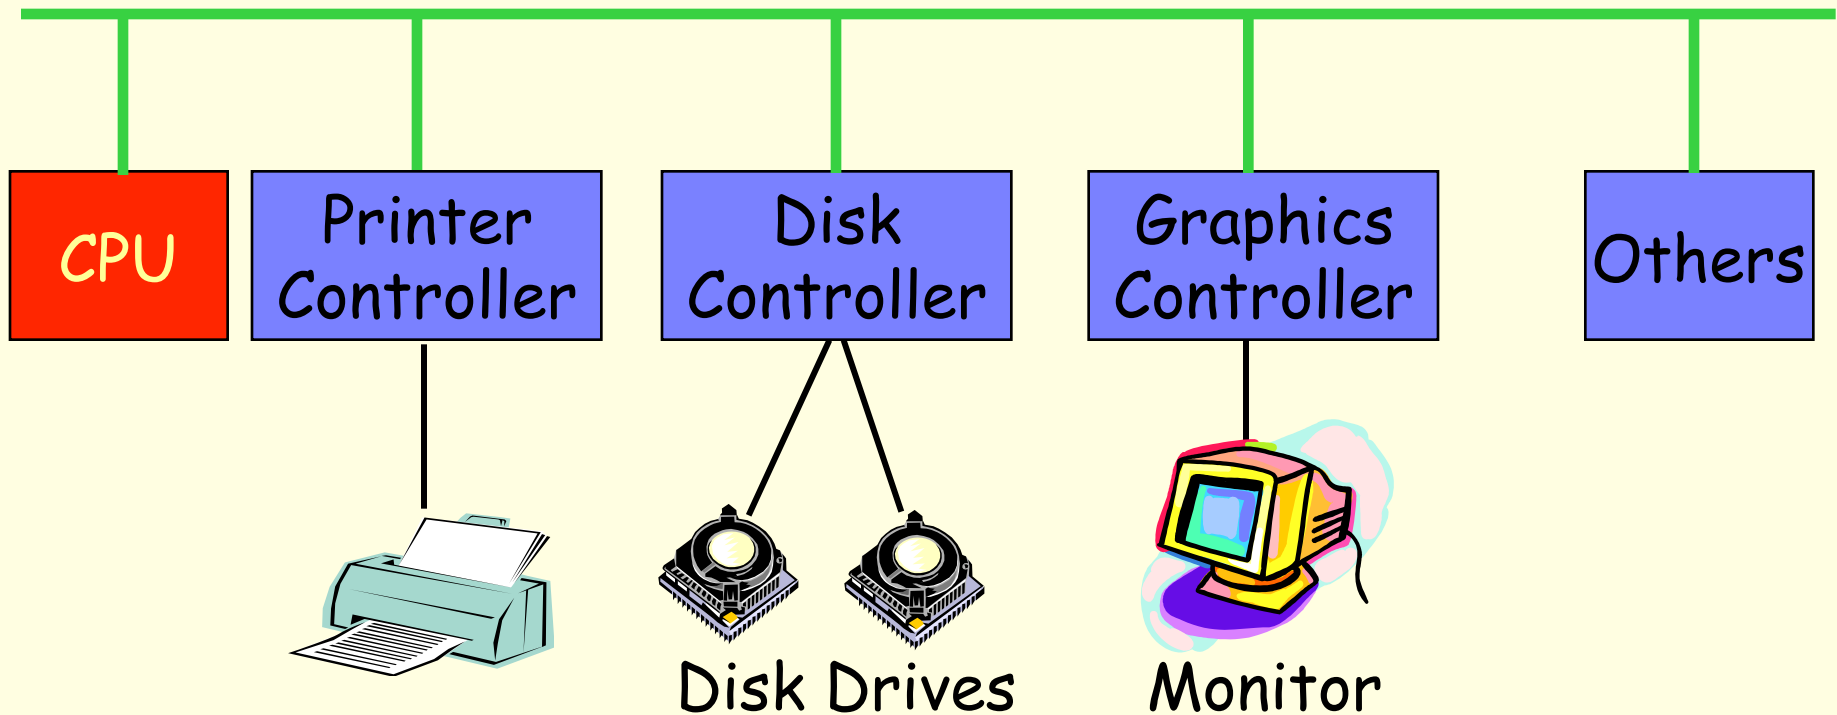
# Pentium I/O Structure



# I/O Controllers


- Sometimes known as I/O Adapters or I/O Modules. Provide the CPU with a programming interface.

## Buses



# I/O Ports


## A Simple I/O Controller




- **Data Ports** are used for passing data to/from the CPU and an I/O device.
- **Control Ports** are used to issue **I/O COMMANDS** (e.g. Write Char) and to check the **STATUS** of a device (e.g. is the device **BUSY** or has an **ERROR** occurred). We write to *specific bits* of a Control Port in order to issue commands, and we read other bits in order to check the status of the device

# I/O Port Addressing 1: Separate I/O Address Space

Main Memory Address Space



I/O Address Space



- In this approach, I/O Ports have their own (very small) I/O address space and the architecture provides some method for accessing them, e.g. special I/O instructions:


**in ax, 20** ; copy 16-bits from I/O port 20 into register AX

**out 35, al** ; copy 8-bits from register AL to I/O port 35

- The Control bus is used to signal if a transfer is for the I/O Address Space or the Main Memory Address Space.
- Pentium provides 64K 8-bit I/O ports numbered 0 to 65535

## I/O Port Addressing 2: Memory-Mapped I/O

Main Memory Address Space



- In this approach, I/O ports "appear" to the programmer as normal memory locations.
- Any instruction that accepts a memory operand can be used to manipulate an I/O port, e.g. MOV, AND, OR, TEST.

**or [80004444H], al**

**; Set bits in I/O Port byte at 80004444H**

- The Pentium supports both Memory-Mapped I/O and I/O Port Addressing via a separate I/O address space.



# Four I/O Schemes

---

- **Programmed I/O**  
Continually “poll” a device’s control port until the device is ready, then initiate transfer.
- **Interrupt-Driven I/O**  
Initiate transfer and then do something else. Device will “interrupt” the CPU when transfer is complete.
- **DMA I/O**  
Initiate large data (block) transfer. Device will transfer block to/from main memory and then “interrupt” the CPU after block is transferred.
- **I/O Processor**  
Delegate complex I/O processing tasks to a dedicated processor.





# Programmed I/O

---

- Check Control Port before reading/or writing to Data Port
- **Example:** Writing a block of data to an I/O Device

```
foreach byte in buffer do  
    loop    Read Status Bit(s) of CONTROL Port  
            if Status Bit(s) indicate ERROR then "Handle" Error  
            exit when Status Bit(s) indicate DEVICE is "READY"  
    endloop  
    Copy byte from Buffer to DATA Port  
    Issue Write Request by writing to Command bits in CONTROL Port  
endfor
```

- **Note:** for some Output (Input) devices the mere act of writing (reading) to (from) the Data Port initiates a new output (input) operation.



# Summary of Programmed I/O

---

- Simple to Program
- Can guarantee Response times
- Poor CPU utilisation (we waste CPU time **Busy Waiting**)
- Multiple Devices are awkward to handle



# Interrupt-driven I/O

---

- Initiate transfer and then do something else. Device will “interrupt” CPU when transfer is complete. On detecting an interrupt, CPU transfers control to device’s interrupt handling procedure (interrupt handler).
- The CPU’s Fetch-Execute cycle now becomes:

**loop**

Fetch Instruction

Decode Instruction

Execute Instruction

**if** INTERRUPT pending **then**

“CALL” Interrupt Handler for this particular Interrupt

**endif**

**endloop**



# Interrupt Processing

## CPU

Complete the current instruction  
Identify the interrupt  
Save "Processor State" e.g. Push  
Registers such as EIP & EFLAGS  
Call Interrupt Handler for  
Identified interrupt

## Interrupt Handler

Process Interrupt (PROGRAMMER)  
Return from Interrupt restoring  
processor state, i.e. Pop  
Registers

Continue with next instruction



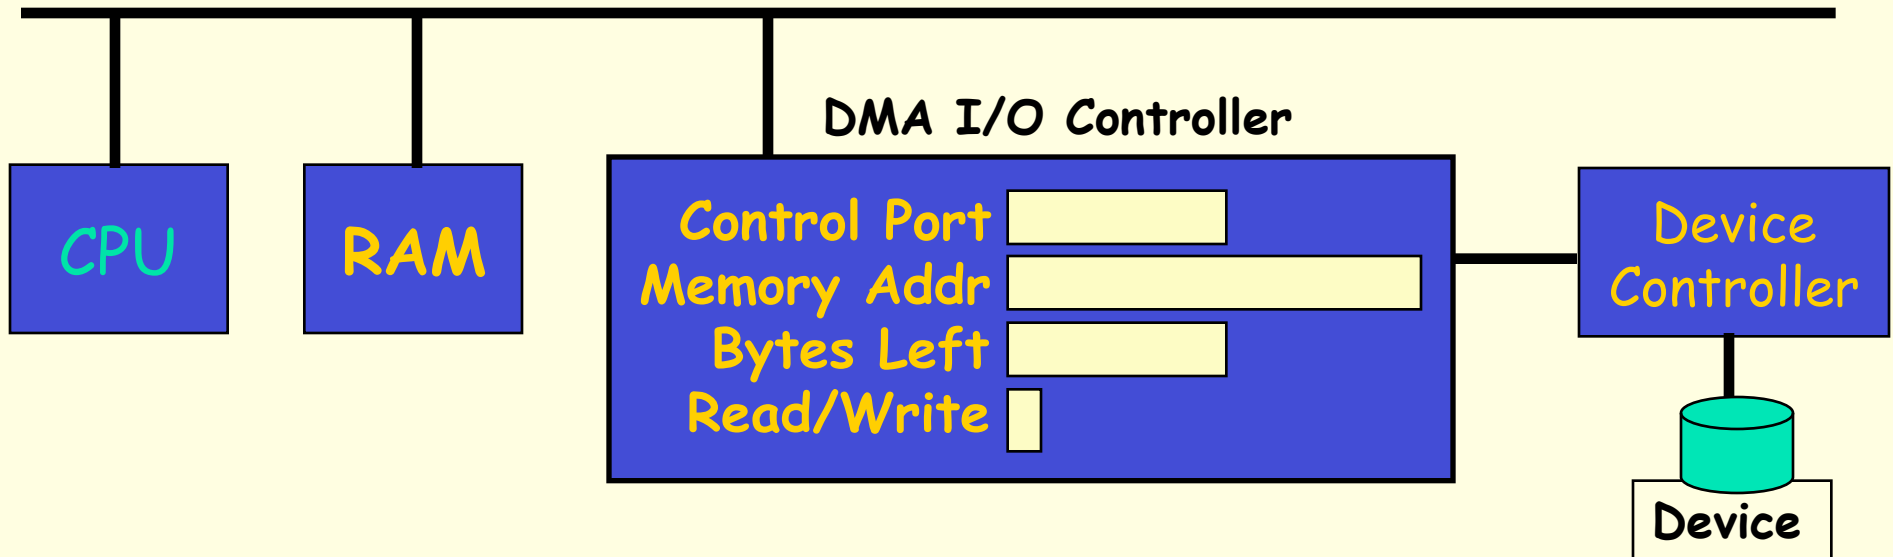
# Summary of Interrupt Driven I/O

- Big step forward compared with programmed I/O
- But Interrupt-Processing Time is relatively Expensive
- Bad for high-speed, high-data volume devices that might lose data if they are not serviced quickly-enough.
- Bad if many Devices continually require attention.

We need to be able to reduce the number of interrupts generated

# DMA I/O for Block Transfers

- CPU writes **start address** of block, **number of bytes** of block and **direction of transfer** to DMA's I/O ports and issues Start Command.
- DMA controller transfers block of data between the Device and Main Memory without the direct intervention of the CPU.
- On completion the DMA Controller interrupts CPU.





# Cycle-Stealing

---

- If the CPU and a DMA controller attempt to access memory at the same time, then the DMA controller will always be given priority (why?). This is known as **Cycle-Stealing**.

|                | Time to perform an N-byte transfer              |
|----------------|---|
| Interrupt I/O: | Time for 1 Interrupt * N                        |
| DMA I/O:       | Time for 1 Interrupt + (N * Memory Access time) |



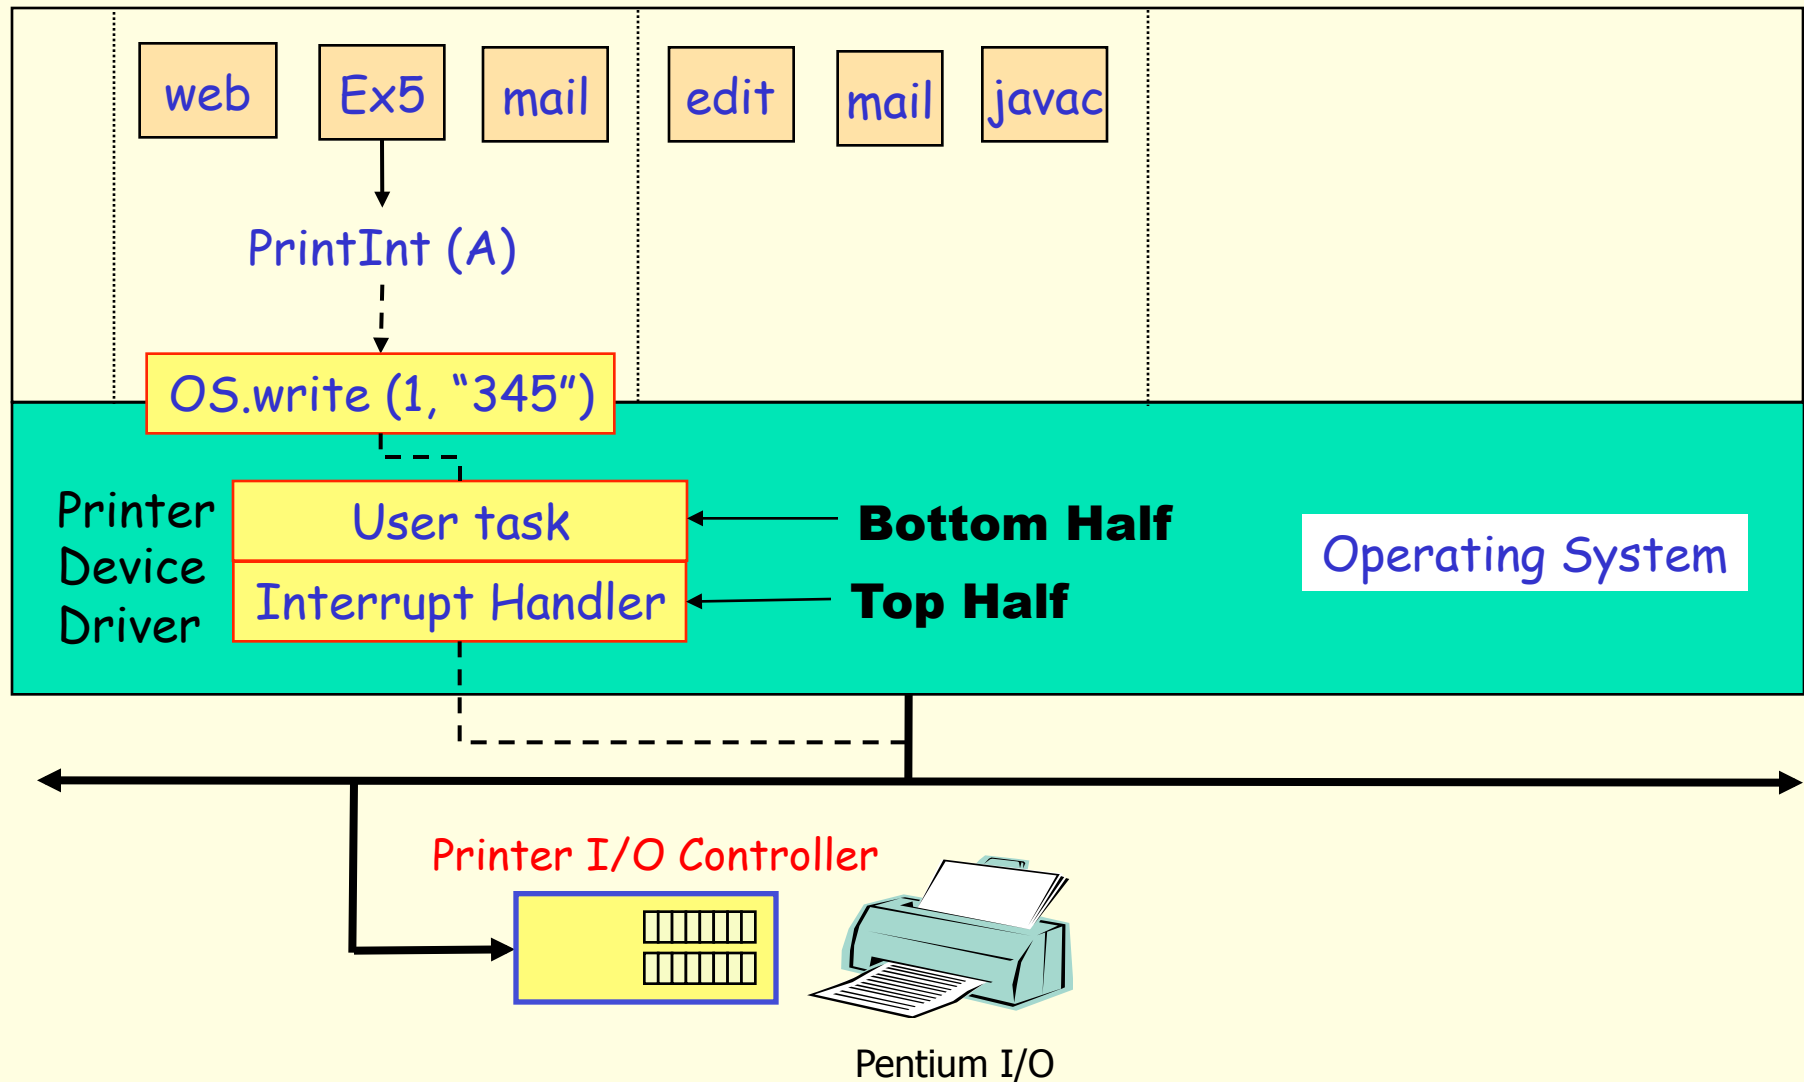
# I/O Processors

---

- Sometimes even DMA I/O is in-adequate.
- A more powerful approach is to use one or more dedicated I/O processors to relieve the CPU of I/O related tasks.
- I/O processors are more capable than DMA controllers and are often full-featured CPU's with their own local memory for buffering data and executing I/O related tasks.



# Device Drivers





## Example: Printer Device Driver (Outline)

- To write a string to the Printer. String terminated by a zero byte.

### Bottom-Half (User task)

**Copy 1st char** from string to Printer's **Data Port**

**Issue Write Request** by writing 1 to Printer's **Control Port** W bit

**OS.SUSPEND** (suspend bottom-half & run another process)

...

Return from Top-Half

### Top-Half (Interrupt Handler)

**IF not *end-of-string*** {

**Copy next char** from String to **Data Port**

**Issue Write Request** by writing 1 to **Control Port's** W bit

} **ELSE** {

**OS.RESUME** bottom-half (when convenient)

}

Return from Interrupt



# Questions

---

- How does the CPU know **which** interrupt handler to call?
- How is an interrupt handler actually **called**?
- What happens when we **return** from an interrupt handler?
- How is data passed between the top-half and bottom-half of the device driver?
- What happens if another interrupt is signalled while the interrupt-handler is executing?



## Our Printer's I/O Ports

|              | 7 | 6 | 5 | 4 | 3 | 2 | 1 | 0 |           |
|--------------|---|---|---|---|---|---|---|---|-----------|
| Control Port | E |   | W |   |   |   |   |   | 80004444H |
| Data Port    |   |   |   |   |   |   |   |   | 80004445H |

### STATUS (checked by reading the Control Port's W and E bits)

|         |                       |
|---------|-----------------------|
| Bit W=0 | Printer is idle       |
| Bit W=1 | Printer is busy       |
| Bit E=0 | No Error (All's well) |
| Bit E=1 | An Error has occurred |

### COMMANDS (issued by writing to the Control Port's W bit)

|         |  |
|---------|--|
| Set W=1 | Causes ASCII Character in Data Port to be printed. On completion I/O controller will reset bit W to 0 and generate an Interrupt (see later). |
|---------|--|



## Bottom-Half of our Printer Driver

```
controlport equ 80004444h ; address of control port
dataport equ 80004445h ; address of data port
string resb 8192 ; 8Kb string buffer
strptr resd 1 ; ptr to chars in string buffer
```

---

```
bottomhalf: mov eax, string ; get address of string
             mov [strptr], eax ; save pointer to 1st char
             mov al, [eax]; ; get 1st char
             jz skip ; skip if char is zero-byte
             mov [dataport], al ; else copy char to data port
             or [controlport], 20H ; and issue write request
             call OS.SUSPEND ; suspend bottom-half
skip: ret ; return from bottom-Half
```

# Interrupt Handler for Printer Driver

```
handler:      sti                                ; re-enable interrupts

              push    eax                        ; save registers used in handler
              inc     [strptr]                  ; advance to next char
              mov     eax, [strptr]             ; copy pointer to register
              mov     al, [eax]                 ; and get char
              jz      endofstr                  ; skip if end of string

              mov     [dataport], al            ; copy char to data port
              or      [controlport], 20h        ; issue write request
              jmp     exit

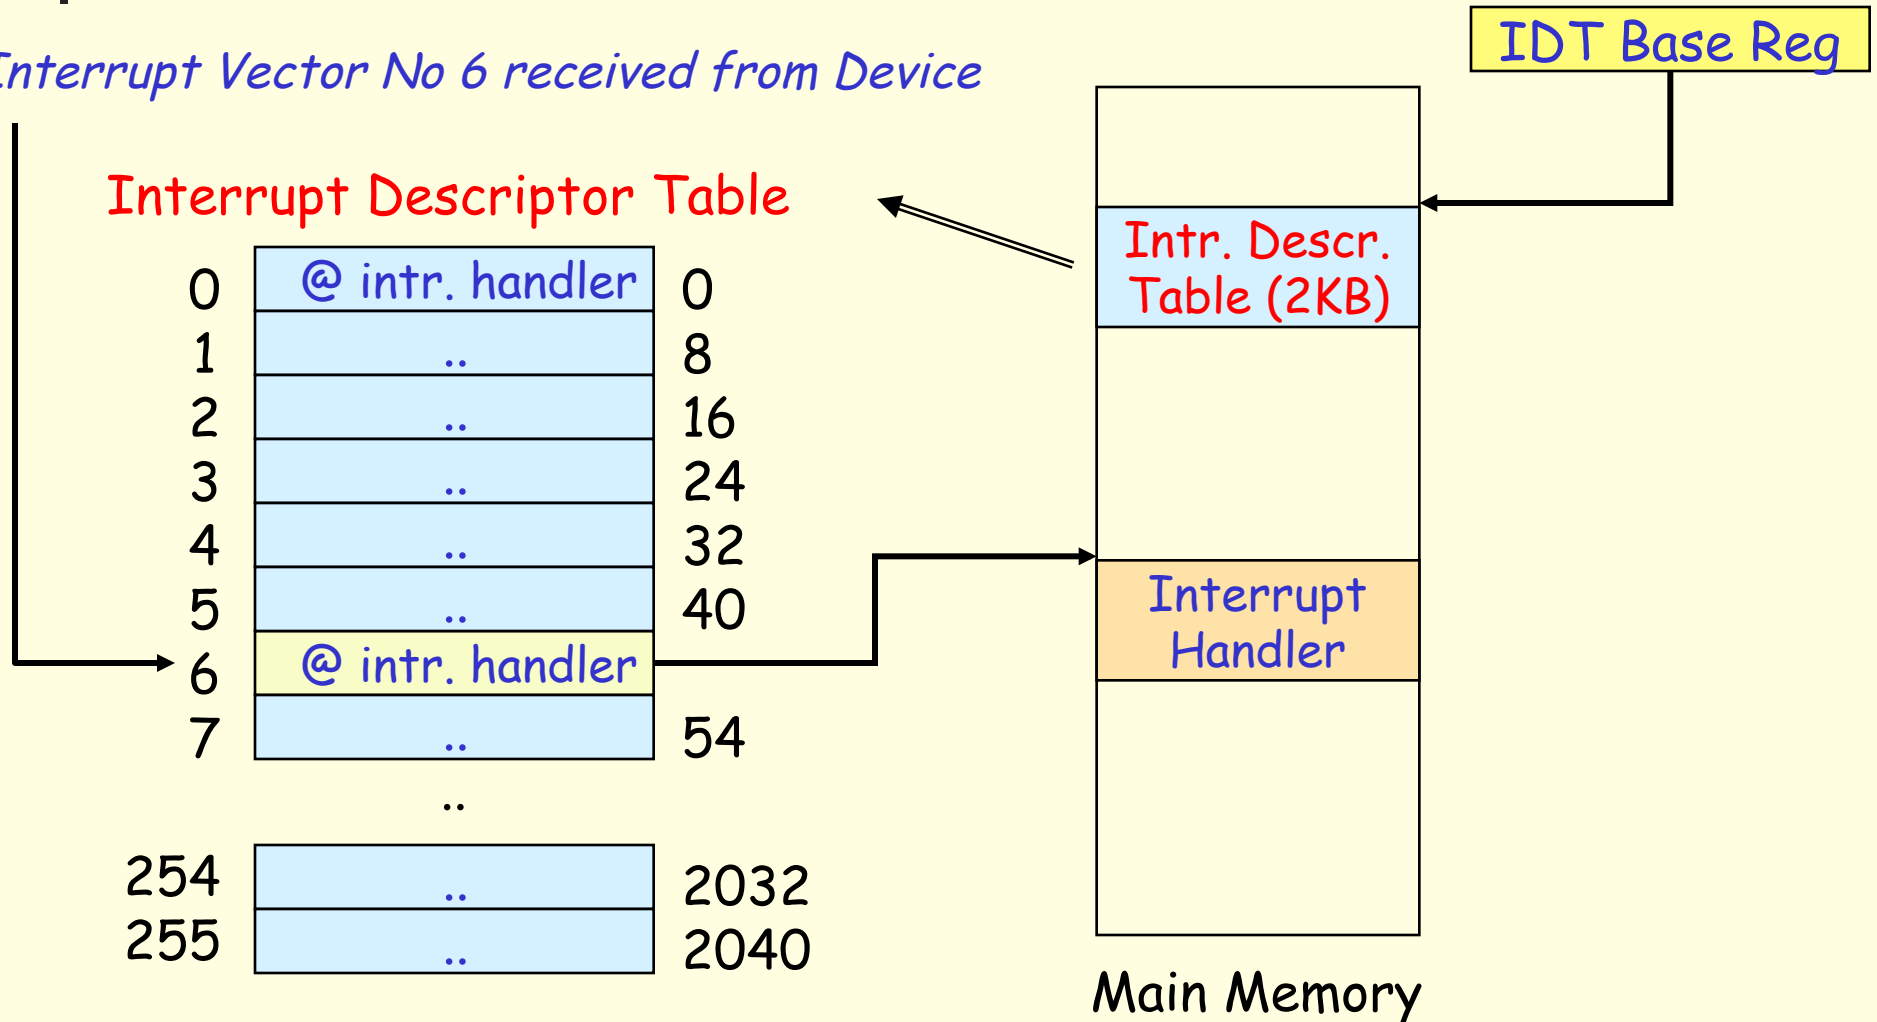
endofstr:     call    OS.RESUME                 ; ask O/S to resume bottom half
exit:        pop     eax

              iret
```

```
graph TD
    endofstr -- solid blue line --> exit
    OS_RESUME[OS.RESUME] -.-> exit
```

# Locating the Interrupt Handler

*Interrupt Vector No 6 received from Device*





# Locating the Interrupt Handler

- When a device wishes to interrupt the CPU, it sends a special "interrupt" signal to the the CPU, along with an number that identifies the interrupt. This number is called the **INTERRUPT VECTOR NUMBER** and on the Pentium is a number in the range 0 to 255 (unsigned byte).
- The Interrupt Vector Number is used to index an **INTERRUPT DESCRIPTOR TABLE** (IDT) with 256 entries. Each entry of this table is a 64-bit "Descriptor" which includes the Interrupt Handler's start address (32-bits).
- *Actually the start address is 48-bits (the other 16-bits define the memory segment of the interrupt handler). Segmentation will be covered in the 2nd year.*
- The start address of the Interrupt Descriptor Table is held in a special CPU register called the **IDT Base Register (IDTR)**.





# Types of Interrupt

- **External I/O device generated Interrupts** - Asynchronous  
I/O Device sends an **INTERRUPT VECTOR NUMBER** to the CPU via Buses.  
Vector Numbers 32-255 are reserved for External Interrupts.

- **CPU-Generated Interrupts (EXCEPTIONS)** - Synchronous  
E.g. Attempt to execute an illegal operation, divide-by-zero.  
Vector numbers 0 to 18 are reserved for CPU Exceptions.

- **Software-Generated Interrupts** - Synchronous  
Interrupts can also be generated explicitly with a instruction, e.g.


`int 80H` ; Calls interrupt handler for intr. vector number 80H

Such interrupts are commonly used to call an operating system method (a so-called **SYSTEM CALL**) since interrupts cause the privilege level to be switched to kernel-mode from user-mode.

# Calling the Interrupt Handler

## When an Interrupt occurs the Pentium CPU will:

- Complete the currently executing instruction
- *Push the EFLAGS Register onto the Stack*
- *Clear the Interrupt Flag bit in the EFLAGS register - this disables further interrupts*
- Push the Return Address (i.e. contents of the EIP register) onto the Stack
- Call the Interrupt Handler using the entry in the Interrupt Descriptor Table



- To return, the Interrupt Handler executes the iret instruction which restores the state (e.g. EFLAGS and EIP) and jumps to the return address on the stack.  
➤ **It is the programmer's responsibility to restore any other registers <<**



# Enabling & Disabling Interrupts

- The Pentium automatically disables interrupts when calling an Interrupt Handler and restores them when returning from the handler - this prevents other devices from interrupting the handler.
- Interrupts can be (re-) enabled more quickly by setting the **Interrupt Enable Flag IF** (bit 9) in the EFLAGS register with the sti instruction:

**sti** ; enables interrupts i.e. sets interrupt enable flag bit in EFLAGS reg.

This is commonly done early within an Interrupt Handler to allow other I/O devices to have their interrupts serviced quickly.

- To explicitly disable interrupts we can use:

**cli** ; disable interrupts i.e. clear IF bit in EFLAGS register

This is useful where the bottom-half of a device-driver can run **concurrently** with the top-half, and needs to update data shared with the top-half. By enclosing the update instructions with CLI and STI we can ensure that the top-half cannot interrupt the top-half.



# Device Driver Summary

- Device-drivers control I/O devices by reading and writing to the **I/O ports** of the device. With memory-mapped I/O this is akin to reading/writing to memory locations.
  - I/O devices signal completion of I/O requests and errors by sending an **Interrupt Vector Number** to the CPU. This causes the CPU to (asynchronously) call (using the Interrupt Descriptor Table) the Device-driver's **INTERRUPT HANDLER**.
  - The **INTERRUPT HANDLER** services the interrupt i.e. checks for errors and copies data to/from the memory area it shares with the Device-driver's **BOTTOM-HALF**.
  - The **BOTTOM-HALF** of the Device driver runs as a thread within the OS and interacts with the device (via I/O ports), the bottom half (via shared memory) and with user-level processes.
- 
- The handling of I/O devices within a computer system plays a critical role in the operation of a computer. Without robust I/O device drivers a computer system will be unstable, subject to random crashes and data corruption. The need for quality "engineering" of I/O device drivers is extremely high. Writing device drivers can be fun but must be taken very seriously.



# Interrupt Vector Numbers

- The Pentium supports up to 256 Interrupts numbered 0 to 255. This number is called the **Interrupt Vector Number**.

## Examples

| Interrupt Vector No | Cause   |
|---------------------|---|
| 0                   | Divide by Zero attempted  |
| 6                   | Attempt to execute an Instruction with an Illegal opcode field    |
| 12                  | Stack Fault Exception   |
| 17                  | General Protection Violation, e.g. writing to read-only "segment" |
| 32-255              | User-defined, used for interrupts from I/O devices                |



# Software Interrupts (System Calls)

- Programs can generate an interrupt with the **INT** instruction e.g.

**int 80H** ; Calls interrupt handler for vector number 80H

- Software interrupts are commonly used for calling operating system functions (**System Calls**) where the address of the function need not be known.
- Software interrupts are also known as TRAPS.

# LINUX System Calls

- Interrupt 80H is used for > 150 LINUX system calls.
- **System Call number** is passed in register EAX.  
Examples: 1=Exit Program, 3=Read from standard input, 4=Write to standard output.
- Parameters 1, 2, 3, 4, 5 are passed in registers ebx, ecx, edx, esi, edi. Result is passed back in register eax or via reference parameters.
- **EXAMPLE:** Write string to standard output

```
mov  eax, 4           ; Linux system call 4: Write( )
mov  ebx, 1           ; Parameter 1: 1=File descriptor for Std output
mov  ecx, AddrOfString ; Parameter 2
mov  edx, LengthOfString ; Parameter 3
int   80H             ; Call Operating System
```

- We also need to save registers before the call, and pop them afterwards.



# FreeBSD System Calls

- Interrupt 80H is used for FreeBSD system calls. Mostly the same as Linux.
- **System Call number** is passed in register `eax`.
- Parameters to FreeBSD system calls are passed on the stack (in reverse order). The result from the call is returned in register `eax`. An extra doubleword should also be pushed before the `INT`.
- **EXAMPLE:** Write string to Standard Output

```
mov  eax, 4                ; System call 4: Write( )
push dword LengthOfString ; Parameter 3
push dword AddrOfString   ; Parameter 2
push dword 1               ; Parameter 1: 1=File descriptor for Std output
push anyDWORD             ; Push any DoubleWord
int   80H                  ; Call Operating System
add   esp, 16              ; Remove Parameters from Stack
```





# MS-DOS System Calls

- Interrupt 21H (Vector number 33) is used for ~ 90 MS-DOS system calls
- **System call number** is passed in register AH. **Examples:** 1=Read character from standard input, 2=Write character to output, 4C=Exit program
- Parameters are passed in other registers (Documentation defines which registers, for which system call)
- **EXAMPLE:** MS-DOS Function 2 writes a char in register DL to the Standard Output

```
mov  ah, 2           ; MS-DOS system call 2: Write character to Std Output
mov  dl, 'a'         ; Parameter 1: Character to be output
int  21H             ; Call Operating System
```

- We also need to push registers before the call, and pop them afterwards.



# That's all Folks !

---

