

# **Introduction to AI**

## Non-monotonic Reasoning

Francesca Toni

(thanks to Marek Sergot and Kostas Stathis)

Section 10.7, Russell and Norvig book

# Outline

- Classical logic: the qualification problem
- Closed World Assumption
- (non-)Monotonicity
- Non-monotonic – defeasible – reasoning
- Negation-as-failure in logic programming (and Prolog)

We will use the Prolog convention on variables/constant/function symbols. Also clauses/rules will be implicitly universally quantified.

# The Qualification Problem (1)

“All birds can fly...”

$\text{flies}(X) \leftarrow \text{bird}(X)$

“...unless they are penguins...”

$\text{flies}(X) \leftarrow \text{bird}(X), \neg \text{penguin}(X)$

“...or ostriches...”

$\text{flies}(X) \leftarrow \text{bird}(X), \neg \text{penguin}(X), \neg \text{ostrich}(X)$

“...or wounded...”

$\text{flies}(X) \leftarrow \text{bird}(X), \neg \text{penguin}(X), \neg \text{ostrich}(X), \neg \text{wounded}(X)$

“...or dead, or sick, or glued to the ground, or...”

Note: the frame problem (planning) is a special case of the qualification problem

# The Qualification Problem (2)

Let **BIRDS** be the set of sentences about flying birds.

Even if we could list all these exceptions, classical logic would still not allow:

$$\text{BIRDS} \cup \text{bird}(\text{tweety}) \not\models \text{flies}(\text{tweety})$$

Note: we would also have to affirm all the *qualifications*, viz.:

- penguin(tweety)
- wounded(tweety)
- sick(tweety)
- ostrich(tweety)
- dead(tweety)
- glued\_to\_the\_ground(tweety)

...

# Classical logic is inadequate

What we want is a new kind of ‘entailment’:

$\text{BIRDS} \cup \text{bird}(\text{tweety}) \models^* \text{flies}(\text{tweety})$

Namely, from BIRDS and  $\text{bird}(\text{tweety})$  it follows *by default* – in the absence of information to the contrary – that  $\text{flies}(\text{tweety})$ .

This kind of reasoning will be *defeasible*.

# Non-monotonic logics

Classical logic is *monotonic*. For a set of sentences  $S$ :

if  $S \models \alpha$  then  $S \cup X \models \alpha$

Namely, new information  $X$  always preserves old conclusions  $\alpha$ .

Reasoning by default is typically *non-monotonic*. We may have that:

$S \models^* \alpha$  but  $S \cup X \not\models^* \alpha$

Examples:

- $\text{DRINKS} \cup \text{coffee} \models^* \text{tastes nice}$
- $\text{DRINKS} \cup \text{coffee} \cup \text{diesel oil} \not\models^* \text{tastes nice}$
- $\text{BIRDS} \cup \text{bird}(\text{tweety}) \models^* \text{flies}(\text{tweety})$
- $\text{BIRDS} \cup \text{bird}(\text{tweety}) \cup \text{penguin}(\text{tweety}) \not\models^* \text{flies}(\text{tweety})$

There have been huge developments in AI over the last 30 years in non-monotonic logic for default reasoning and other applications.

# The 'Closed World Assumption'

Consider the set of sentences that could be describing a database DB:

has-office-in(IBM, Winchester)	city(Winchester)
has-office-in(IBM, London)	city(London)
has-office-in(IBM, Paris)	city(Paris)
has-office-in(MBI, London)	city(NewYork)
capital-city(London)	company(IBM)
capital-city(Paris)	company(MBI)

Does MBI have an office in Paris?

Does IBM have an office in New York?

*We do not know.* But it is usual in databases (and many other contexts) to make a *Closed World Assumption* (CWA):

if  $\alpha$  is not in the DB, assume  $\neg\alpha$

# Credulous vs Sceptical reasoning: “The Nixon diamond”

- Quakers are typically pacifists.
- Republicans are typically not pacifists.
- Richard Nixon is a Quaker
- Richard Nixon is a Republican

Do we conclude that Nixon is a pacifist or not?

*Sceptical (or cautious) reasoning:*

No, since we cannot choose between two possible, conflicting default conclusions.

*Credulous (or brave) reasoning:*

Yes, to both, since we have reason to believe both.

Which form of reasoning to choose? It depends on the application.



# Another example

- Alcoholics are typically adults
- Adults are typically healthy

Do we conclude:

- Alcoholics are typically healthy?

# Normal logic programs

A normal logic program is a set of clauses of the form:

$$A \leftarrow L_1, \dots, L_n \quad (n \geq 0)$$

where  $A$  is an atom and each  $L_i$  is a literal.

A literal is either an atom (a 'positive literal') or of the form

$\text{not } B$

where  $B$  is an atom. ( $\text{not } B$  is a 'negative literal').


The atom  $A$  is the head of the clause; the literals  $L_1, \dots, L_n$  are the body of the clause. When the body is empty ( $n = 0$  above) the arrow  $\leftarrow$  is usually omitted.

$\text{not}$  is negation as failure (NAF):

$\text{not } B$  succeeds when all attempts to prove  $B$  fail. (There are more precise ways of saying this, but this will do here).

# Negation-as-failure: Operational Semantics (1)

We have already seen the computation of a goal (query)  $G = L_1, \dots, L_m$  as a series of derivation steps:



- $\theta_i$  is the mgu at the  $i$ -th derivation step
- The answer computed is (the restriction to the vars of  $G$  of) the composition of all these mgus:

$$\theta = \theta_1 \circ \dots \circ \theta_n$$

# Negation-as-failure: Operational Semantics (2)

Now there are two kinds of derivation steps:

(a) select a positive literal  $L_j = B$  from the current goal:

$$\begin{array}{l} \leftarrow L_1, \dots, L_{j-1}, B, L_{j+1}, \dots, L_n \\ \theta_i \mid \begin{array}{l} \text{match } B \text{ with } B' \leftarrow M_1, \dots, M_k, \text{ with } B\theta_i = B'\theta_i \\ \leftarrow (L_1, \dots, L_{j-1}, M_1, \dots, M_k, L_{j+1}, \dots, L_n)\theta_i \end{array} \end{array}$$

(b) select a negative literal  $L_j = \text{not } B$  from the current goal:

$$\begin{array}{l} \leftarrow L_1, \dots, L_{j-1}, \text{not } B, L_{j+1}, \dots, L_n \\ \theta_i = \{\} \mid \begin{array}{l} \text{subcomputation:} \\ \text{all ways of computing goal } B \text{ must fail (finitely)} \\ \leftarrow (L_1, \dots, L_{j-1}, L_{j+1}, \dots, L_n) \end{array} \end{array}$$

Note that the NAF sub-computation just checks not B: it does not generate bindings (substitutions) for variables.

# Selection of sub-goals

- Sub-goals can be selected in any order. The answers are the same, whatever the selection rule is, though the efficiency of the computation can change.
- A language like Prolog selects always the leftmost sub-goal in the current goal.
- Strictly, the selection of sub-goals must be safe: it must not pick a negative literal containing a variable. (Most Prolog systems do not implement this!)

# Example


P:  $p \leftarrow q$ , **not**  $r$

$r \leftarrow b$

$r \leftarrow$  **not**  $q$

$q \leftarrow$

Goal:  $p$



**computed answer:  $\{\}$**

# Negation-as-failure is non-monotonic

Consider the set of sentences S:

$p(X) \leftarrow q(X), \text{ not } r(X)$

$q(a)$

$r(b)$

Then  $S \vdash_{\text{NAF}} p(a)$

but  $S \cup \{r(a)\} \not\vdash_{\text{NAF}} p(a)$

In fact, negation-as-failure builds in a kind of Closed World Assumption whereby

$S \cup \{r(a)\} \vdash_{\text{NAF}} \neg p(a)$

# Rules & Exceptions using NAF

It is reasonably straightforward to formulate general rule and exception structures using NAF

## Example

- Typically (by default, unless there is reason to think otherwise,...) a bird can fly.
- Except that dead birds cannot fly

$\text{can fly}(X) \leftarrow \text{bird}(X), \text{not abnormal\_bird}(X)$

$\text{abnormal\_bird}(X) \leftarrow \text{dead}(X)$

## Example

People are innocent by default (unless they can be proven to be guilty)

$\text{innocent}(X) \leftarrow \text{not guilty}(X)$



# The Nixon diamond using NAF (1)

- Quakers are typically pacifists.
- Republicans are typically not pacifists.
- Richard Nixon is a Quaker
- Richard Nixon is a Republican

Do we conclude that Nixon is a pacifist or not?

$\text{pacifist}(X) \leftarrow \text{quacker}(X), \text{not abnormal\_quacker}(X)$

$\text{not\_pacifist}(X) \leftarrow \text{republican}(X), \text{not abnormal\_republican}(X)$

$\text{abnormal\_quacker}(X) \leftarrow \text{not\_pacifist}(X)$

$\text{abnormal\_republican}(X) \leftarrow \text{pacifist}(X)$

$\text{quacker}(\text{nixon})$

$\text{republican}(\text{nixon})$

# The Nixon diamond using NAF (2)

By shortening predicates (pacifist becomes p etc)

$p(X) \leftarrow q(X), \text{ not } ab\_q(X)$

$n\_p(X) \leftarrow r(X), \text{ not } ab\_r(X)$

$ab\_q(X) \leftarrow n\_p(X)$

$ab\_r(X) \leftarrow p(X)$

$q(\text{nixon})$

$r(\text{nixon})$

Goal (for example. Here nixon becomes nix):  $p(\text{nix})$

$\leftarrow p(\text{nix})$

$\leftarrow q(\text{nix}), \text{ not } ab\_q(\text{nix})$

$\leftarrow \text{not } ab\_q(\text{nix})$



$\leftarrow ab\_q(\text{nix})$

$\leftarrow n\_p(\text{nix})$

$\leftarrow r(\text{nix}), \text{ not } ab\_r(\text{nix})$

$\leftarrow \text{not } ab\_r(\text{nix})$



$\leftarrow ab\_r(\text{nix})$

...

**Infinite computation!**

# Answer set programming (1)

Given a set of definite clauses  $P$ :

- the *Herbrand universe* is the set of all ground terms obtained from constants and function symbols in  $P$ , and the *Herbrand base* (**HB**) is the set of all atoms from predicate symbols in  $P$  and terms in the Herbrand universe
- A *Herbrand model* is a subset of the Herbrand base that renders  $P$  true
- The meaning of  $P$  is the *least Herbrand model* (**LHM**) of  $P$

Note: the LHM can be computed bottom-up (a-la-forward chaining)

# Answer set programming (2)

- Given a (ground) normal logic program  $P$  and  $S \subseteq \text{HB}$ , the *reduct of  $P$  by  $S$*  (referred to as  $P^S$ ) is obtained in two steps:
  1. Eliminate all rules with **not  $p$**  in the body, for every  **$p \in S$**
  2. Eliminate all negative literals from the body of all remaining rules
- $P^S$  is a set of definite clauses.
- $S \subseteq \text{HB}$  is an *answer set* of  $P$  iff the LHM of  $P^S$  is  $S$
- Several efficient answer set solvers exist

# Nixon diamond using Answer set programming (1)

$p(X) \leftarrow q(X), \text{ not } ab\_q(X)$

$n\_p(X) \leftarrow r(X), \text{ not } ab\_r(X)$

$ab\_q(X) \leftarrow n\_p(X)$

$ab\_r(X) \leftarrow p(X)$

$q(nix)$

$r(nix)$

Here  $HB = \{q(nix), r(nix), p(nix), \text{not\_}p(nix), ab\_q(nix), ab\_r(nix)\}$ . Let  $P$  be the set of all ground instances of these rules over  $HB$ .

Consider  $S = \{p(nix), ab\_r(nix), q(nix), r(nix)\}$ .

**Is  $S$  an answer set? YES.** To see this, let us construct  $P^S$

1. Eliminates  $n\_p(nix) \leftarrow r(nix), \text{ not } ab\_r(nix)$

2. Gives  $p(nix) \leftarrow q(nix),$

$ab\_q(nix) \leftarrow n\_p(nix), \quad ab\_r(nix) \leftarrow p(nix),$

$q(nix), \quad r(nix)$

The LHM of  $P^S$  is  $\{q(nix), r(nix), p(nix), ab\_r(nix)\} = S$ , as required.

## Nixon diamond using Answer set programming (2)

$S = \{q(nix), r(nix), n\_p(nix), ab\_q(nix)\}$  is also an answer set

- Multiple “extensions” = answer sets
- Credulous (whatever holds in some answer set)  
vs sceptical (whatever holds in all answer sets)

# Summary

- We have discussed:
- Classical logic: the qualification problem
- Closed World Assumption
- (non-)Monotonicity
- Non-monotonic – defeasible – reasoning
- Negation-as-failure for defeasible reasoning
- Answer set programming
- Examples