

# Operating Systems

## Introduction

MSc CO502  
Autumn Term Weeks 7-11

**Morris Sloman & Roman Kolcun**

m.sloman@imperial.ac.uk  
Room 572

# Course Objectives

What is an operating system, and how it supports the implementation of software on a computer.

Understand the features and mechanisms that underlie operating systems, including:

- process and thread management and synchronization
- memory management
- security
- input-output
- file systems

Linux characteristics as a case study

# Outline

## **Morris Sloman** (12 lectures)

- Overview: function and structure
- Processes and Threads: concepts and scheduling
- Process Coordination: synchronization & deadlocks
- Memory Management: allocation and virtual memory
- Security: authentication and access control

## **Roman Kolcun** (6 lectures)

- Input/Output: device drivers, disk management & scheduling
- File Systems: files and directory structures
- Virtual Machines Systems (if time)

# Course Structure

Four lectures + 2 tutorials per week (Weeks 7 – 11)

**Lectures:** Mondays 9am LT145 , Tuesdays 2pm LT145,  
Fridays 2-4pm LT 311

**Tutorials:** Mondays 10am LT 145, Fridays 4pm R341/2

Course slides are on Cate

Acknowledgements:

Slides based on material by Peter Pietzuch, Cristian Cadar  
and Julie McCann

# Recommended Books

- 1. Modern Operating Systems: Global Edition**, A. Tanenbaum, H. Bos, 4th edition, Pearson, 2015
- 2. Operating Systems – Internals and Design Principles**, W. Stallings, 8th Edition, Pearson, 2014
- 3. Operating System Concepts**, A. Silberschatz, P. Galvin, G. Gagne, 8th Edition, John Wiley & Sons, 2014

Note: Earlier editions of these are OK and may be more readily available

☛ Important: Do not just rely on these slides!

# OS Overview

# Computer Architecture Overview

## Processor

- Controls computer hardware
- Executes instructions and programs

## Memory

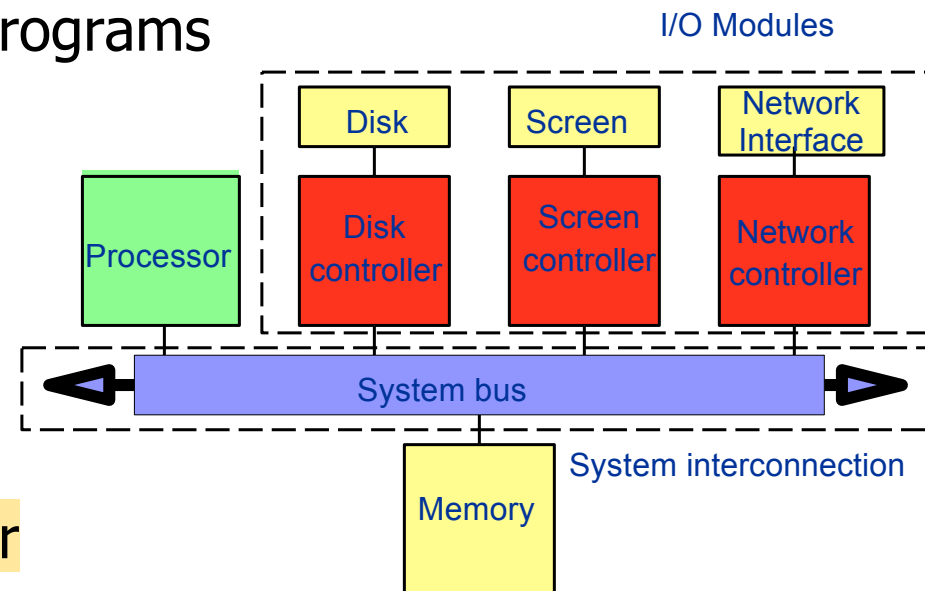
- Stores data and programs

## I/O modules

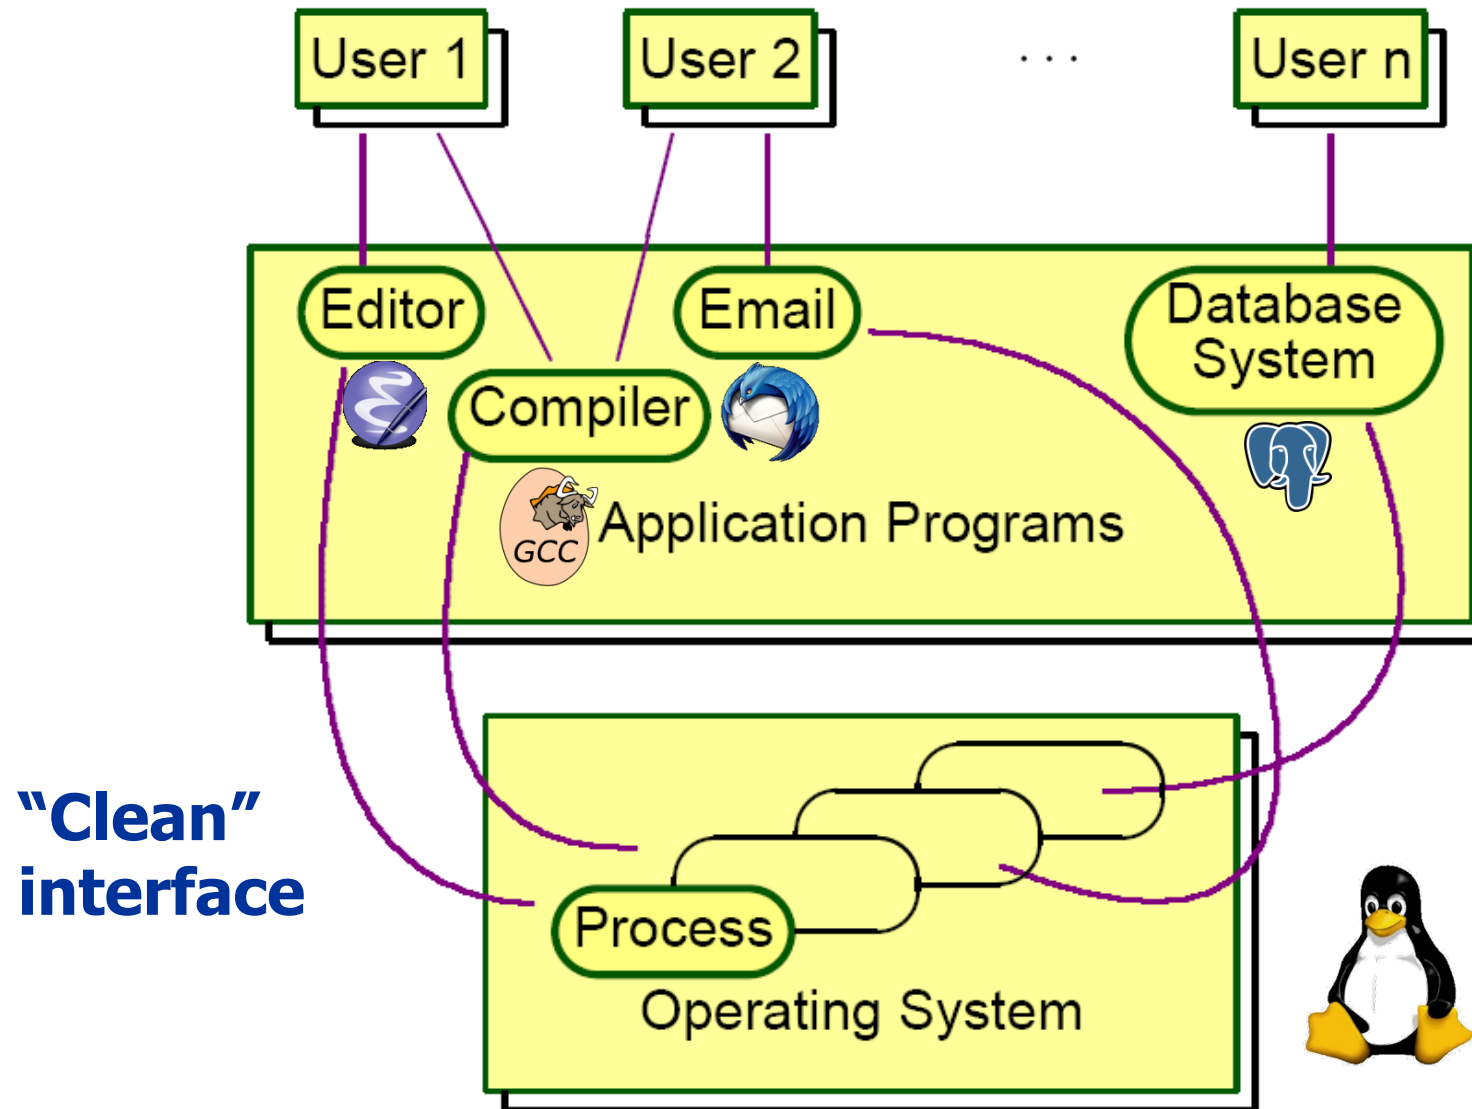
- Read and write from I/O devices
- Intelligence in I/O controller

## System interconnection

- Connects different hardware components via bus
- Provides communication between hardware components

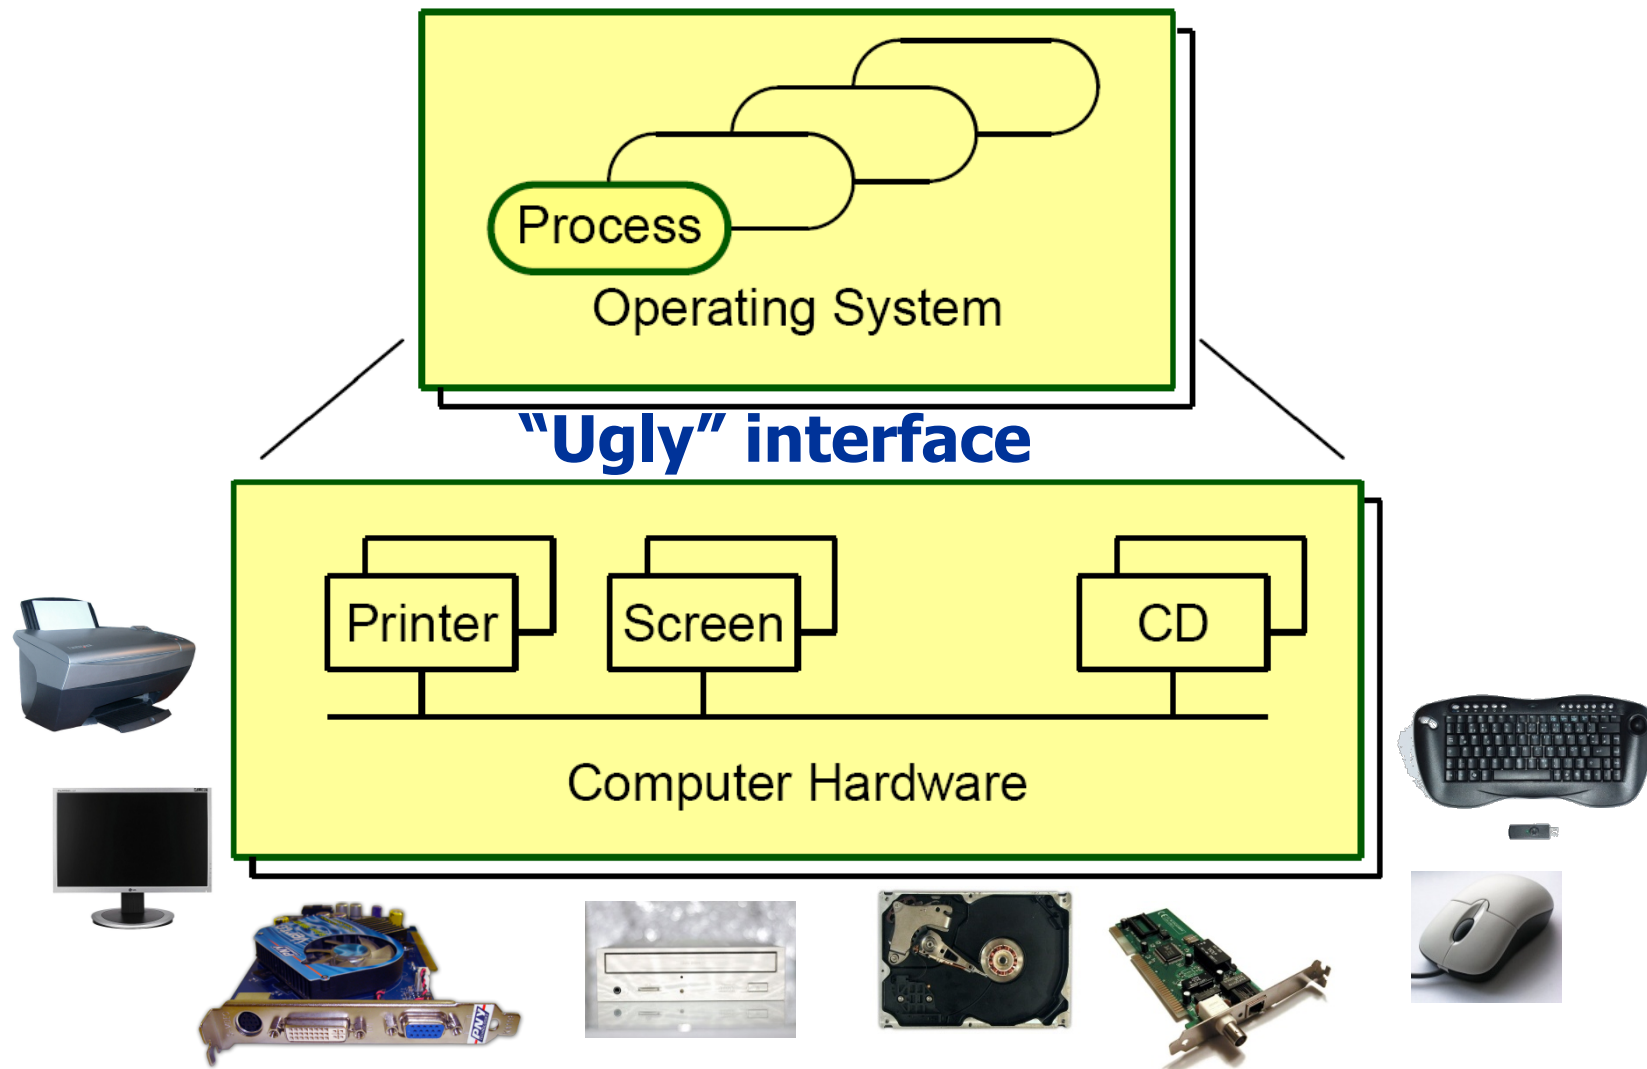


# Operating Systems – Top Level View





# Operating Systems – Bottom Level View



# 1. Resource Management

Making efficient use of (limited) available resources

- Optimise utilisation of processor, memory, disks, network etc....

Sharing resources among multiple users

- Schedule access, fair allocation
- Prevent interference

# Resources

## Processors

- Divide number and/or time

## Memory

- RAM, cache, disks, ...

## Input/Output devices

- Screens, printers, network interface, ...

## Internal devices

- Clocks, timers, accelerometers ...

## Long-term storage (files)

- Disks, storage cards, DVD, tapes, ...

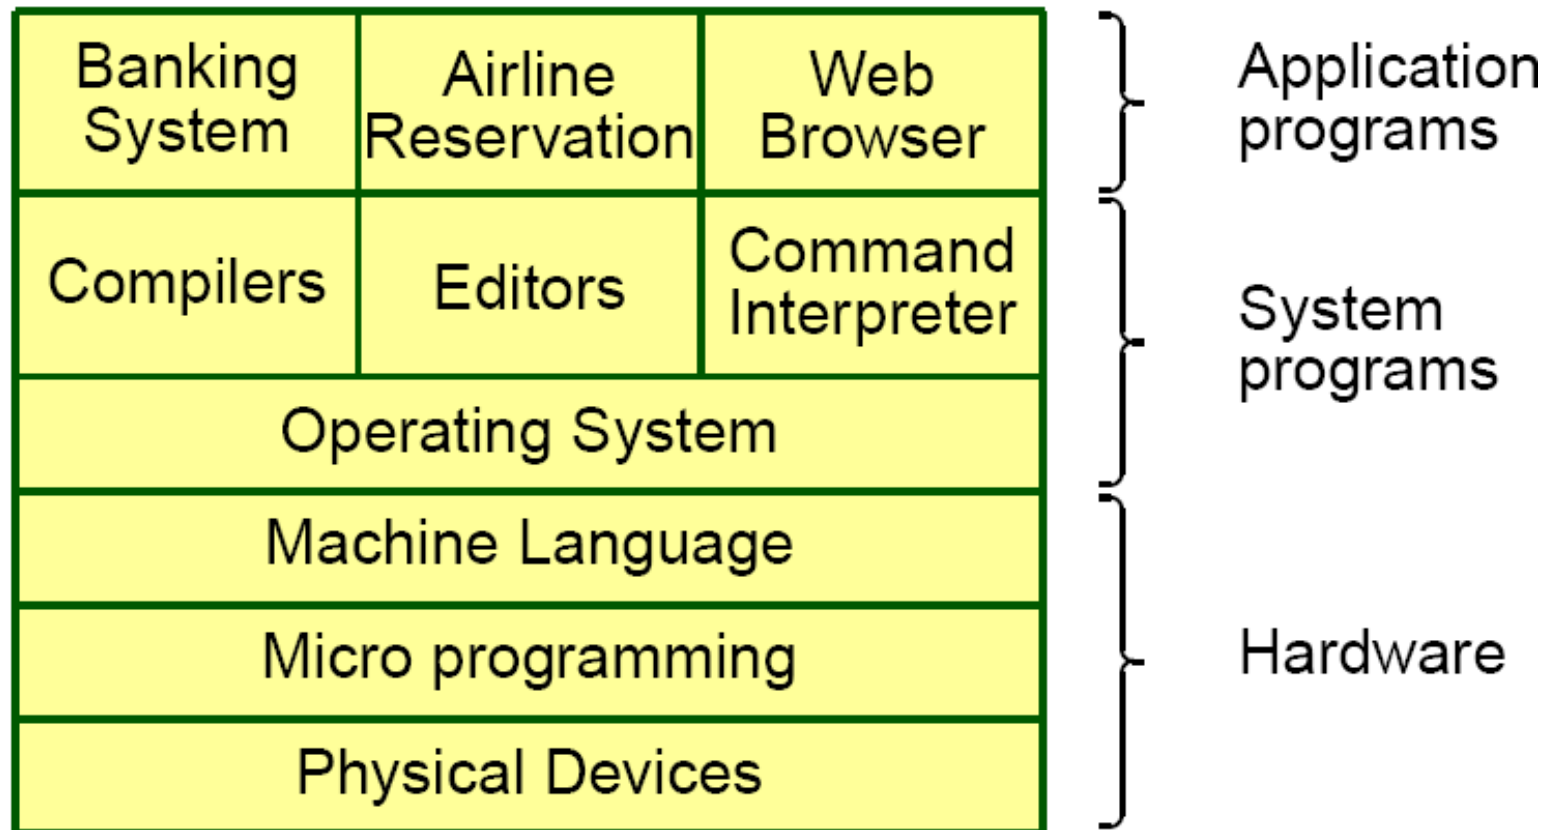
## Software

- Browsers, editors, e-mail clients, databases, .....

## 2. Providing Clean Interfaces

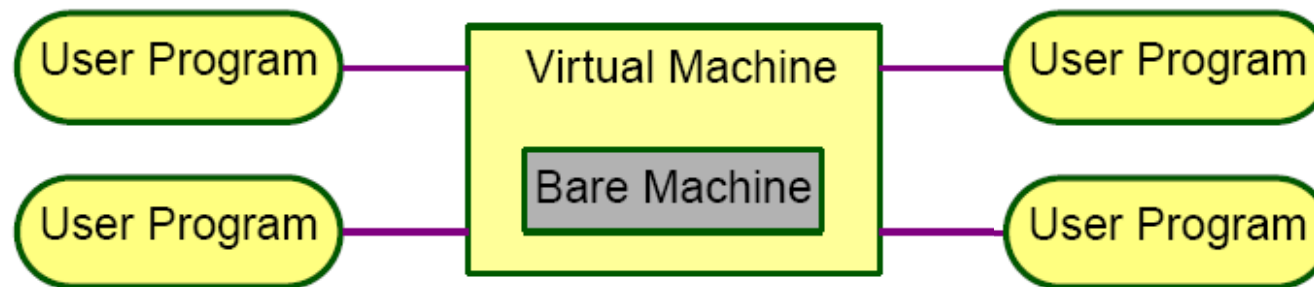
OS converts raw hardware into usable computer system

- Hides complexity of lower levels from higher levels



# Virtual Machine Abstraction

- Details of hardware kept hidden from programs
- Only OS can allow access to hardware resources
- User request should be abstract
  - e.g. no need to know how files stored on disk



# Virtual Machine Facilities

**Simplified I/O:** Device independence; open a file on disk, CD, screen is one operation.

**Virtual Memory:** Larger than real or partitioned.

**Filing System:** Long term storage, on disk or tape, accessed by symbolic names.

**Program Interaction and communication:** Pipes, semaphores, locks, monitors.

**Network communication:** Message passing

**Protection:** Prevent programs accessing resources not allocated to them.

**Program Control:** User interaction with programs, command language, shells.

**Accounting & Management Information:** Usage of processors, memory, file storage etc.

# OS Characteristics: Sharing

## Sharing of data, programs and hardware

- Time multiplexing and space multiplexing

## Resource allocation

- Efficient and fair use of memory, CPU time, disk space, ...
- Simultaneous access to resources
  - Processor, Disks, RAM, code, network, ...
- Mutual exclusion
  - Protect multiple programs from uncontrolled access to shared resources.
  - Prevent multiple writes to same data structure or file.
- Protection against corruption
  - Accidental or malicious

# OS Characteristics: Concurrency I

## Several simultaneous parallel activities

- Overlapped I/O & computation
- Multiple users and programs run in parallel

## Switch activities at arbitrary times

- Guarantee fairness and prompt response
- Differential responsiveness e.g. interactive vs. batch

## Safe concurrency

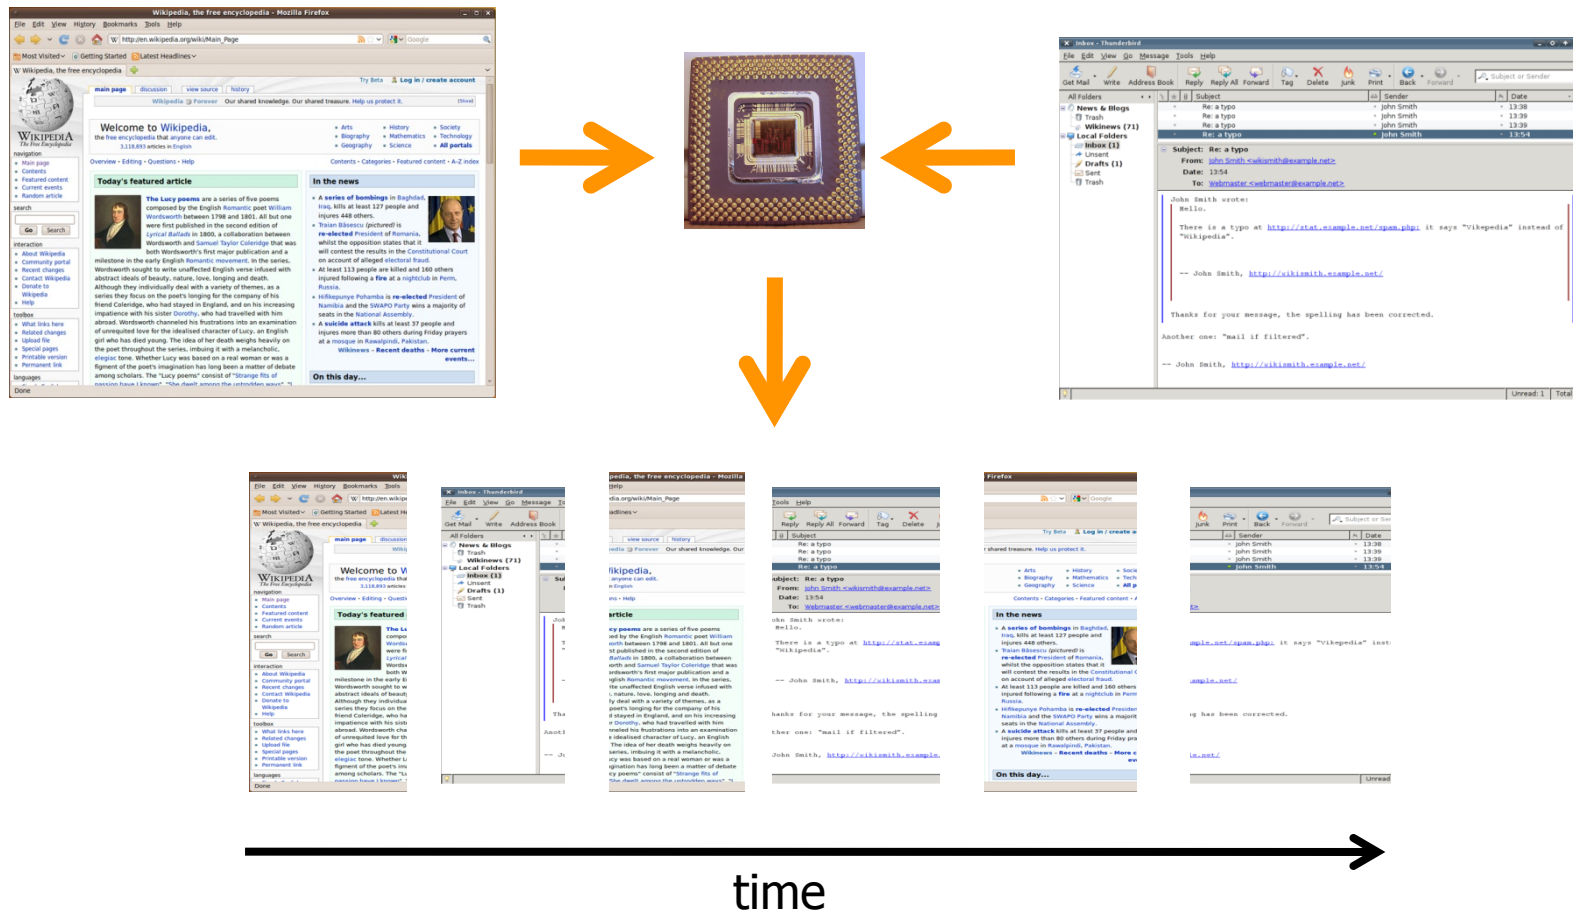
- Synchronisation of actions
  - Avoids long waiting cycles; gives accurate error handling
- Protection from interference
  - Each process has its own space



# OS Characteristics: Concurrency II

## Time-slicing

- Switch application running on physical CPU every 50ms



What other switching mechanisms are feasible?

# OS Characteristics: Non-determinism

## Non-determinism

- Results from events occurring in unpredictable order
  - e.g. timer interrupts, user input, program error, network packet loss, disk errors, . . .
- Makes programming OS hard!

# OS Characteristics: Storing Data

Long term storage: File systems for disks, DVDs, memory cards .....

- Easy access to files through user-defined names
  - Directory structure, links, shared disks
- Access controls
  - Read, write, delete, execute or copy permissions
- Protection against failure (backups)
  - Daily/weekly/monthly, partial/complete
- Storage management for easy expansion
  - Add disks without need for re-compilation of OS

# Operating System Zoo

**Desktop/Laptop** (e.g. Windows, Mac OS X, Linux)

- Typically 2-8 cores  
+ high resolution screen

**Server OS** (e.g. Linux, Windows Server 20XX, Solaris, FreeBSD,)

- Share hardware/software resources e.g. internet servers
- Typically many multicore processors + large disks

**Smartphones** (e.g. iOS, Android)

- Simpler CPUs, starting to be sophisticated

**Real-time OS**

- Guaranteed time constraints

**Embedded OS** (e.g. QNX, VXWorks)

- Transport, communications, banking, homes etc.
- Only trusted software

**Smart card OS**

- Usually single function
- Many have JVM
- OS is primitive

**Sensor Network OS** (e.g. TinyOS)

- Resource/energy conscious

# OS Structure

**Monolithic OS kernels** (e.g. Linux, BSD, Solaris, ...)

- Single black box

**Microkernels** (e.g. Symbian, L4, Mach, ...)

- Little as possible in kernel (fewer bugs)

**Hybrid kernels** (e.g. Windows NT, Mac OS X, ...)

- Take a guess... 😊

# Monolithic Kernels

Kernel is single executable with own address space

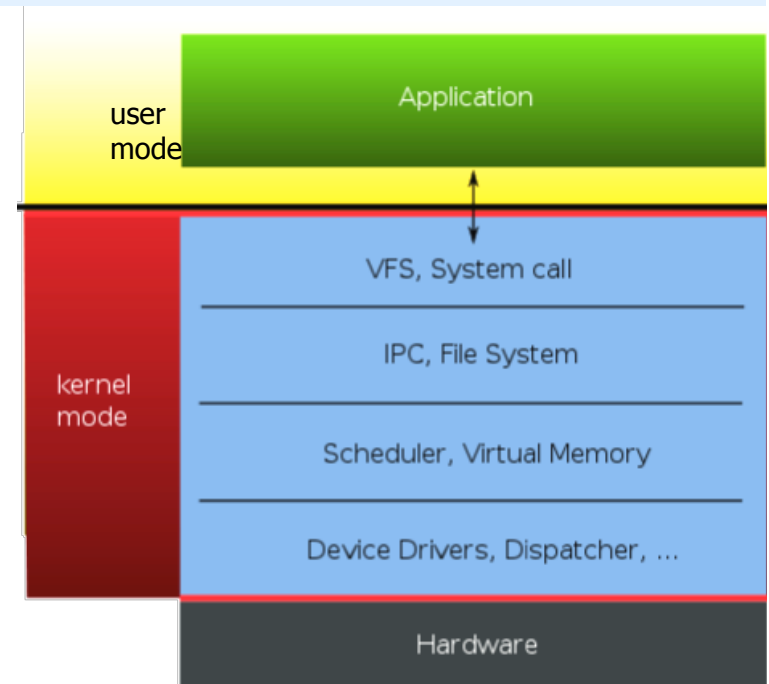
- Structure implied through pushing parameters to stack and trap (systems calls)
- Most popular kernel style

## Advantages

- Efficient calls within kernel
- Easier to write kernel components due to shared memory

## Disadvantages

- Complex design with lots of interactions
- No protection between kernel components



# Microkernels

Minimal “kernel” with functionality in user-level servers

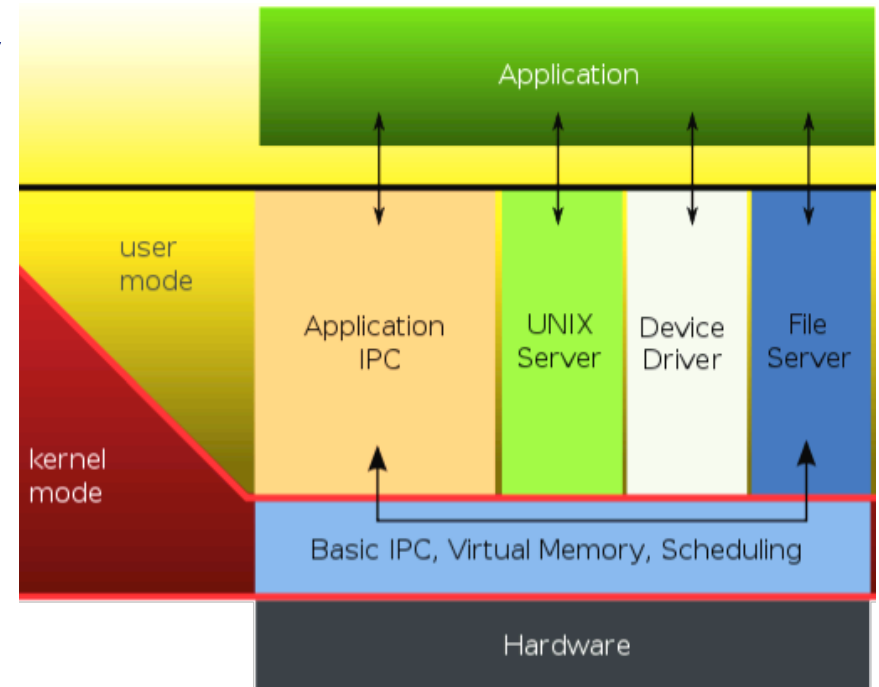
- Kernel does IPC (message-passing) between servers
- Servers for device I/O, file access, process scheduling, ...

## Advantages

- Kernel itself not complex → less error-prone
- Servers have clean interfaces
- Servers can crash and restart without bringing kernel down

## Disadvantages

- Overhead of IPC within kernel high



# Hybrid Kernels

Combines features of both monolithic and microkernels

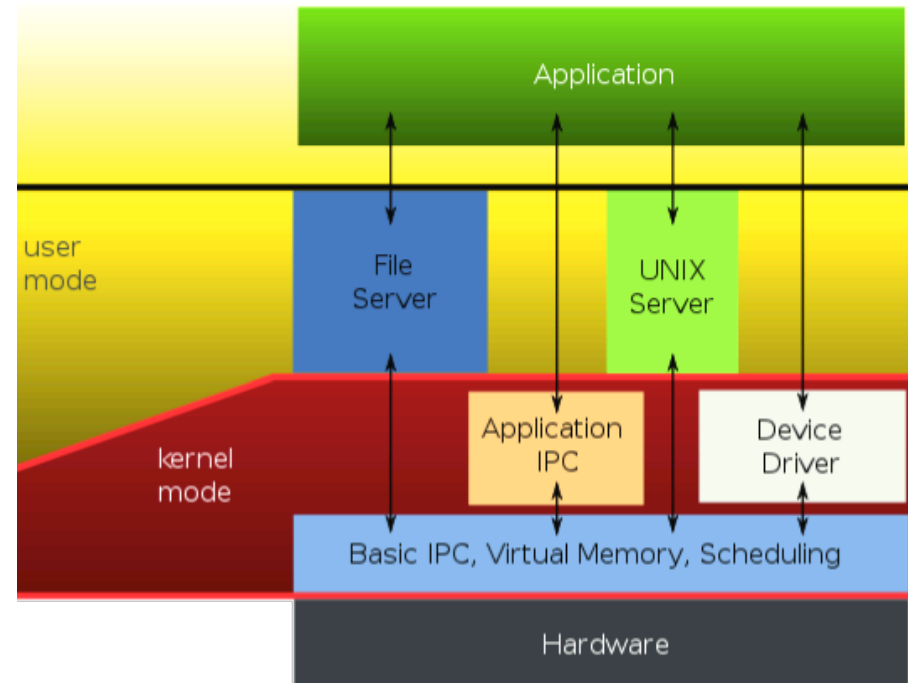
- Often a design philosophy

## Advantages


- More structured design

## Disadvantages

- Performance penalty for user-level servers







# Introduction to Linux

# Linux History and Motivation

Variant of Unix like FreeBSD, System V, Solaris etc.

- Ken Thomson left Multics (Bell Labs)
  - *Uniplexed* information and computing service
- Dennis Ritchie got interested

Late 80's: 4.3 BSD and System V r3 dominant

- Systems call libraries reconciliation POSIX

1987 Tanenbaum released MINIX microkernel

- Tractable by single person (student)

Linus Torvalds, frustrated, built fully-featured yet monolithic version → Linux

- Major goal was interactivity, multiple processes and users
- Code contributed by world-wide community

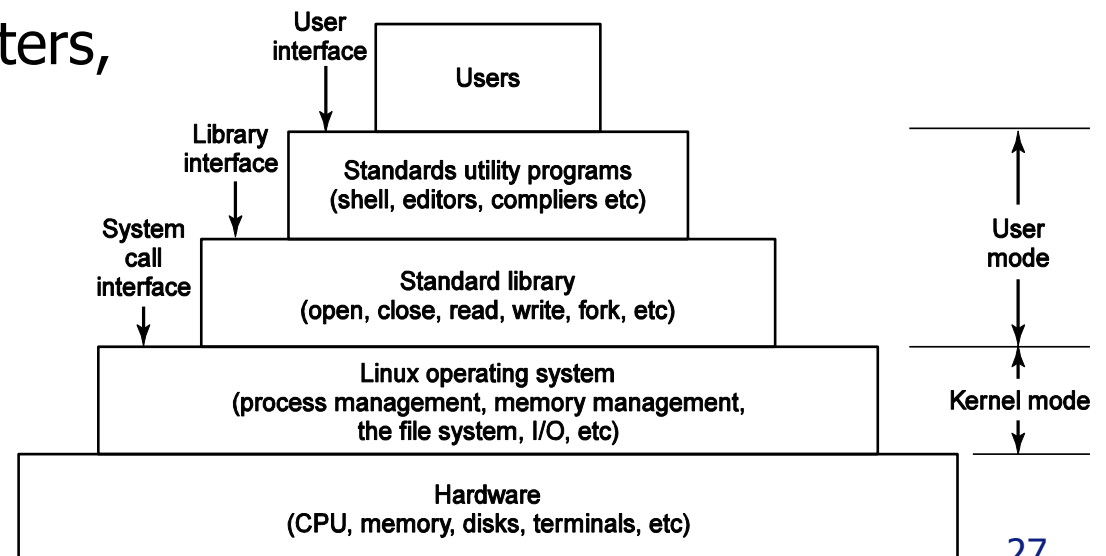
# Structure and Interfaces

## System calls

- Implemented by putting arguments in registers (or stack)
- Issue trap to switch from user to kernel

## Rich set of programs (through GNU project)

- e.g. shells (bash, ksh, ...), compilers, editors, ...
- Desktop environments: GNOME, KDE, ...
- Utility programs: file, filters, editors, compilers, text processing, sys admin, etc



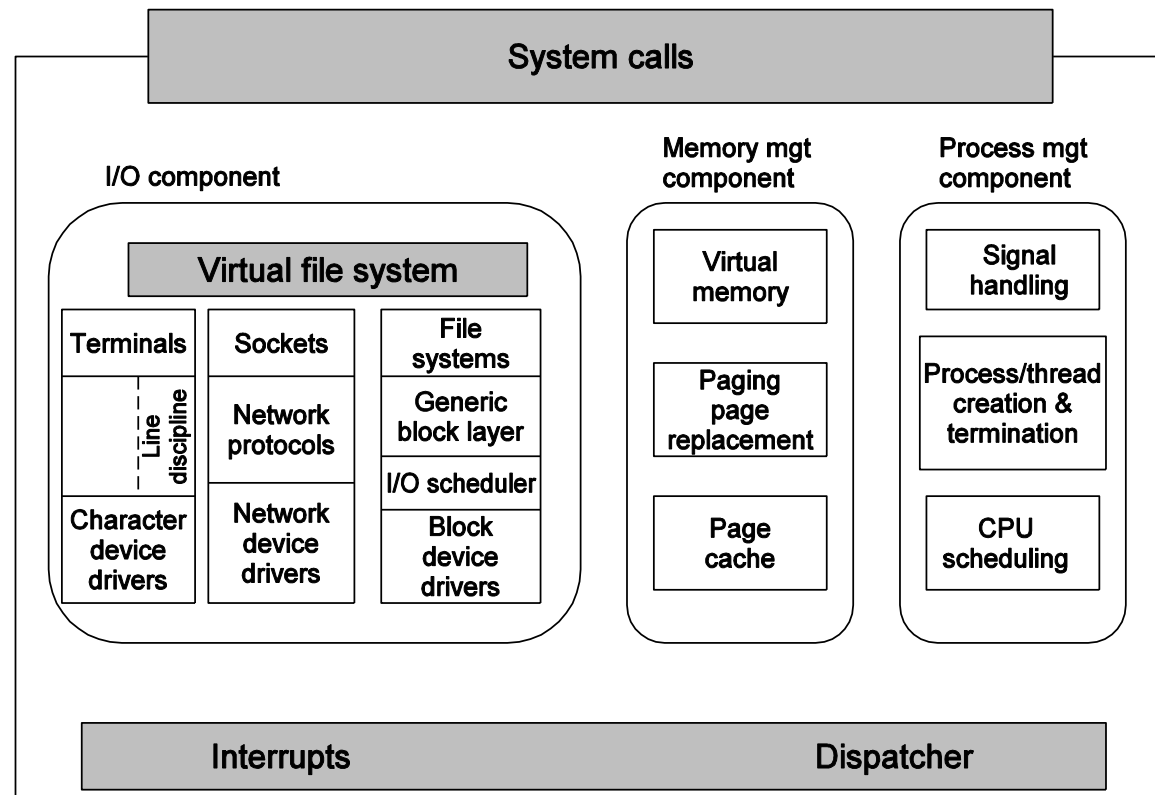
# Kernel Structure

Interrupt handlers primary means to interact with devices

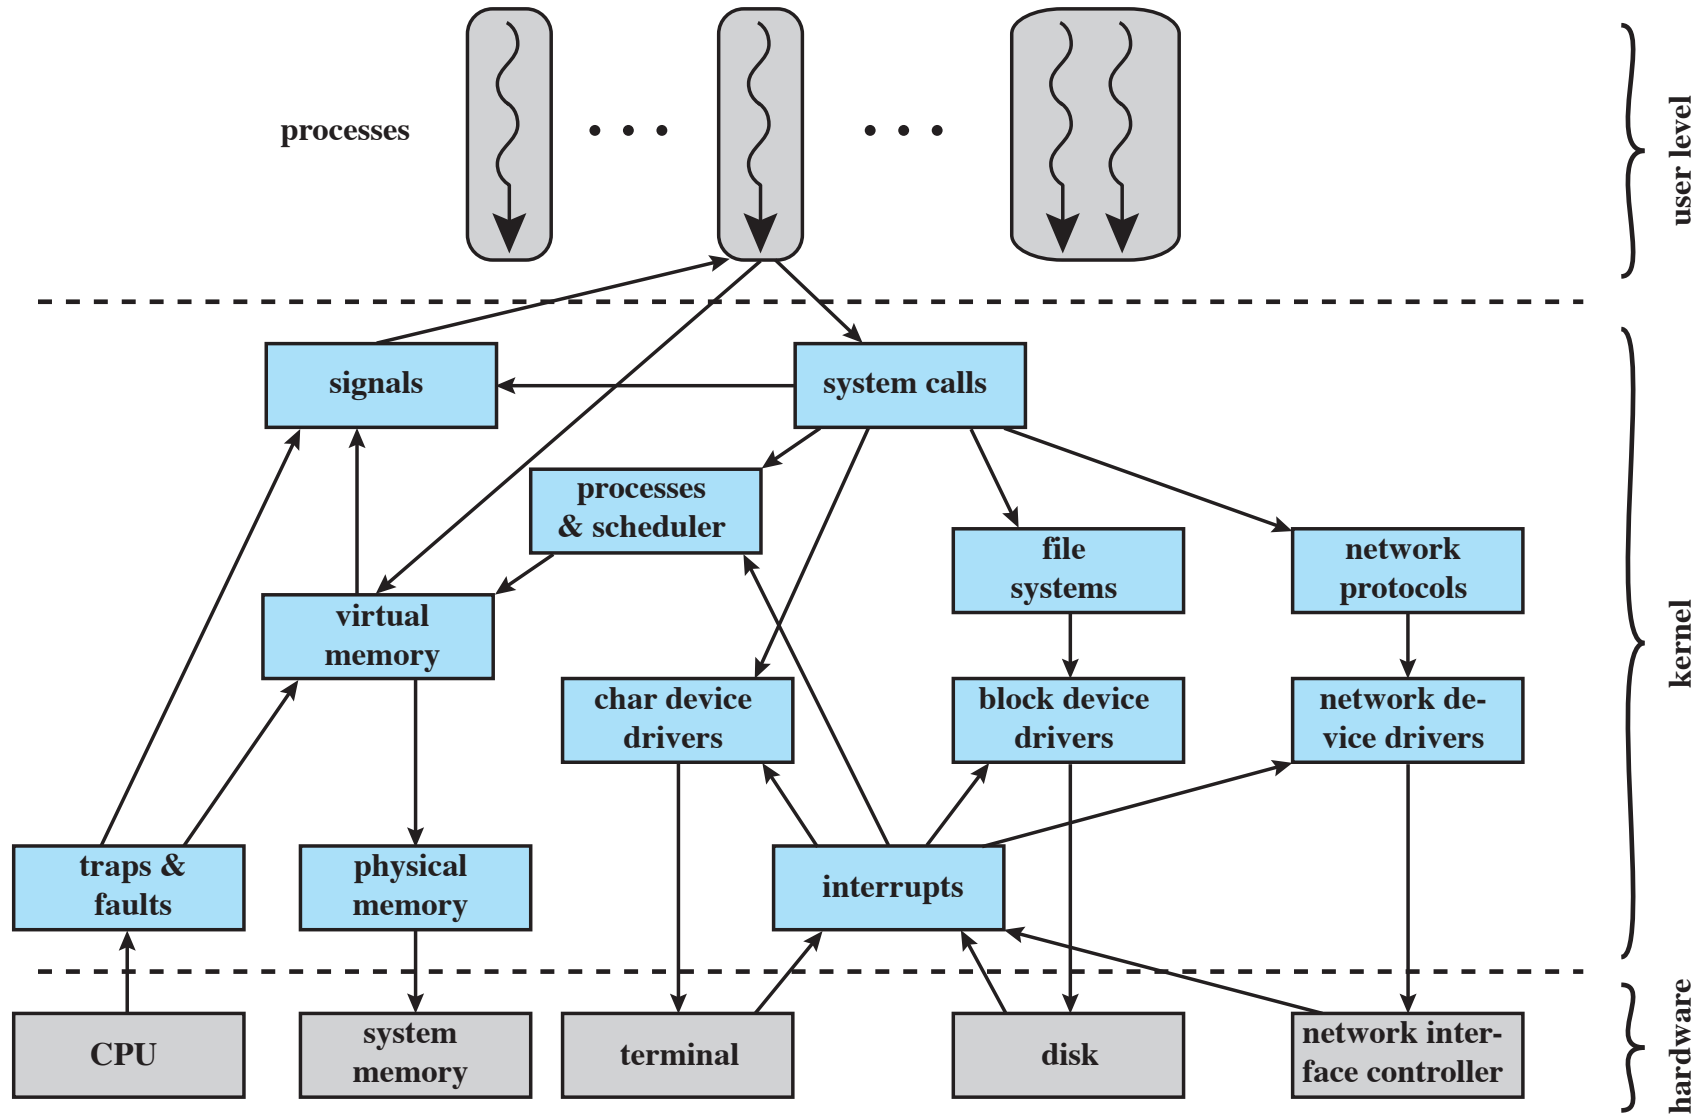
- Kicks off dispatching
  - Stop process, save state and start driver and return
- Dispatcher written in assembler

IO scheduler orders disk operation

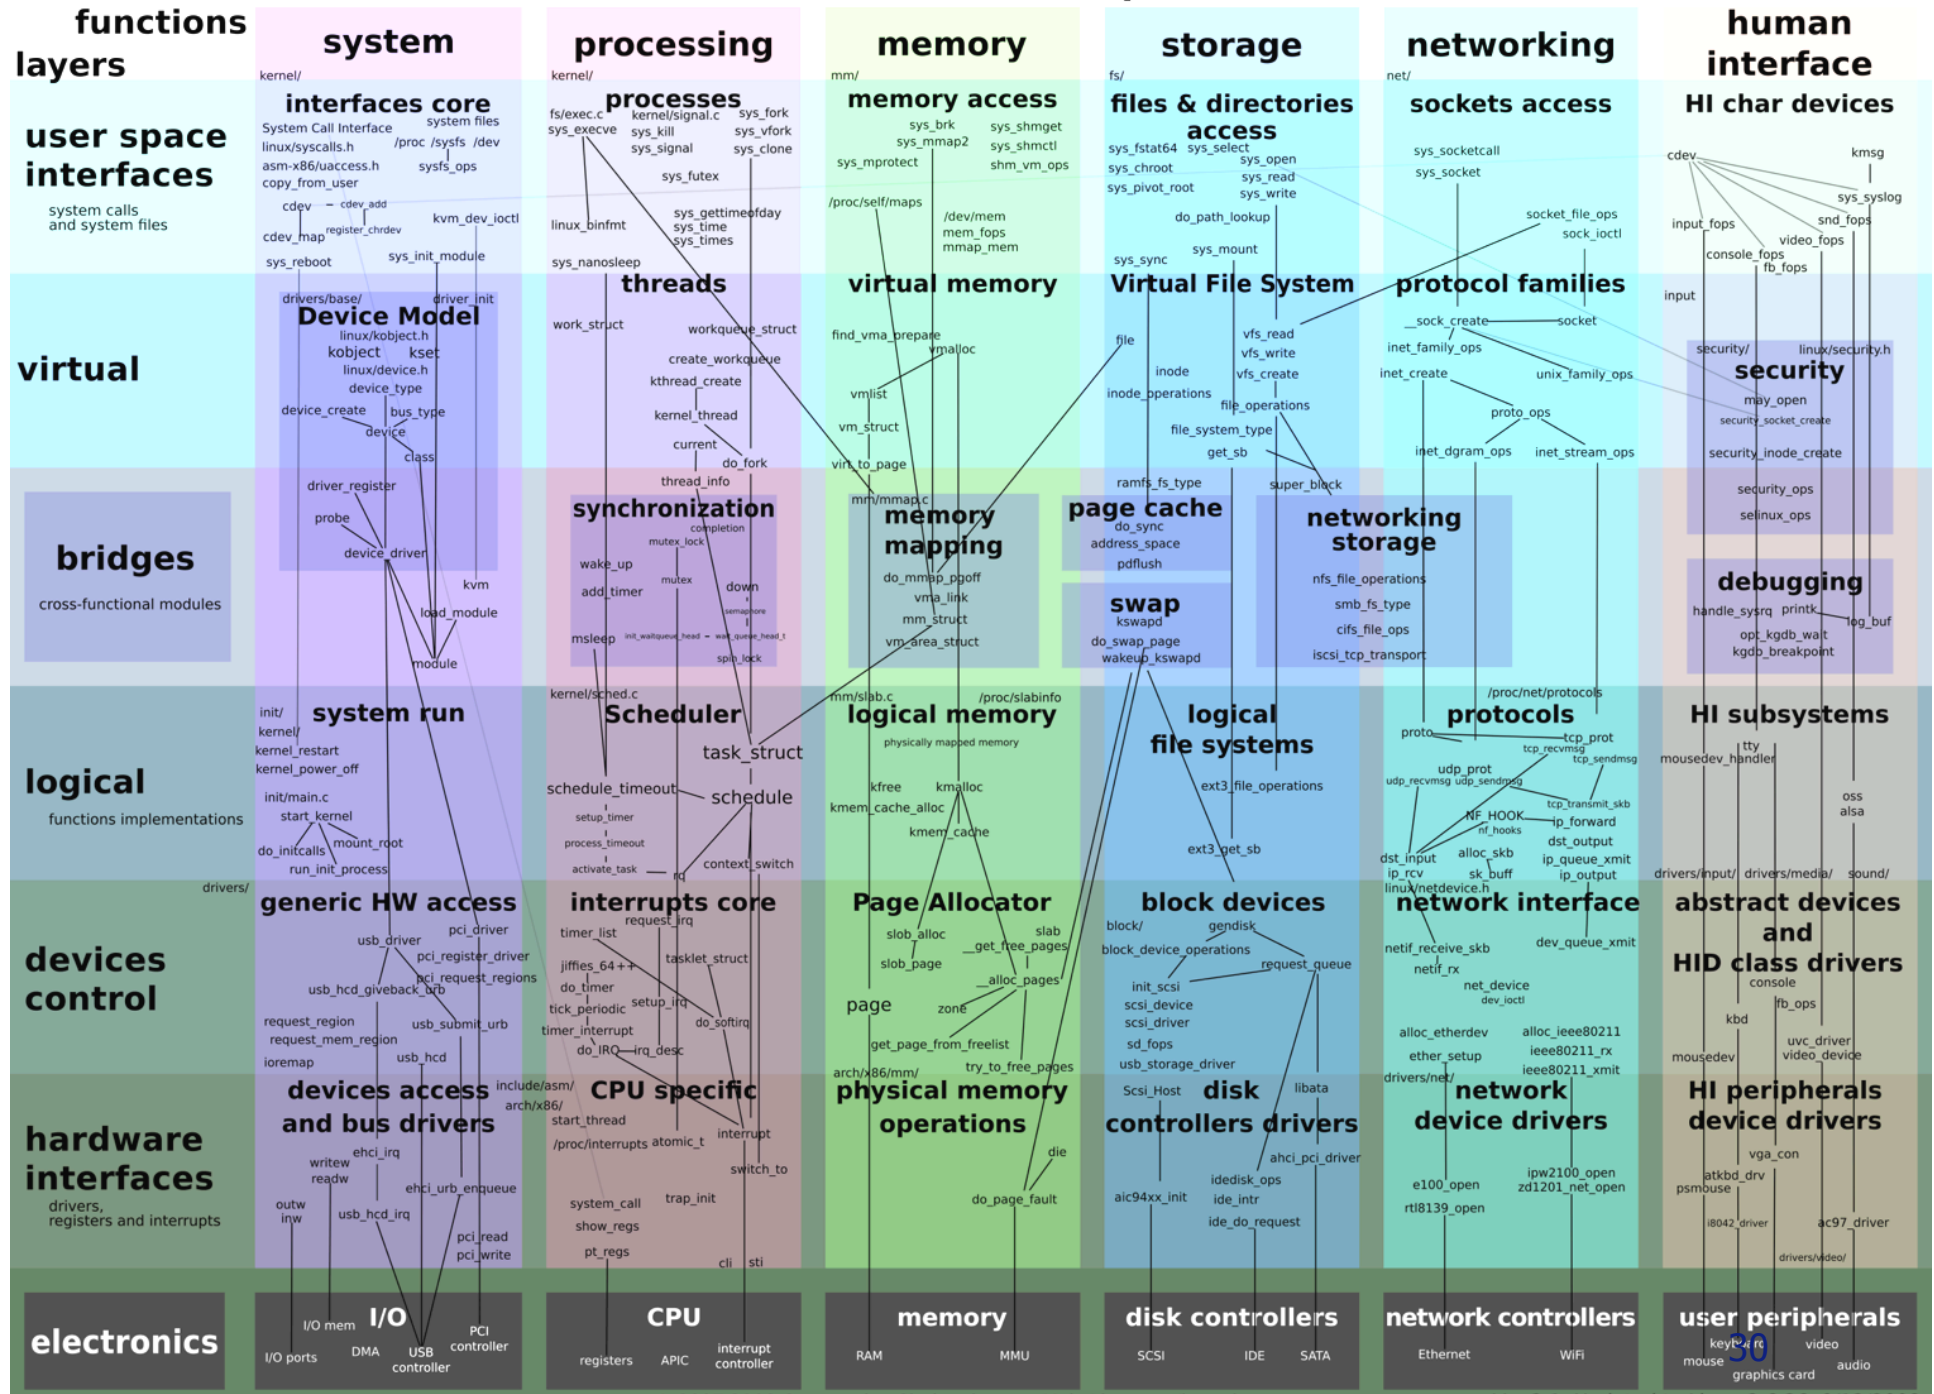
*Monolithic:*  
Static in-kernel  
components  
and dynamically  
loadable modules  
with shared internal  
data structures



# Linux Kernel Components



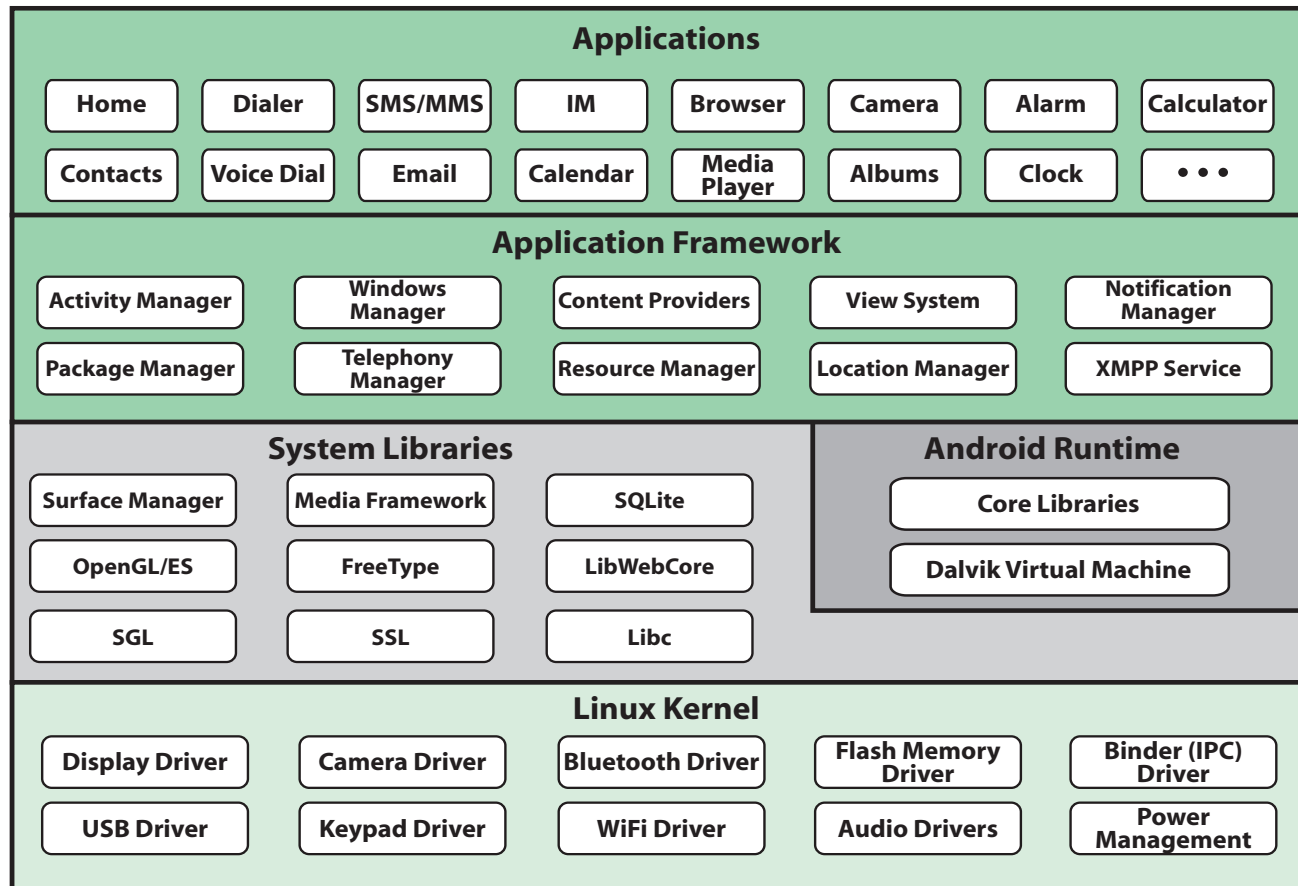
## Linux kernel map







# Android Operating System

- A Linux-based system originally designed for touchscreen mobile devices such as smartphones and tablet computers
- The most popular mobile OS
- Development was done by Android Inc., which was bought by Google in 2005
- 1<sup>st</sup> commercial version (Android 1.0) was released in 2008
- Most recent version is Android  $\geq$  4.3 (Jelly Bean)
- The Open Handset Alliance (OHA) was responsible for the Android OS releases as an open platform
- The open-source nature of Android has been the key to its success

# Android Software Architecture



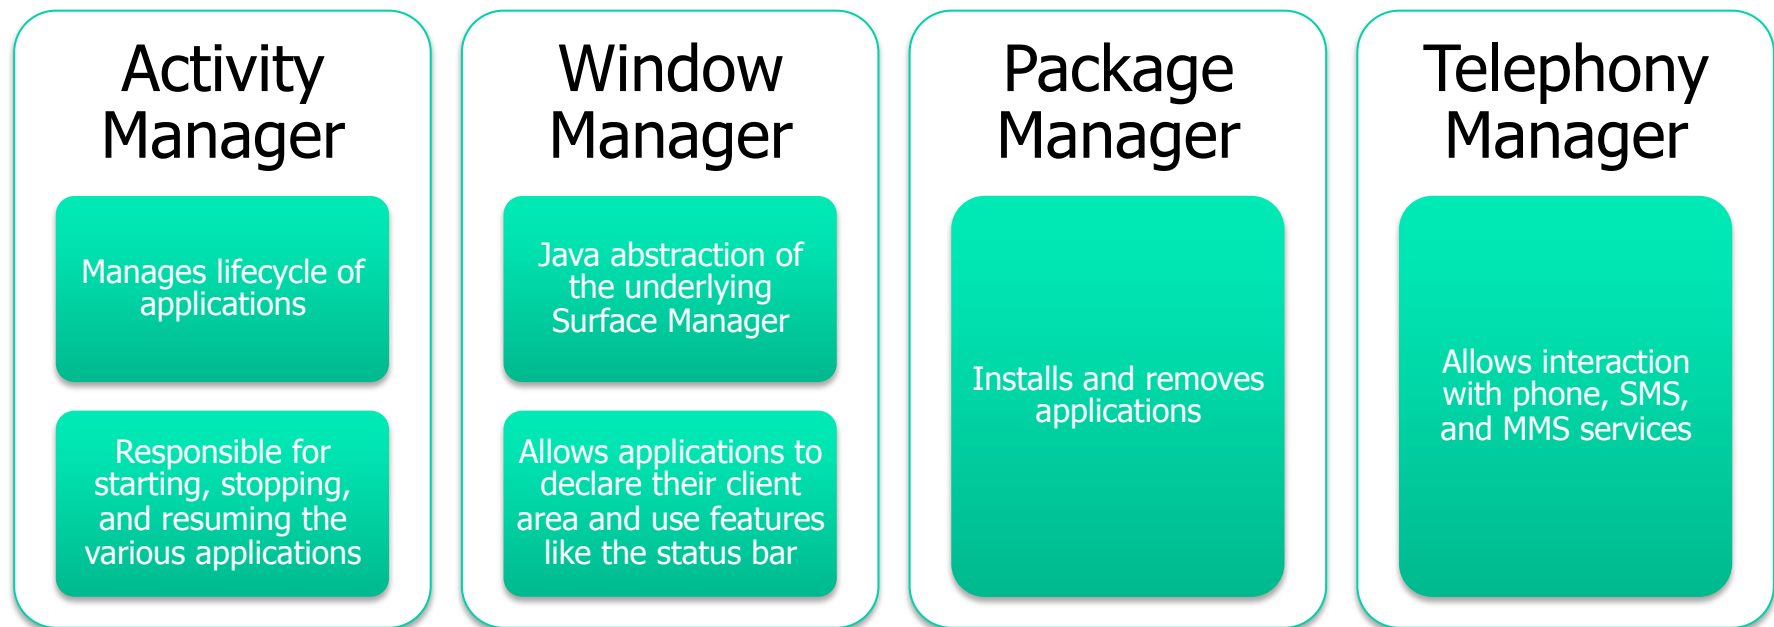
## Implementation:

-  Applications, Application Framework: Java
-   System Libraries, Android Runtime: C and C++
-  Linux Kernel: C



# Application Framework

- Provides high-level building blocks accessible through standardized API's that programmers use to create new apps
  - architecture is designed to simplify the reuse of components
- Key components:



# Application Framework (cont.)

Key components: (cont.)

- Content Providers
  - these functions encapsulate application data that need to be shared between applications such as contacts
- Resource Manager
  - manages application resources, such as localized strings and bitmaps
- View System
  - provides the user interface (UI) primitives as well as UI Events
- Location Manager
  - allows developers to tap into location-based services, whether by GPS, cell tower IDs, or local Wi-Fi databases
- Notification Manager
  - manages events, such as arriving messages and appointments
- XMPP
  - provides standardized messaging functions between applications

# System Libraries

Collection of useful system functions written in C or C++ and used by various components of the Android system

Called from the application framework and applications through a Java interface

Exposed to developers through the Android application framework

Some of the key system libraries include:

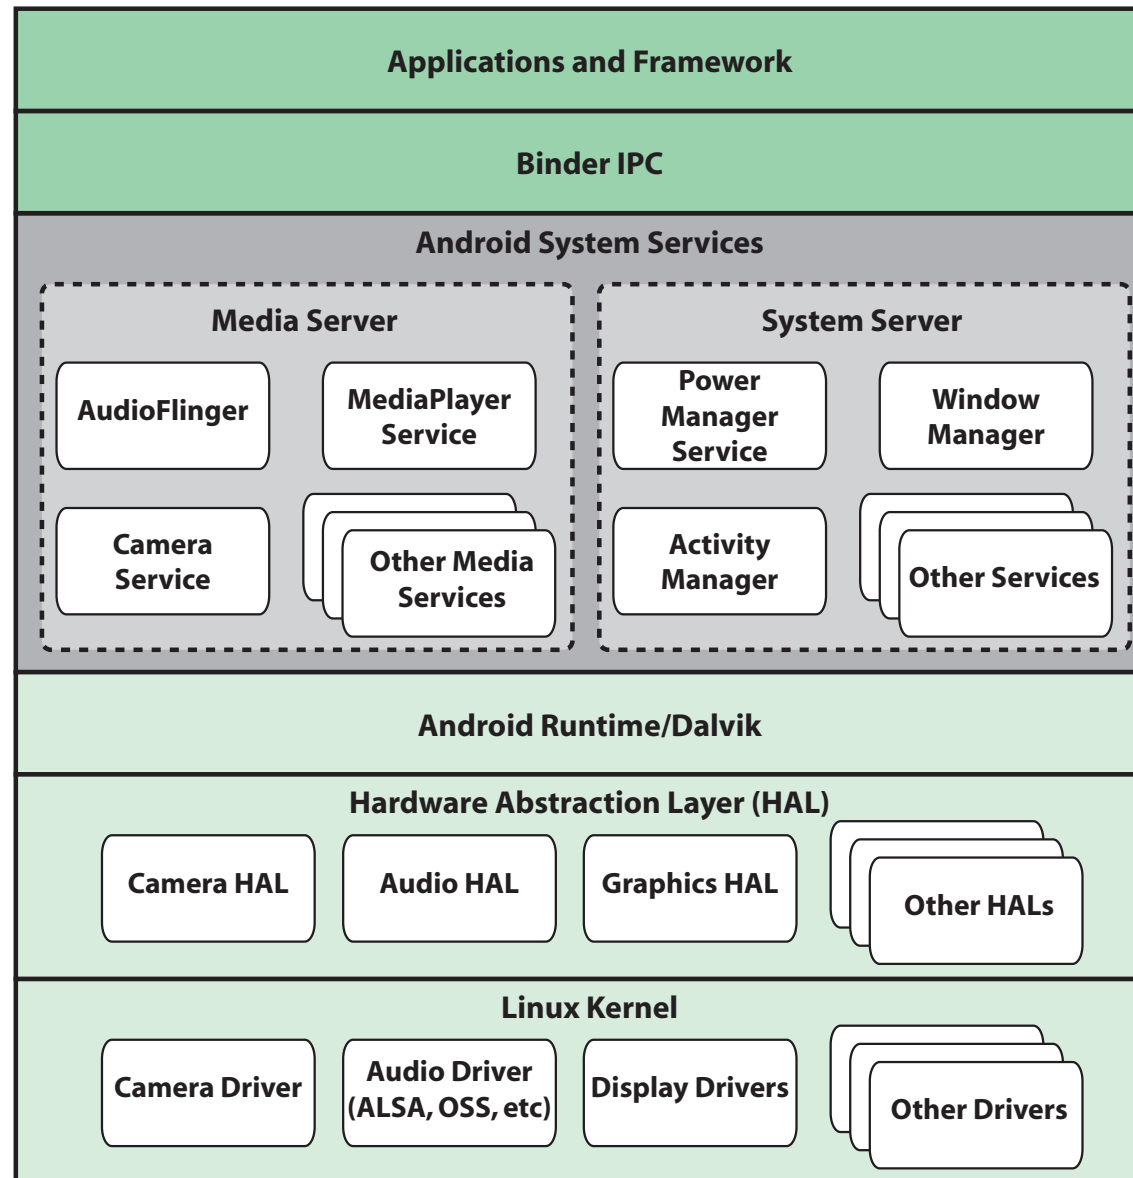
- Surface Manager
- OpenGL
- Media Framework
- SQL Database
- Browser Engine
- Bionic LibC

# Android Runtime



- Every Android application runs in its own process with its own instance of the Dalvik virtual machine (DVM)
- DVM executes files in the Dalvik Executable (.dex) format
- Component includes a set of core libraries that provides most of the functionality available in the core libraries of the Java programming language
- To execute an operation the DVM calls on the corresponding C/C++ library using the Java Native Interface (JNI)

# Android System Architecture



# Activities

- An activity is a single visual user interface component, including things such as menu selections, icons, and checkboxes
- Every screen in an application is an extension of the Activity class
- Use Views to form graphical user interfaces that display information and respond to user actions

# Power Management


## Alarms

- Implemented in the Linux kernel and is visible to the app developer through the AlarmManager in the RunTime core libraries
- Implemented in the kernel so that an alarm can trigger even if the system is in sleep mode
  - this allows the system to go into sleep mode, saving power, even though there is a process that requires a wake up

## Wakelocks

- Prevents an Android system from entering into sleep mode
- These locks are requested through the API whenever an application requires one of the managed peripherals to remain powered on
- An application can hold one of the following wakelocks:
  - Full\_Wake\_Lock
  - Partial\_Wake\_Lock
  - Screen\_Dim\_Wake\_Lock
  - Screen\_Bright\_Wake\_Lock

# Evolution of OS Code Sizes



source: Wikipedia 2010

## Code bloat

- Is lines of code useful comparison for complexity?
  - e.g. Linux scheduler (50K LoC); Vista scheduler (75K LoC)



# Summary

## OS Functions

- Simplify programming: device abstraction; virtual machine; memory management, file systems.
- Support concurrency, resource sharing & synchronisation

## Kernel Structure

- Monolithic, Micro & Hybrid.

## Operating System complexity