


CO580 Algorithms

Dr Timothy Kimber

January 2015



Course Outline

This NEW course will cover

- The role of algorithms in computing
- How to design algorithms
- How to evaluate algorithms
- (Lots of) algorithms used in specific settings

This course will be

- Practical – it should take your programming to the next level
- Somewhat mathematical

The main language used will be [Java](#), although a good algorithm can usually be implemented in many languages (even Prolog!)

Course Outline

The lecturer

- I am a Teaching Fellow in DoC
- I have a PhD in Computational Logic
- I teach Prolog to the MAC and Specialism classes and have also taught Java programming

The structure

- 27 hours of lectures and lab-based tutorials (weeks 2–10)
- One or two assessed exercises (10%)
- A 2-hour written examination *next term* (90%)


Books

- *Introduction to Algorithms*, Cormen et al., 3rd edn, 2009.
- *Algorithms*, Sedgewick & Wayne 4th edn, 2011.

What Is An Algorithm?

Algorithm a procedure for solving a **mathematical problem** in a finite number of steps that frequently involves repetition of an operation; *broadly* : a step-by-step procedure for solving a problem or accomplishing some end (especially by a computer)

<http://www.merriam-webster.com/dictionary/algorithm>



Muḥammad ibn Mūsā al-Khwārizmī (780–850)

Some Mathematical Problems

Example (*AgeCalc*)

Given: a person's current age *Age*; the current year *Year*; another year *AnotherYear*

Find: the person's age in *AnotherYear*

A Solution: $\text{Age} + (\text{AnotherYear} - \text{Year})$

Example (*Greatest Common Divisor*)

Given: two integers *X* and *Y*

Find: an integer *Z* such that: *Z* is a divisor of both *X* and *Y*; there is no integer $Z' > Z$ that is also a divisor of both *X* and *Y*

A Solution: Euclid's Algorithm (300 BC)

Algorithms, Programs, Correctness

- So, an algorithm is the underlying solution that is **implemented by** a program.
- We will see how algorithms can be analysed independently from any particular implementation and what makes a “good” algorithm.

Properties of Algorithms

- You have been focusing on writing programs that satisfy one key requirement: that they implement a **correct** algorithm.

Definition (Correct Algorithm)

A procedure is a *correct algorithm* for a problem iff, for every input instance of the problem, it halts with the correct output [Cormen p6]

Beyond Correctness

- This course goes beyond correctness to look at another vital property of algorithms: the **resources** they use
- Resources:
 - space** main memory used
 - time** number of CPU cycles used
- Commonly we are most interested in time (speed)
- Space and time are often **traded off** against one another. Using memory to store extra information can save a lot of time

Time

- So, **how long** will my algorithm take to 'run'?
- Consider this example:

Example (*List Search*)

Input: the sequence (list) $L = \langle a_1, \dots, a_N \rangle$ of integers, and an integer k

Output: *True* if k is in L , *False* otherwise

Simplification: assume L is ordered

Simple List Search

Simple Search (Input: list L and value k)

- For each element e in L
 - If $e == k$ Output *True* and HALT
- Output *False* and HALT

Questions to answer:

- Is Simple Search correct?
- How long will it take to run?

Simple Search

$$k = 21$$

L	5	6	7	21	23	29	50	
-----	---	---	---	----	----	----	----	--

Simple Search

$$k = 21$$

L	5	6	7	21	23	29	50	
-----	---	---	---	----	----	----	----	--

Simple Search

$k = 21$

L	5	6	7	21	23	29	50	
-----	---	---	---	----	----	----	----	--

Simple Search

$$k = 21$$

L	5	6	7	21	23	29	50	
-----	---	---	---	----	----	----	----	--

Simple Search

$$k = 21$$

L	5	6	7	21	23	29	50	
-----	---	---	---	----	----	----	----	--

Simple Search

$k = 21$

L	5	6	7	21	23	29	50	
-----	---	---	---	----	----	----	----	--

- Output: True
- A very simple example, but how we are going to [analyse](#) it is the important part

Java Implementation

```
private static boolean search(int[] a, int k) {  
    int N = a.length;  
    for (int i = 0; i < N; i++) {  
        if (a[i] == k) { return true; }  
    }  
    return false;  
}
```

- So, how long will this Java implementation take to run?

Java Implementation

```
private static boolean search(int[] a, int k) {  
    int N = a.length;  
    for (int i = 0; i < N; i++) {  
        if (a[i] == k) { return true; }  
    }  
    return false;  
}
```

- So, how long will this Java implementation take to run?
- Running time will depend on: compiler, processor, ..., size of input, and **type of input**
- We need to decide what input case to consider
- What would be **best, average or worst** case inputs?

Simple Search: Worst Case Analysis

- For worst case input Simple Search is a 'linear' algorithm
- We construct a formula to represent the running time for an input of 'size' N : $T(N) = \sum_{i=1}^n c_i t_i$
- The cost (time taken) for line i is represented by c_i . For a given processor etc. each c_i is **constant** and **small**.

	<i>// COST</i>	<i>TIMES</i>
<code>private static boolean search(int[] a, int k) {</code>		
<code>int N = a.length;</code>	<i>// c1</i>	<i>1</i>
<code>for (int i = 0; i < N; i++) {</code>	<i>// c2</i>	<i>N + 1</i>
<code>if (a[i] == k)</code>	<i>// c3</i>	<i>N</i>
<code>{ return true; }</code>	<i>// c4</i>	<i>0</i>
<code>}</code>		
<code>return false;</code>	<i>// c5</i>	<i>1</i>
<code>}</code>		

Runtime Analysis

- So, the running time for a **worst case input of size N** is

$$T(N) = c_1 + (N + 1)c_2 + Nc_3 + c_5$$

- or


$$T(N) = (c_1 + c_2 + c_5) + (c_2 + c_3)N$$

- or

$$T(N) = aN + b$$


- We can **empirically** determine a and b for a specific program, machine. This allows performance to be predicted, improved etc.
- However, the general formula applies for **any implementation**
- The important (but unsurprising) result is that $T(N)$ is a **linear function** of N . The running time is directly proportional to N .

The Algorithm Battlefield



- The running time of an algorithm is some function of N
- For small N all algorithms are
 - fast
 - roughly the same

The Algorithmm Battlefield




For large N :

- Algorithm performance resolves into classes according to the highest order term in N (order N , order N^2 etc.)
- $T(N) = aN^2 + bN \gg T'(N) = cN$ regardless of a, b, c (Why?)
- This provides clear goals for algorithm design

Asymptotic Notation


- “For large N ” means as N grows without bound or **asymptotically**.
- The asymptotic growth of a function (such as $T(N)$ for some algorithm) is classified using Θ , O and Ω notations.
- These notations define **bounds** on the growth of a function $f(N)$ with respect to some other function $g(N)$.
- $f(N)$ is
 - $O(g(N))$ if $g(N)$ is an asymptotic **upper** bound for $f(N)$
 - $\Omega(g(N))$ if $g(N)$ is an asymptotic **lower** bound for $f(N)$
 - $\Theta(g(N))$ if $g(N)$ is an asymptotically **tight** bound for $f(N)$
- The asymptotic bound of its running time (most often $O(g(N))$) is the most common way to characterise an algorithm

Big O: Upper Bound




$$O(g(N)) = \left\{ f(N) \mid \begin{array}{l} \text{there are positive constants } c \text{ and } n_0 \\ \text{such that } 0 \leq f(N) \leq c g(N) \text{ for all } N \geq n_0 \end{array} \right\}$$

Big Omega: Lower Bound



$$\Omega(g(N)) = \left\{ f(N) \mid \begin{array}{l} \text{there are positive constants } c \text{ and } n_0 \\ \text{such that } 0 \leq c g(N) \leq f(N) \text{ for all } N \geq n_0 \end{array} \right\}$$

Big Theta: Tight Bound



$$\Theta(g(N)) = \left\{ f(N) \mid \begin{array}{l} \text{there are positive constants } c_1, c_2 \text{ and } n_0 \\ \text{such that} \\ 0 \leq c_1 g(N) \leq f(N) \leq c_2 g(N) \text{ for all } N \geq n_0 \end{array} \right\}$$

Asymptotic Runtime Analysis

For Simple Search

- The worst case performance is given by $T(N) = aN + b$
- This function is $\Omega(N)$, $O(N)$ and $\Theta(N)$, written (abusively)
 - $T(N) = \Omega(N)$
 - $T(N) = O(N)$
 - $T(N) = \Theta(N)$
- Therefore, the running time for **any input** is also $O(N)$
- Is the running time for any input also $\Omega(N)$ and $\Theta(N)$?

In General

- $f(N) = \Theta(N) \iff f(N) = O(N) \text{ and } f(N) = \Omega(N)$
- $f(N) = O(N^x) \implies f(N) = O(N^y) \text{ for all } y > x$

Highest Order Terms

If running time is independent of N , we say the algorithm **runs in constant time** and write

- $T(N) = \Theta(1)$

As we would expect, in a **polynomial** the term with the largest exponent dominates.

Definition (Polynomial)

A **polynomial of degree d** (for $d \geq 1$) is a function $p(N)$ of the form

$$p(N) = a_0 + a_1N + a_2N^2 + \cdots + a_dN^d$$

in which $a_d \neq 0$. The polynomial is **asymptotically positive** iff $a_d > 0$.

If $T(N)$ is an asymptotically positive polynomial of degree d , then

- $T(N) = \Theta(N^d)$

Highest Order Terms

Exponential functions include a term of the form a^N .

- If $a > 1$ then the function grows faster than **any polynomial**

Polylogarithmic functions include a term of the form $(\log_2 N)^k$.

- Recall: $b^{\log_b a} = a$
- k is a constant
- Changing base just multiplies by a constant
- Any positive polynomial grows faster than **any polylogarithm**

Binary List Search

- So, we have a $O(N)$ search algorithm. Can we do any better?
- Do we have to look at every element? Since list is ordered we can eliminate whole chunks at a time.

Binary Search (Input: list L , value k)

- Repeat
 - Choose an element e near the middle of the list
 - If $e == k$, Output *True* and HALT
 - Otherwise, if $k > e$, continue searching elements after e
 - Otherwise continue searching elements before e
 - Until there is nothing left to search
 - Output *False* and HALT
-
- Is Binary Search correct?

Binary List Search

$$k = 1$$

L	5	6	7	21	23	29	50	
-----	---	---	---	----	----	----	----	--

Binary List Search

$$k = 1$$

L	5	6	7	21	23	29	50	
-----	---	---	---	----	----	----	----	--

Binary List Search

$k = 1$

L	5	6	7	21	23	29	50	
-----	---	---	---	----	----	----	----	--

Binary List Search

$k = 1$

L	5	6	7	21	23	29	50	
-----	---	---	---	----	----	----	----	--

Binary List Search

$$k = 1$$

L	5	6	7	21	23	29	50	
-----	---	---	---	----	----	----	----	--

Binary List Search

$k = 1$

L	5	6	7	21	23	29	50	
-----	---	---	---	----	----	----	----	--

Binary List Search

$$k = 1$$

L	5	6	7	21	23	29	50	
-----	---	---	---	----	----	----	----	--

Binary List Search

$k = 1$

L	5	6	7	21	23	29	50	
-----	---	---	---	----	----	----	----	--

- Output: False

Java Implementation

```
private static boolean binarySearch(int k, int[] a) {  
    int l = 0, m, n = a.length;           // array indexes  
  
    while (n > l) {  
        m = l + (n - l) / 2;  
  
        if (k == a[m]) { return true; }  
        else if (k > a[m]) { l = m + 1; }  
        else { n = m; }  
    }  
  
    return false;  
}
```


- What is the worst case performance?

Binary List Search: Java


 $k = 1$

	0	1	2	3	4	5	6	
a	5	6	7	21	23	29	50	


Binary List Search: Java




Binary List Search: Java




Binary List Search: Java




Binary List Search: Java




Binary List Search: Java




Binary List Search: Java




Binary List Search: Java




Binary List Search: Java




Binary List Search: Java



Binary List Search: Java



Binary List Search: Java



- false

Java Implementation

```
private static boolean binarySearch(int k, int[] a) {  
    int l = 0, m, n = a.length;           // array indexes  
  
    while (n > l) {  
        m = l + (n - l) / 2;  
  
        if (k == a[m]) { return true; }  
        else if (k > a[m]) { l = m + 1; }  
        else { n = m; }  
    }  
  
    return false;  
}
```

- What is the worst case performance?

Binary Search: Worst Case Analysis

	<i>// COST</i>	<i>TIMES</i>
<code>private static boolean binarySearch(int k, int[] a) {</code>		
<code>int l = 0, m, n = a.length;</code>	<i>// c1</i>	1
<code>while (n > l) {</code>	<i>// c2</i>	$t(N) + 1$
<code>m = l + (n - l) / 2;</code>	<i>// c3</i>	$t(N)$
<code>if (k == a[m])</code>	<i>// c4</i>	$t(N)$
<code>{ return true; }</code>	<i>// c5</i>	0
<code>else if (k > a[m])</code>	<i>// c6</i>	$t(N)$
<code>{ l = m + 1; }</code>	<i>// c7</i>	0
<code>else</code>		
<code>{ n = m; }</code>	<i>// c8</i>	$t(N)$
<code>}</code>		
<code>return false;</code>	<i>// c9</i>	1
<code>}</code>		

- For now just let number of loop iterations required to search a list of size N be $t(N)$

Recurrence Equations


- $t(N)$ can be expressed in the form a **recurrence equation**:

$$t(N) = \begin{cases} 1 & , \text{ if } N = 1 \\ 1 + t(N') & , \text{ if } N > 1 \end{cases}$$

- N' represents the number of elements left to search if we start with N and execute the loop once
- Not quite a recurrence yet - we need to replace N' with an expression involving N
- What is N' (in the worst case)?

Worst Case for N'


$k = 1$



- Observe: m is always placed at $1 + \lfloor N/2 \rfloor$
- if $k < a[m]$ we have $\lfloor N/2 \rfloor$ elements left
- if $k > a[m]$ we have
 - $\lfloor N/2 \rfloor$ elements left if N is odd
 - $\lfloor N/2 \rfloor - 1$ elements left if N is even
- So the most elements we have left to search is $N' = \lfloor N/2 \rfloor$

Worst Case for N'


$k = 1$



- Observe: m is always placed at $1 + \lfloor N/2 \rfloor$
- if $k < a[m]$ we have $\lfloor N/2 \rfloor$ elements left
- if $k > a[m]$ we have
 - $\lfloor N/2 \rfloor$ elements left if N is odd
 - $\lfloor N/2 \rfloor - 1$ elements left if N is even
- So the most elements we have left to search is $N' = \lfloor N/2 \rfloor$

Worst Case for N'


$k = 100$



- Observe: m is always placed at $1 + \lfloor N/2 \rfloor$
- if $k < a[m]$ we have $\lfloor N/2 \rfloor$ elements left
- if $k > a[m]$ we have
 - $\lfloor N/2 \rfloor$ elements left if N is odd
 - $\lfloor N/2 \rfloor - 1$ elements left if N is even
- So the most elements we have left to search is $N' = \lfloor N/2 \rfloor$

Worst Case for N'

$k = 100$



- Observe: m is always placed at $1 + \lfloor N/2 \rfloor$
- if $k < a[m]$ we have $\lfloor N/2 \rfloor$ elements left
- if $k > a[m]$ we have
 - $\lfloor N/2 \rfloor$ elements left if N is odd
 - $\lfloor N/2 \rfloor - 1$ elements left if N is even
- So the most elements we have left to search is $N' = \lfloor N/2 \rfloor$

Solving the Recurrence

- So, the actual recurrence is

$$t(N) = \begin{cases} 1 & , \text{ if } N = 1 \\ 1 + t(\lfloor N/2 \rfloor) & , \text{ if } N > 1 \end{cases}$$

- Now we need to [solve the recurrence](#), so that we can determine the actual value of $t(N)$ when $N > 1$
- Simplification: assume N is some power of 2
- Then $\lfloor N/2 \rfloor = N/2$

Solving the Recurrence

- One technique is to **telescope** the recurrence
- Substituting terms on the right hand side we have

$$t(N) = 1 + t(N/2) \quad (1)$$

$$t(N) = 1 + 1 + t(N/4) \quad (2)$$

$$t(N) = 1 + 1 + 1 + t(N/8) \quad (3)$$

- We need to know how many 1s we have when the expression reaches

$$t(N) = 1 + \cdots + t(1)$$

- From (1–3) above we can see the answer is j where $2^j = N$
- So, $j = \log_2(N)$ and $t(N) = 1 + \log_2(N)$, for all N
- If we drop the assumption $t(N)$ will still contain a $\log_2 N$ term

Binary Search: Worst Case Analysis

```

private static boolean binarySearch(int k, int[] a) { // COST    TIMES
    int l = 0, m, n = a.length;                      // c1      1
    while (n > l) {                                    // c2      t(N) + 1
        m = l + (n - l) / 2;                          // c3      t(N)
        if (k == a[m])                                // c4      t(N)
            { return true; }                          // c5      0
        else if (k > a[m])                            // c6      t(N)
            { l = m + 1; }                             // c7      0
        else                                           // c8      t(N)
            { n = m; }
    }
    return false;                                     // c9      1
}

```

$$T(N) = (c_1 + c_2 + c_9) + (c_2 + c_3 + c_4 + c_6 + c_8)(1 + \log_2(N))$$

$$T(N) = a \log_2(N) + b$$

- So, $T(N) = \Theta(\log_2(N))$

Comparing the Algorithms

- Suppose we have implementations of the two algorithms:
 - Simple Search (excellent programmer): uses $5N$ instructions
 - Binary Search (average programmer): uses $100\log_2(N)$ instructions
- Simple Search uses 5000 instructions to search 1000 numbers.
- Binary Search can search $2^{5000/100} = 1.1 \times 10^{15}$ numbers in the same time! (if we had enough memory)
- In 1971 the first microprocessor ran at 740 kHz. If that machine ran Binary Search it could search a list of 10 million numbers faster than our lab machines (3.4 GHz) using Simple Search.
- Of course, both computers would still take much less than a second.
- However, 'naive' algorithms for some problems are often $O(N^2)$ or worse. In this case running times can easily be hours or days, even on today's hardware.

Divide and Conquer

- The Binary Search algorithm uses a **divide and conquer** approach.
- Divide and conquer reduces the main problem into a number of smaller **subproblems** (one in this case).
- The solutions to the subproblems may or may not need to be combined to produce the final solution.
- Divide and conquer often produces logarithmic performance.
- Divide and conquer is a technique that can be widely employed when designing algorithms. We shall see other such general **algorithmic schemes**.