

# **Computer Networks and Distributed Systems**

## **Interaction Implementation**

Course 527 – Spring Term 2014-2015

**Anandha Gopalan**


[a.gopalan@imperial.ac.uk](mailto:a.gopalan@imperial.ac.uk)

<http://www.doc.ic.ac.uk/~axgopala>

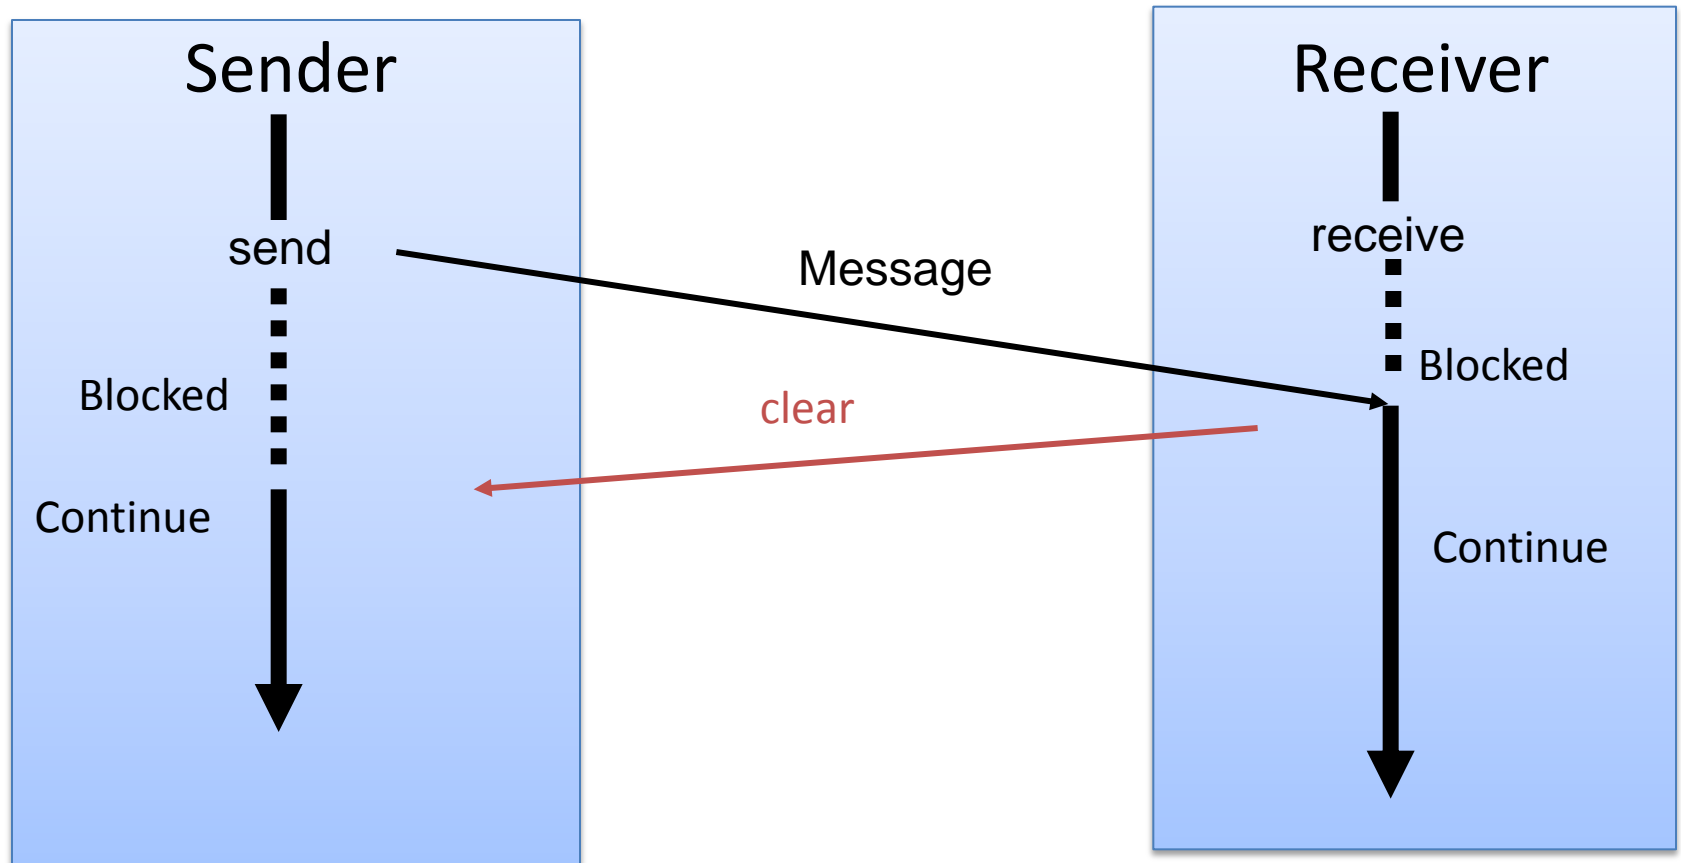
# Outline

- Message passing
- RPC implementation
  - Binding
  - Concurrency
  - Error Control
- Heterogeneity
  - External Representations
  - Transformations

# Implementing Asynchronous Send



# Implementing Synchronous Send

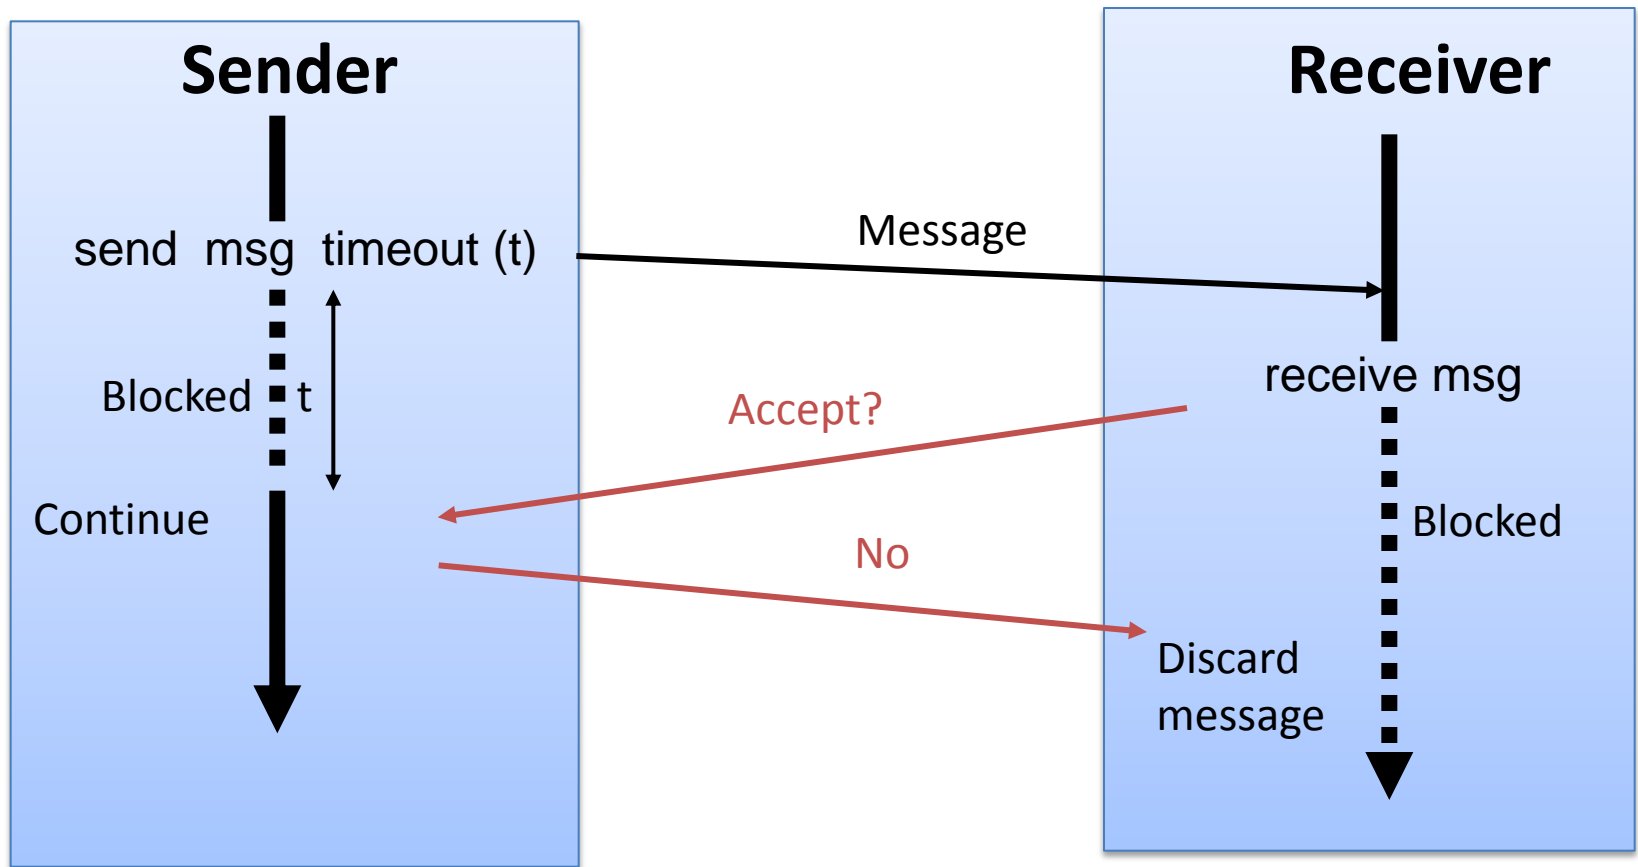


Clear is a runtime system message – not sent by application process

# Exercise

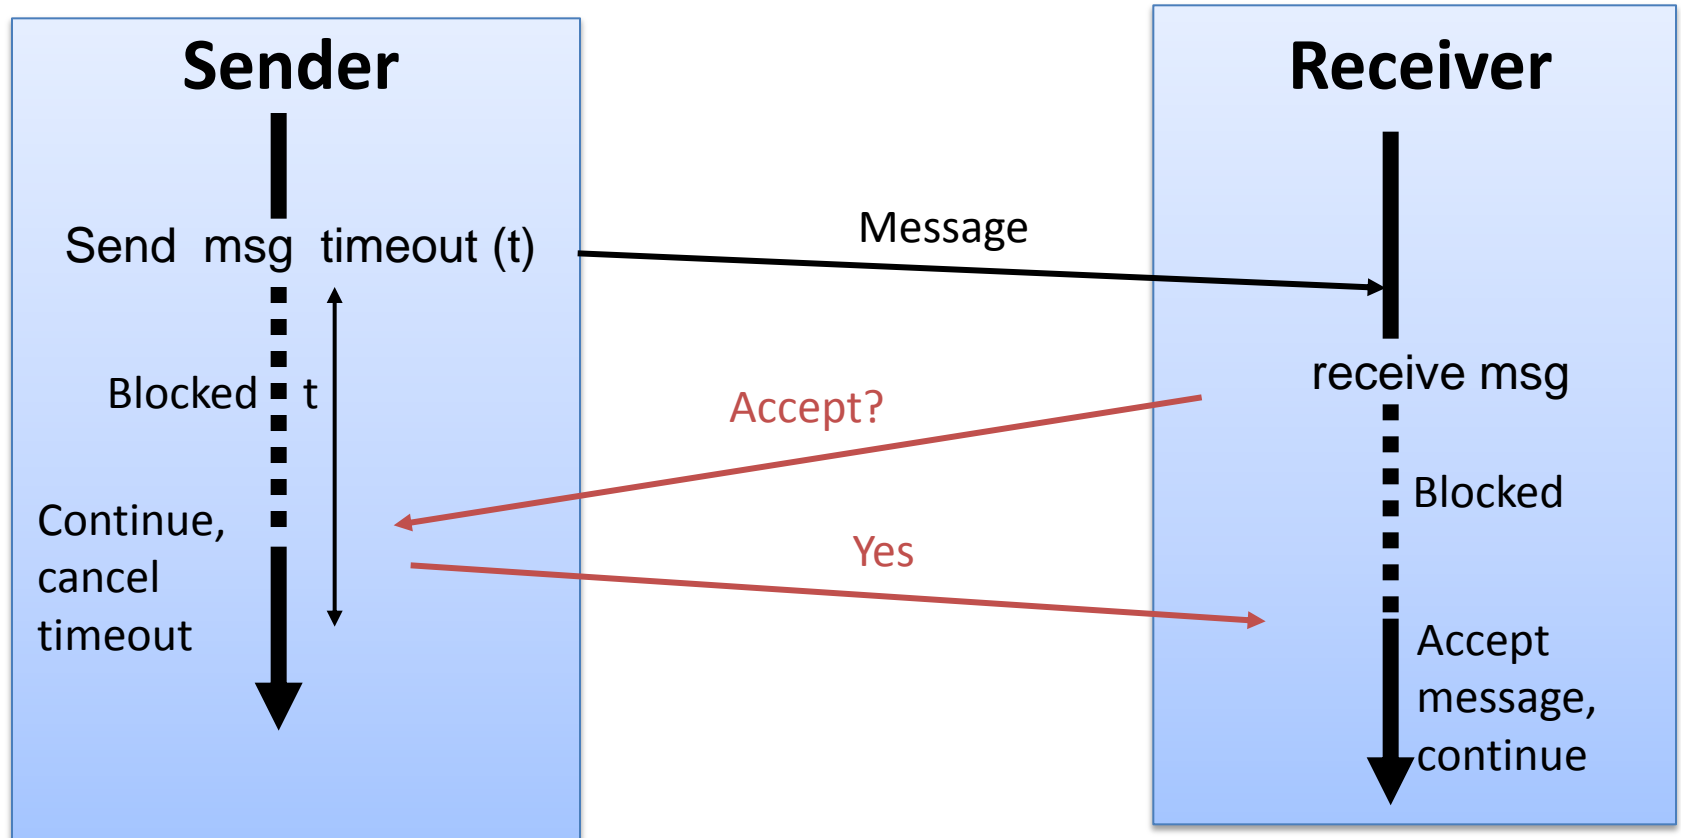
- Modify the synchronous protocol to cater for a timeout on the send i.e. `send msg delay (t)`
- The sender continues after the timeout if the message has not yet been received – this implies the receiver should not get the message if the timeout expires
- Show the message exchanges that would occur:
  - i. if the sender's timeout expires
  - ii. if the sender's timeout does not expire

# Synchronous Send: timeout expires



**Accept?** & **No** are sent by runtime system, not application processes

# Synchronous Send: timeout does not expire



**Accept?** & **Yes** are sent by runtime system, not application processes

# Binding

- Binding is the assignment of a reference value (e.g. address or object reference) to a placeholder (e.g. message port or object reference variable)
- It is similar to opening a connection in the communication system or opening a file in an OS
- **First Party Binding**
  - Client initiates binding as in Java and CORBA



# Interface Type Checking

- Client interface must be type compatible with server interface i.e. same interactions and signatures (set of parameters + data types)
- Client and server likely to be compiled independently and at different times
  - Use same interface type definition to generate client and server interface
  - Permit server to be subtype of client interface
  - Check for structural compatibility at run-time

# Interface Type Checking

- Use same interface type definition to generate client and server interface
  - Client and server hold identity of interface derived from interface definition module
  - Generate Interface identity by
    - checksum over source
    - name + timestamp of last modification or compilation
  - At bind time, check type identities are equal
  - Strong type compatibility

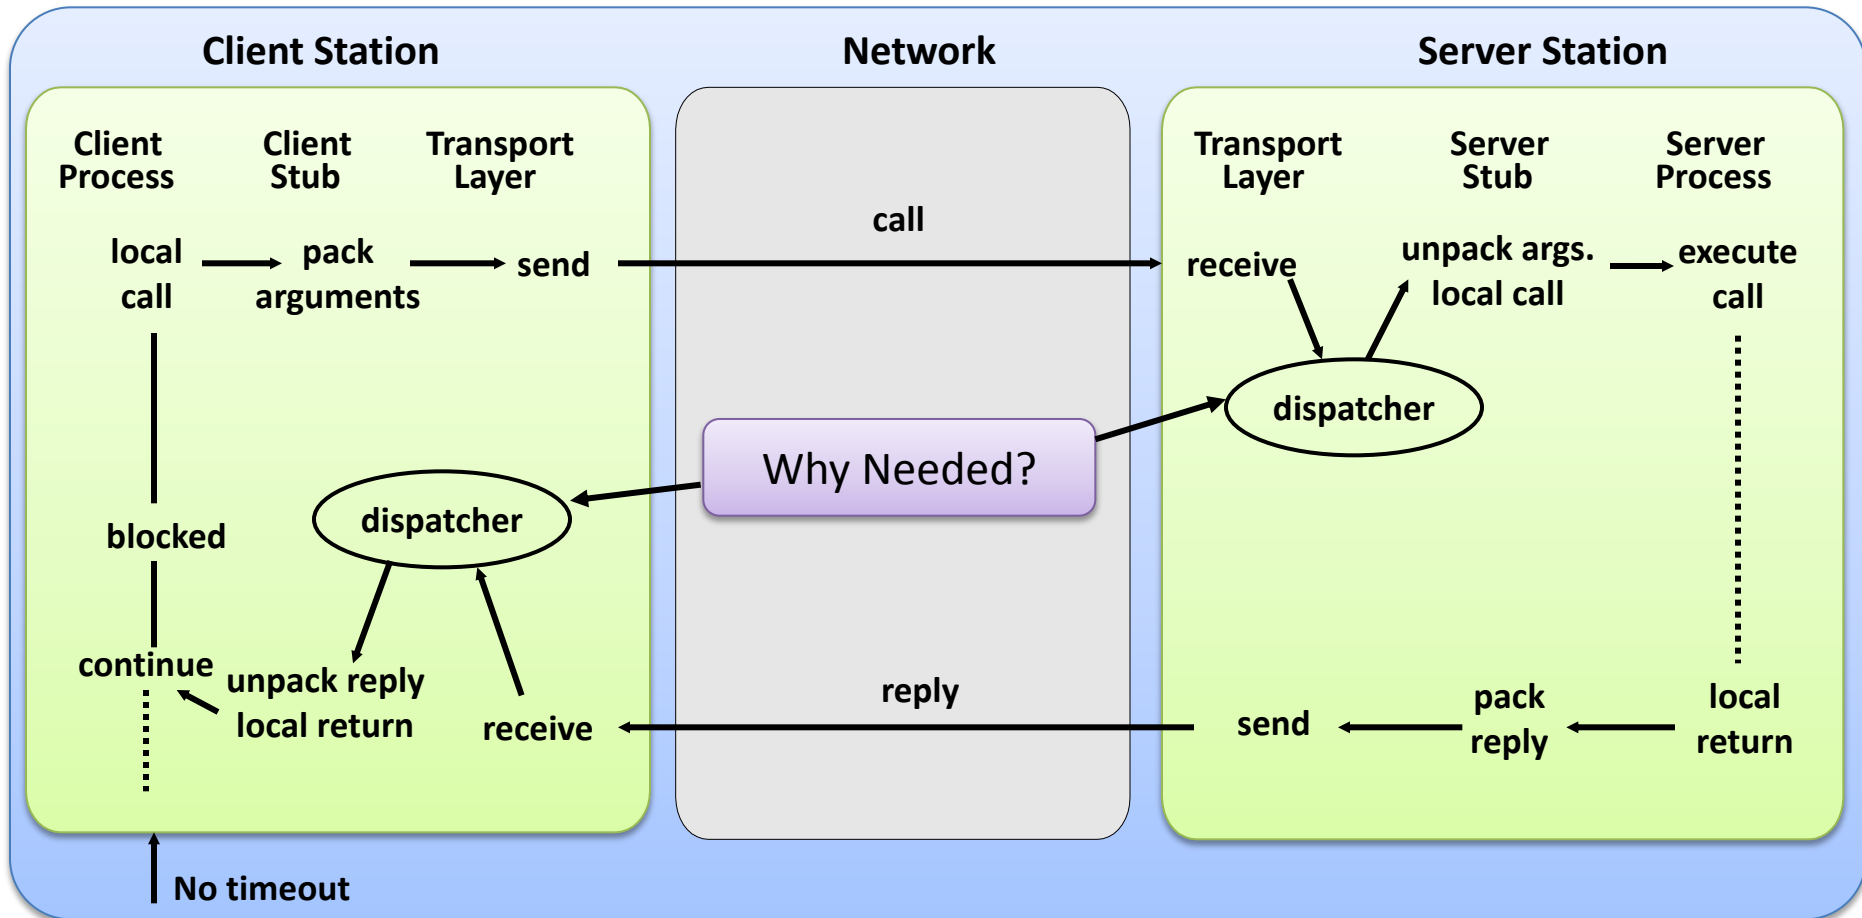
# Interface Type Checking

- Permit server to be subtype of client interface  
i.e. provides *additional* operations which are not used by client, but must not extend operations in original interface
- Maintain run-time representation of interface and check for structural compatibility at bind time
  - *Weak* type compatibility e.g. the following two interfaces are structurally equivalent

```
interface A {  
    opa1 (in string a1,  
        in short a2 , out long a4);  
    opa2 (in string a4);  
}
```

```
interface B {  
    opb1 (in string b1,  
        in short b2 , out long b3);  
    opb2 (in string b4)  
}
```

# Remote Procedure Call



At most once semantics  
client receives reply - procedure executed exactly once  
on failure i.e. no reply received – don't know

# Dispatcher

- Server needs dispatcher to map incoming calls onto relevant procedure
- Dispatcher in client passes incoming reply message to relevant stub procedure
- Interface compiler generates a number (or name) for each procedure in interface – inserted into call message by client stub procedure
- Dispatcher at server receives all call messages and uses procedure number (name) to identify called procedure

# RMI Dispatcher

- Java uses reflection and a generic dispatcher so no need for skeletons
- Client proxy (stub) includes information about a method in request message by creating instances of Method class containing
  - Class, types of arguments, type of return value, type of exceptions
  - Proxy marshalls object of class method, array of argument objects

# RMI Dispatcher

- Dispatcher receives request,
  - unmarshalls method object,
  - uses method information to unmarshall arguments
  - converts remote object reference to local object reference
  - calls method object's invoke method supplying local object reference and arguments
  - when method executed, marshalls result or exceptions into reply message and sends it back to client
- See <http://docs.oracle.com/javase/1.5.0/docs/api/java/lang/reflect/Method.html>

# RPC Binding

- A name server registers exported interfaces and is queried to locate a server when an interface is imported
- Server
  - Calls `export(interface type, server name, nameserver)`
  - Dispatcher address added by stub and passed to Transport
  - Server Transport
    - Generates unique exportid & sends a register message to name server containing type, name, exportid



# RPC Binding

- Client

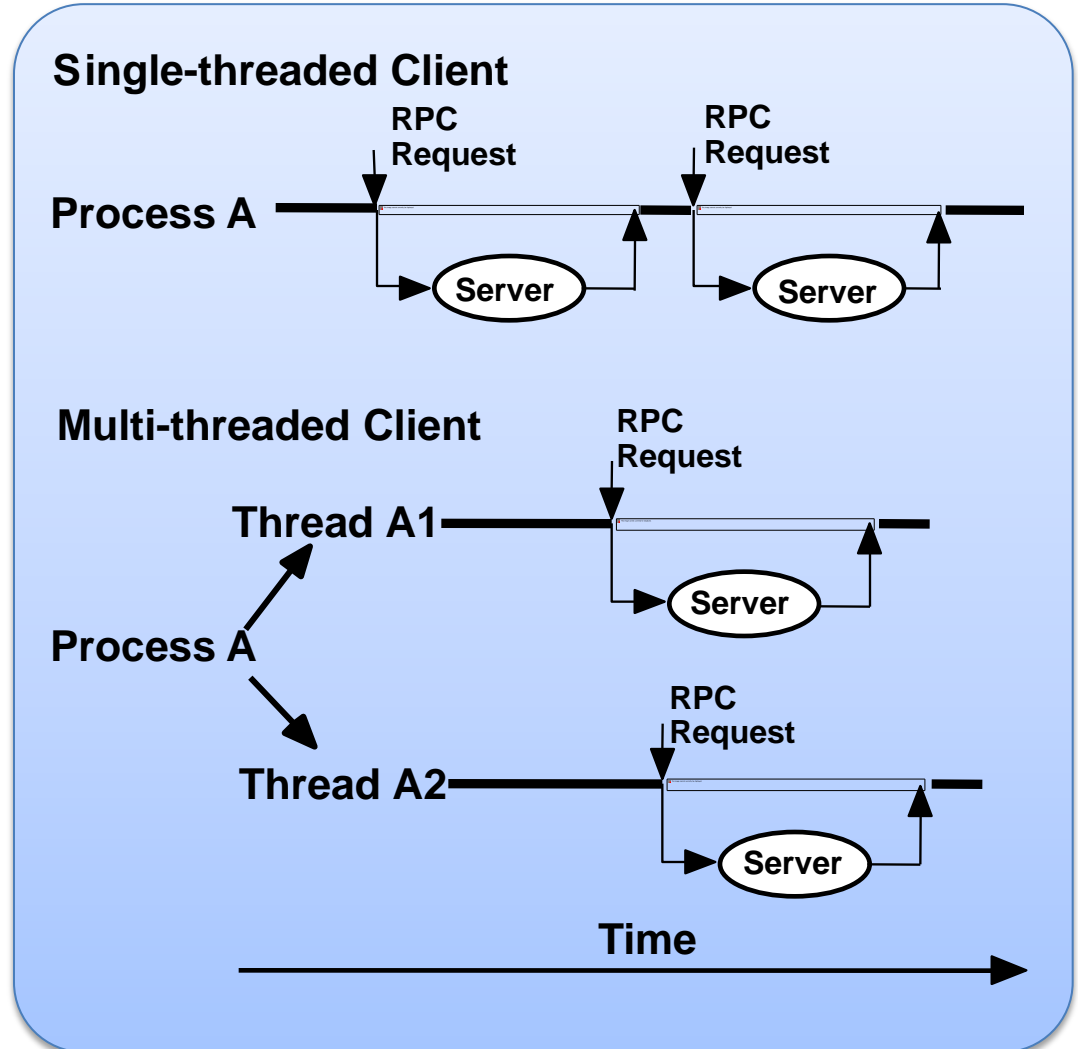
- Calls `import(interface type, server name, nameserver)`
- Dispatcher address added by stub and passed to Transport
- Client Transport
  - Send query message with type & name to nameserver; Reply contains type and address of server instance
  - Query server to check validity of type, name and exportid; Return interface reference (address) or error

# Failures

- Server Failure
  - Use exportid to detect failed server
  - On restart – exports interface again - generates a new exportid
  - All messages to server include exportid
  - Dispatcher aborts calls with incorrect exportid
- Client Failure
  - Orphans – client fails after making call but before receiving response
  - No ack to response
  - Server either implements a form of 'rollback' or does nothing


# Client Threads

- In a single-threaded program which does RPCs to different servers, the RPCs must be done serially
- Each RPC blocks the program for at least  $2 \times$  the network delay. Throughput is adversely affected
- Using threads, remote invocations (RPC or object invocation) may be performed concurrently by a single client process




# Server Concurrency

- Multi-threading can improve server responsiveness since if requests are processed concurrently, long requests will not block short requests



# Client Concurrency


No dead-locks with callbacks if client multi-threaded



# Server Implementation Options

- **Server is single active process**
  - Dispatcher processes one request at a time and calls the relevant stub procedure which calls the actual procedure. Problems?
- **Thread-per-Request**
  - Dispatcher creates a new thread to handle each request. Problems?
- **Thread Pool**
  - A fixed number of threads are generated at start-up and free threads are allocated to requests by the dispatcher
  - Concurrency but lower creation overheads
- **Thread-per-Session**
  - A thread is created at connection set up to process all requests from the particular client. Problems?

# RPC Error Control



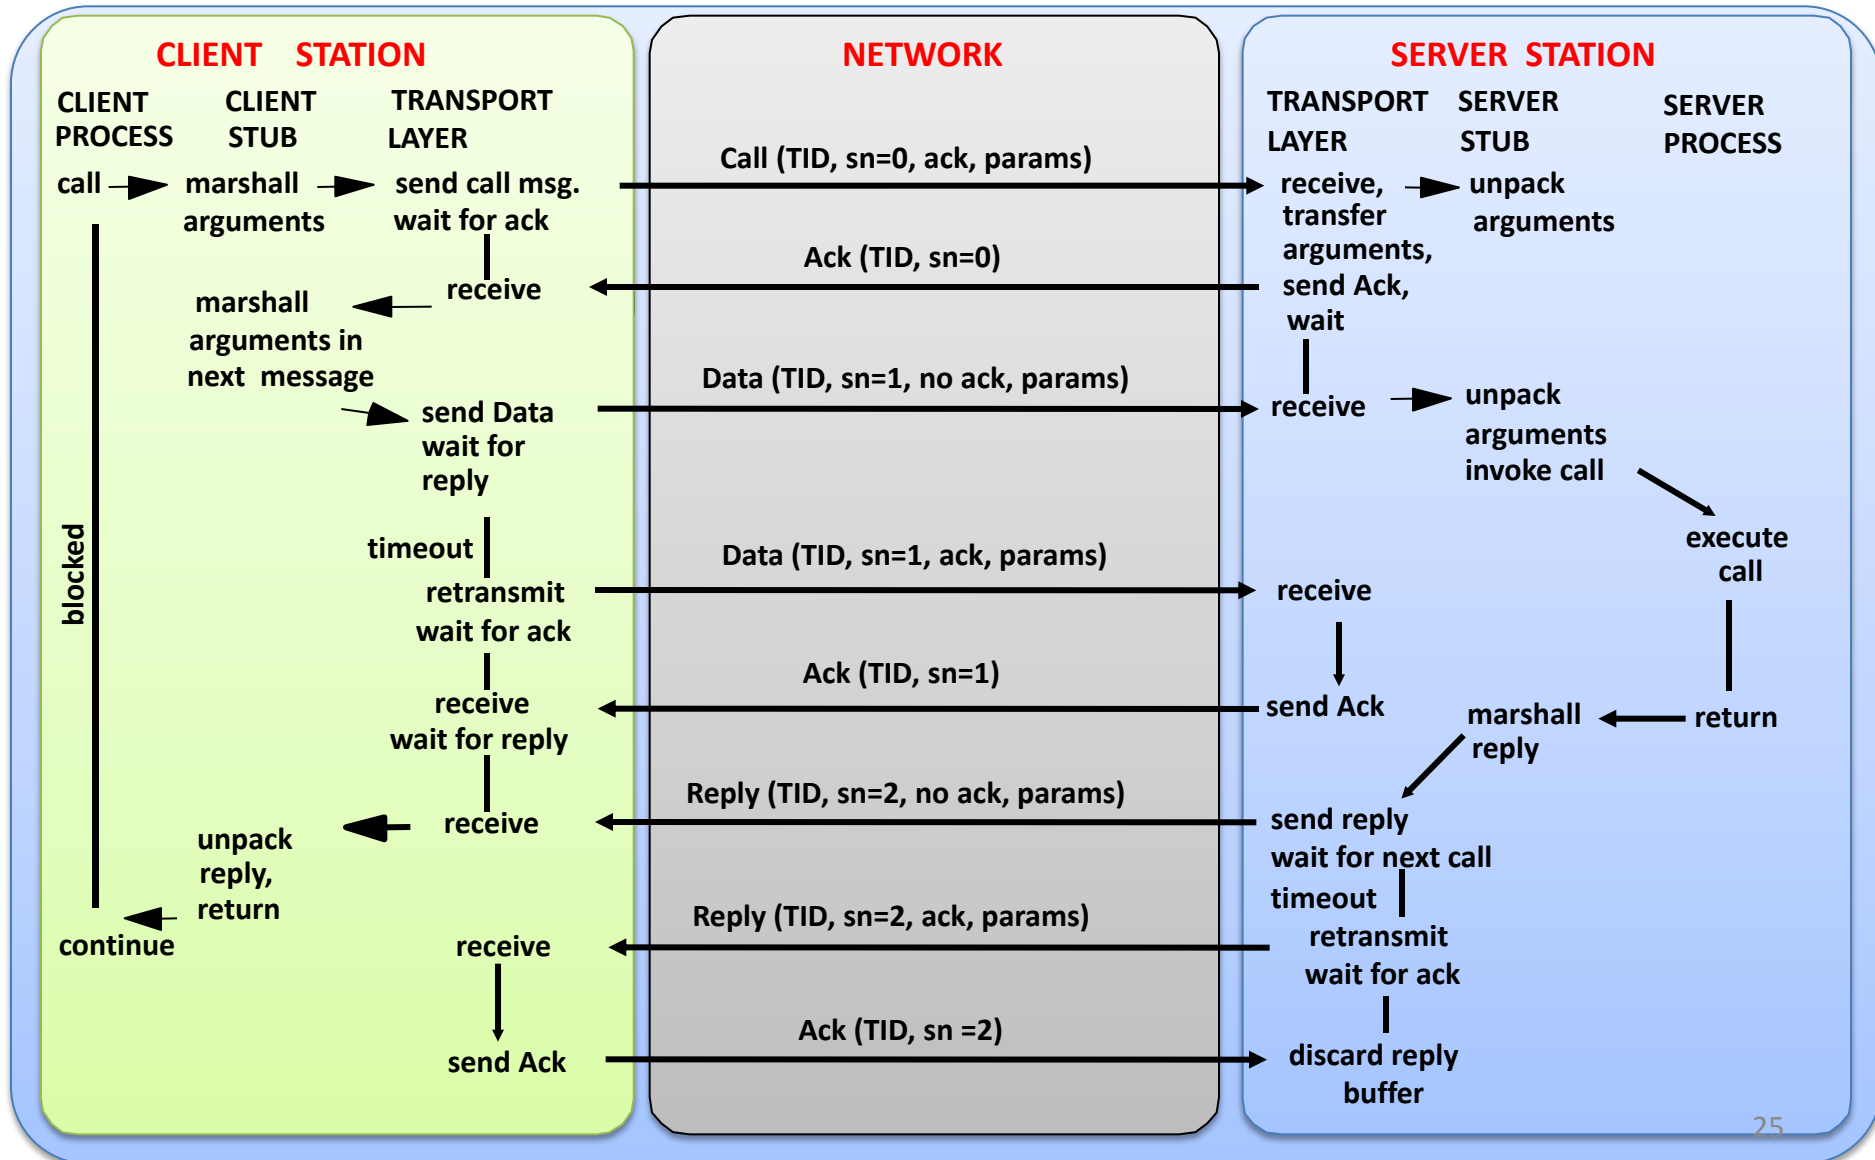
## Error Control

- After sending message set timeout
- Retransmit if no ACK
- Save reply until ACK received in case call repeated

## How can this be optimised?

Must also cater for long parameters requiring multiple messages to transfer

# RPC Implementation

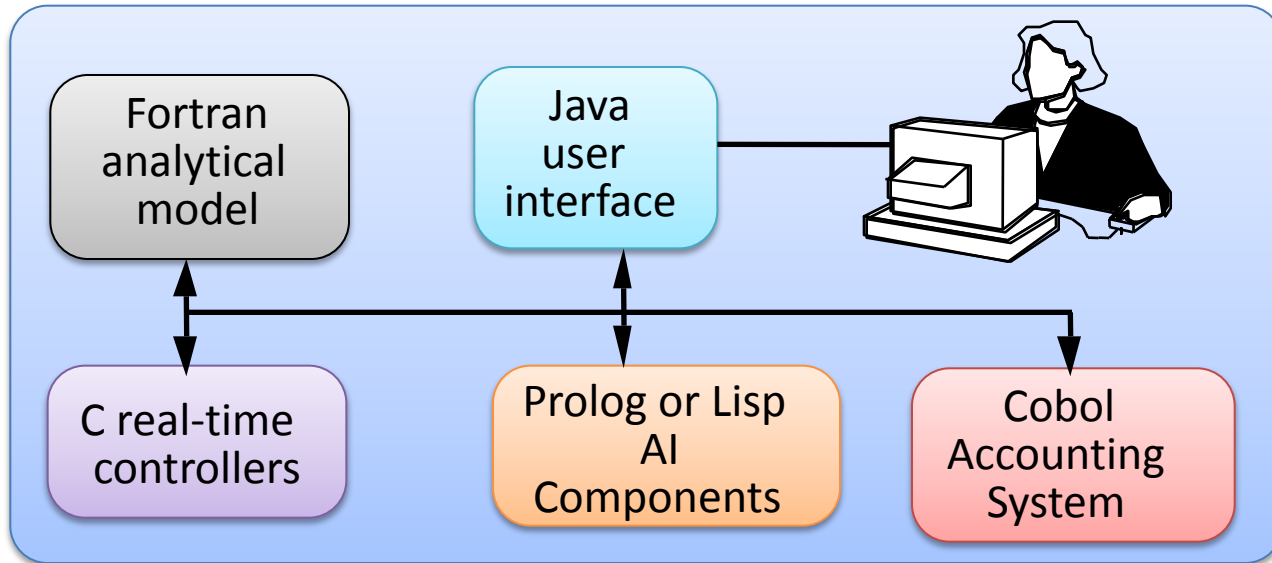




# RPC Parameters

- TID = Transaction identifier plus interface export identifier
- sn = message sequence number
- ack = please acknowledge message
- no ack = no acknowledgement expected
- params = in or out parameters

# Language Heterogeneity

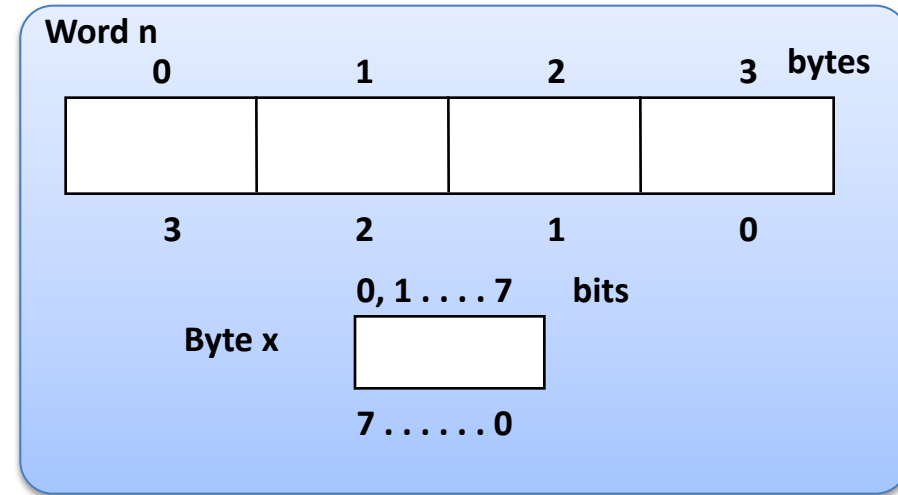


- Data structure representation differences:
  - Array implementation
  - Record implementation
  - Alignment of bytes on words etc.
  - No equivalent data structure e.g. no records in Fortran, no lists in C
- What can be done about this?

# Processor Heterogeneity

Computers differ in representation of:

- Characters - Ascii, Ebcdic, graphics.....
- Integers - 1 or 2's complement
- length
- Reals: mantissa & exponent length, format, base 2, 16 ...
- Bit and byte addressing within a word




Need to transform representations when transferring data

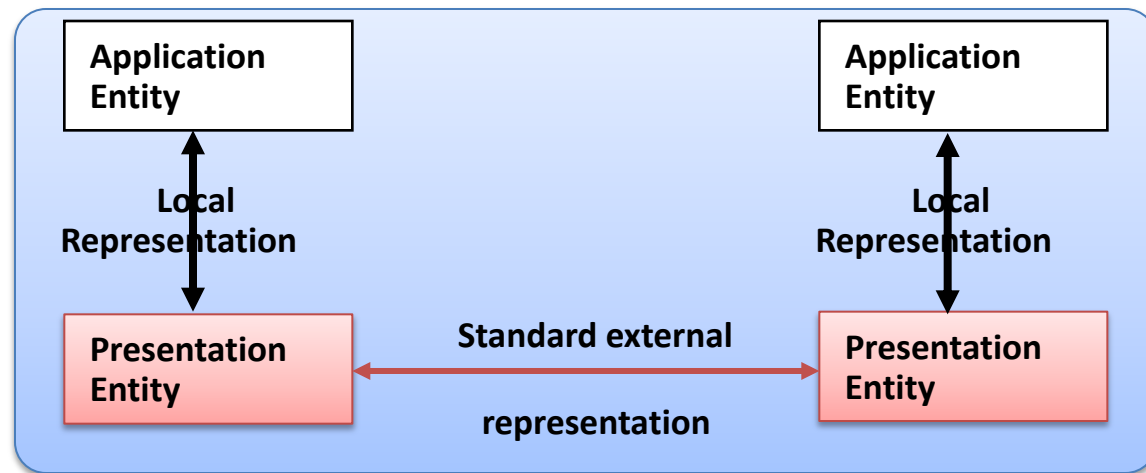
$N * (N-1)$  translators

for  $N$  machines

**What can be done  
about this?**



# Standard External Data Representation (XDR)



- Standard network wide external data representation (XDR) reduces number of translators  $\rightarrow 2N$  translators (to and from external standard) for  $N$  different machine types. Transformation must:
  - preserve meaning – can be difficult
  - resolve syntax differences

# Standard External Data Representation (XDR)

- Each Machine knows only about its own data representation and external representation
- Overhead of conversion when communicating between machines of same type
- What to do if only a few different machines?

# Data Encoding e.g. XDR Characteristics

- Fixed Length
  - 16 or 32 bit integers
  - more efficient transformation
  - maximum value limitation → truncation
- Variable length
  - E.g. strings of printable characters to represent numbers
  - Requires length indicator or end delimiter
  - No value limitation
  - Inefficient
  - Packed binary → discard leading 0's
  - Length field usually fixed length or extensible in bytes
    - most significant bit set → another byte follows

length	value	length	value
0 1 0	1 1	1 0 0	1 0 1 0
2	3	4	10

# Data Encoding e.g. XDR Characteristics

- Implicit Type
  - Types must be known in advance at receiver e.g. ports, object method parameters
  - Fewer overheads
- Explicit Tag or Type Identifier
  - Increased overheads
  - Information to perform transformation is self contained in message
  - Position independent
  - Needed for variant types
  - Can perform dynamic type checking



# Extensible Markup Language ( XML)

- Text based, explicit tags → human readable
- Very verbose, not human friendly → really aimed at machine processing
- Data items tagged with 'markup' strings describing logical structure
- Use start and end tags rather than length
- Extensible – users can define own tags
- Used for internet interactions and data storage e.g. XML databases
- Very inefficient encoding but can be compressed



# XML Elements and Attributes

**Element:** container for data – enclosed by start and end tag

**Attribute:** used to label data – usually name/value

```
<person id="123456789"> ← Attribute
    <name>Smith</name>
    <place>London</place> ← Place element
    <year>1934</year>
    <!-- a comment -->
</person >
```

Person Element

# XML Namespace

- Namespace used to scope names
- A set of names for a collection of element types and attributes
- Referenced by a url
- Specify namespace by a xmlns attribute
- Can use namespace name as prefix for names

```
<person pers:id="123456789" xmlns:pers =  
  "http://www.example.com/person">  
  <pers:name> Smith </pers:name>  
  <pers:place> London </pers:place >  
  <pers:year> 1934 </pers:year>  
</person>
```




Namespace attribute

# XML Schema

- Defines elements and attributes that can appear in a document
- Defines element nesting, number, ordering, whether empty or can include text
- For each element defines type and default value

```
<xsd:schema xmlns:xsd = URL of XML schema definitions >
  <xsd:element name= "person" type ="personType" />
  <xsd:complexType name="personType">
    <xsd:sequence>
      <xsd:element name = "name" type="xs:string"/>
      <xsd:element name = "place" type="xs:string"/>
      <xsd:element name = "year" type="xs:positiveInteger"/>
    </xsd:sequence>
    <xsd:attribute name= "id" type = "xs:positiveInteger"/>
  </xsd:complexType>
</xsd:schema>
```

# Representation Transformation



- What problems could occur when doing transformations, e.g. with numbers?

# Semantics of Representation

- Two representations can have similar syntax but different meaning
  - E.g. complex numbers – (float x, y ) = rectangular or polar coordinates → transformation is application dependent
- Type may have no meaning outside own context
  - E.g. pointer, file name
- Procedures passed as parameters
  - Cannot transfer code to different computer for execution
- What should be done?

# Example of Use of Encode


```
struct rec {  
    int a;  
    boolean b;  
};  
  
struct form {  
    int x;  
    float y;  
    rec z [3]; /* assume 3 elements */  
};  
  
form obj = (5, 23.75, 10, true, 5, false, 7, true)
```

→ can be “flattened” for transfer:  
where I = int, F = float, B = boolean

I	F	I	B	I	B	I	B
		└──────────┘					
x	y	3 elements of z					

# Structural Information

- Structural information must be maintained
  - Structural information represented internally by pointers (addresses)
  - must be flattened into a linear message



**Recursively Flattened Tree**

{ { {0, 0, value C} {0, 0 value D} value B} {0 {0, 0 value F} value E} value A}

# Transferring Cyclic Structures

- Use Encode and Decode procedures for primitive types and simple constructed types
- Structural information must be flattened:
  - Number sub-objects
  - Transform pointers into handles (i.e. number) of sub-objects

Sub object 1

Handle 2

Null

Contents

Sub object 2

Handle 3

Handle 1

Contents

Sub object 3

Handle 4

Handle 2


Contents

Sub object 4

Null

Handle 3

Contents





# Java Object Serialization

- Java objects can be passed as arguments and results in RMI. Object is an instance of a Java serializable class

```
public class Person implements Serializable {  
    private String name;  
    private String place;  
    private int year;  
    public Person (String aName, String aPlace, int aYear) {  
  
        name = aName;  
        place = aPlace;  
        year = aYear;  
    }  
    // methods for accessing instance variables  
}
```

# Java Object Serialization

```
FileOutputStream fos = new  
    FileOutputStream("t.tmp");  
  
ObjectOutputStream out =  
    new ObjectOutputStream(fos);  
  
out.writeObject  
    (new Person("Anandha", "London", 2015));  
  
out.close();
```

<http://docs.oracle.com/javase/7/docs/api/java/io/ObjectOutputStream.html>  
<https://docs.oracle.com/javase/7/docs/platform/serialization/spec/protocol.html>

# Java Object Serialization

- Java objects can contain references to other objects
  - All referenced objects are serialized together
- References are converted to handles i.e. internal references to object within the serialized form
- Each object is serialized only once – detect multiple references to same object

# Java Object Serialization

- Serialization
  - Write class information
  - Write types and names of instance variables
  - If instance variables are of a new class, then write their class information followed by types and names of instance variables
  - Uses reflection – ability to enquire about properties of a class e.g. names and types of instance variables and methods

# Java Object Serialization

- |  |   |                                       |
|--|---|---------------------------------------|
| 1. Specifies it's a serialization protocol | 10. Field type (Int)                                      | 20. End of block data for this object |
| 2. Serialization version                   | 11. Length of field name                                  | 21. No more superclasses              |
| 3. Specifies new Object                    | 12. Field name ( <b>year</b> )                            | 22. Int value                         |
| 4. Specifies new Class                     | 13. Field type (String)                                   | $(7 * 16^2 + 13 * 16 + 15 = 2015)$    |
| 5. Length of class name                    | 14. Length of field name                                  | 23. String                            |
| 6. Name of class ( <b>Person</b> )         | 15. Field name ( <b>name</b> )                            | 24. Length of value                   |
| 7. serialVersionUID                        | 16. String instance (type + length + "java/lang/string;") | 25. Value ( <b>Anandha</b> )          |
| 8. Flags – object supports serialization   | 17. Field type (String)                                   | 26. String                            |
| 9. Number of fields in class               | 18. Length of field name                                  | 27. Length of value                   |
|  | 19. Field name ( <b>place</b> )                           | 28. Value ( <b>London</b> )           |

```
ac ed 00 05 73 72 00 06 50 65 72 73 6f 6e b1 49
3f a5 9f 2f a4 0e 02 00 03 49 00 04 79 65 61 72
4c 00 04 6e 61 6d 65 74 00 12 4c 6a 61 76 61 2f
6c 61 6e 67 2f 53 74 72 69 6e 67 3b 4c 00 05 70
6c 61 63 65 71 00 7e 00 01 78 70 00 00 07 df 74
00 07 41 6e 61 6e 64 68 61 74 00 06 4c 6f 6e 64
6f 6e
```

Hex dump of the t.tmp file from earlier

- Using `od -t x1 t.tmp`

# Summary

- Message passing systems map closely onto the underlying communication services, however RPCs and Object invocation are more complex to implement
- They require binding implementation and have to cater for failures of client, server, name servers or communication system
- RPCs and invocations can either be implemented by an optimised special purpose protocol or by a general purpose Transport protocol such as TCP

# Summary

- Translation to a standard external representation should be optional to avoid unnecessary overheads
- Typed interfaces do not need explicit tags in the XDR
- Some types cannot be transferred e.g. memory addresses
- Complex data types must be “flattened” for transfer to a remote machine (or to disc store) and addresses transformed to local references (e.g. array index)