

Interactive Computer Graphics: Lecture 4

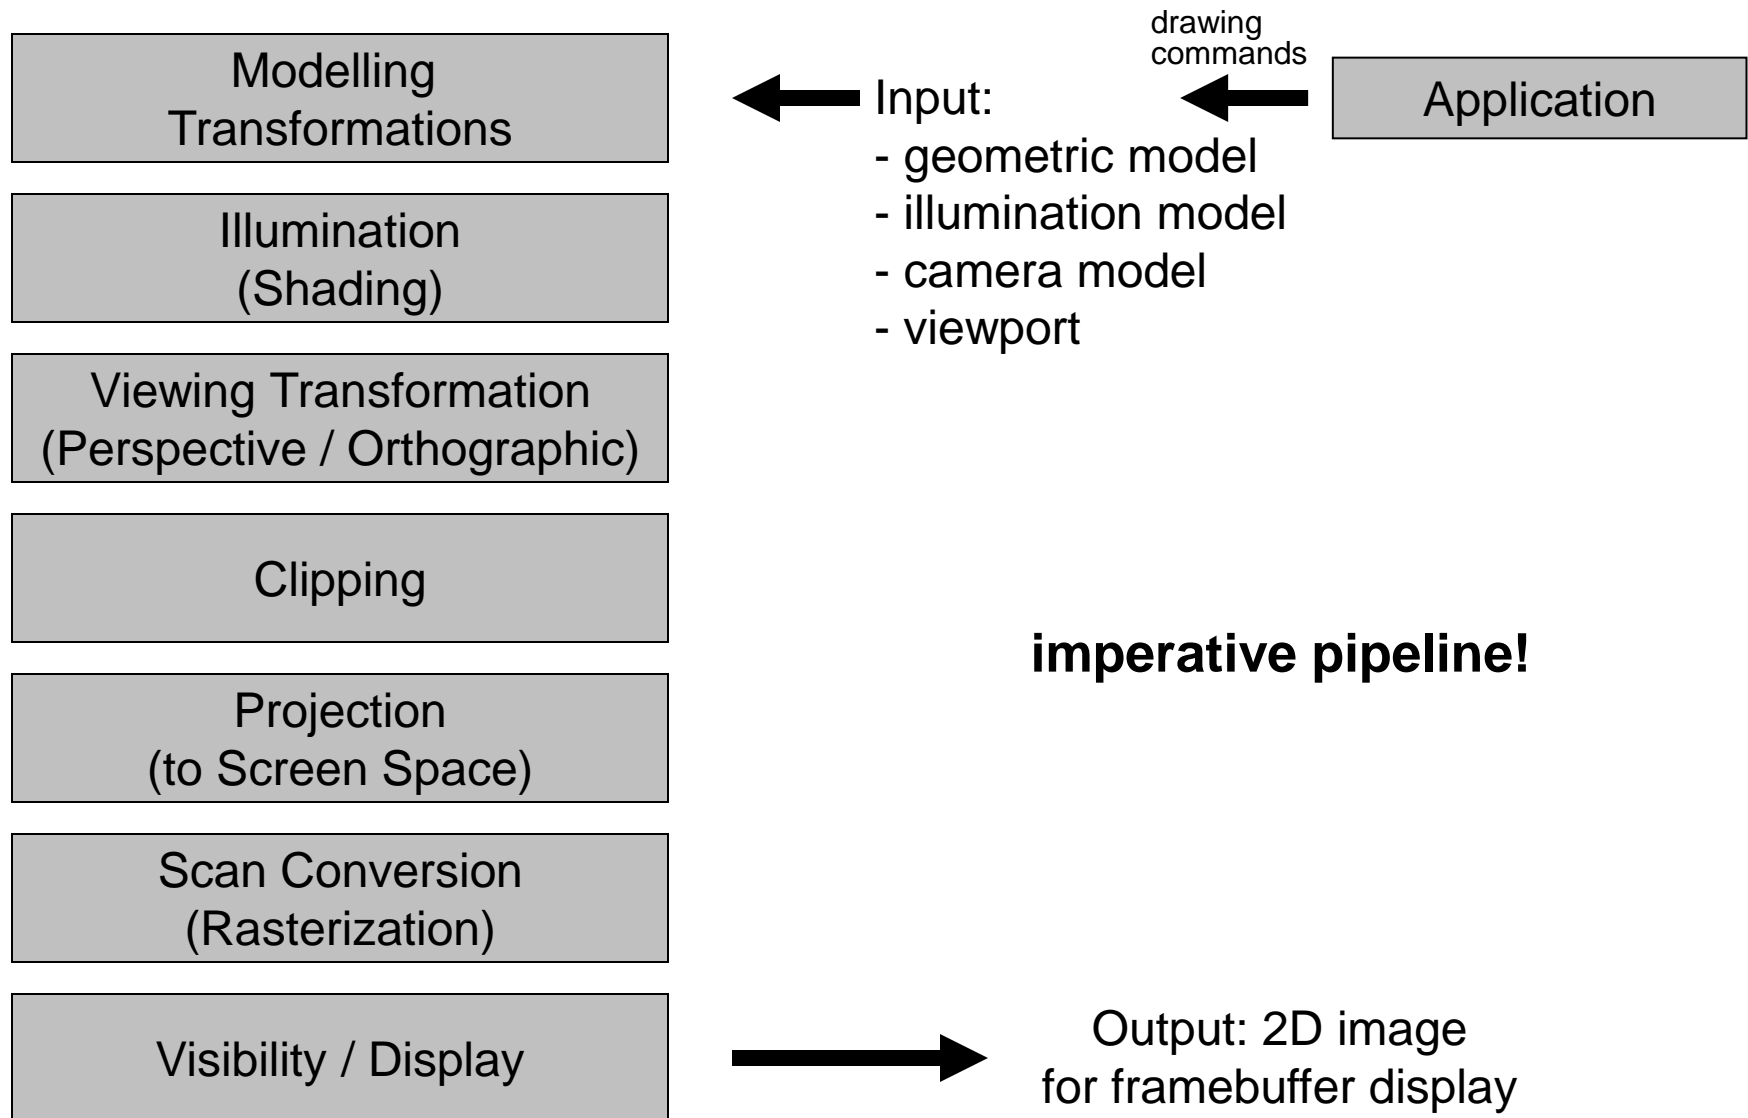
The Graphics Pipeline: OpenGL and GLSL

Thanks to Markus Steinberger and
Dieter Schmalstieg

The Graphics Pipeline: High-level view

- Declarative (What, not How)
 - For example virtual camera with scene description, e.g. scene graphs
 - Every object may know about every other object
 - Renderman, Inventor, OpenSceneGraph,...
- Imperative (How, not What)
 - Emit a sequence of drawing commands
 - For example: draw a point (vertex) at position (x,y,z)
 - Objects can be drawn independant from each other
 - OpenGL, PostScript, etc.

The Graphics Pipeline



The Graphics Pipeline

Modelling
Transformations

Illumination
(Shading)

Viewing Transformation
(Perspective / Orthographic)


Clipping

Projection
(to Screen Space)

Scan Conversion
(Rasterization)

Visibility / Display

- 3D models are defined in their own coordinate system
- Modeling transformations orient the models within a common coordinate frame (world coordinates)



The Graphics Pipeline

Modelling
Transformations

Illumination
(Shading)

Viewing Transformation
(Perspective / Orthographic)


Clipping

Projection
(to Screen Space)

Scan Conversion
(Rasterization)

Visibility / Display

- Vertices are lit (shaded) according to material properties, surface properties and light sources
- Uses a local lighting model



The Graphics Pipeline

Modelling
Transformations

Illumination
(Shading)

Viewing Transformation
(Perspective / Orthographic)

Clipping

Projection
(to Screen Space)

Scan Conversion
(Rasterization)

Visibility / Display

- Maps world space to eye (camera) space (matrix evaluation)
- Viewing position is transformed to origin and viewing direction is oriented along some axis (typically z)



The Graphics Pipeline

Modelling
Transformations

Illumination
(Shading)

Viewing Transformation
(Perspective / Orthographic)


Clipping

Projection
(to Screen Space)

Scan Conversion
(Rasterization)


Visibility / Display

- Portions of the scene outside the viewing volume (view frustum) are removed (clipped)
- Transform to Normalized Device Coordinates



Eye space

NDC



The Graphics Pipeline

Modelling
Transformations

Illumination
(Shading)

Viewing Transformation
(Perspective / Orthographic)


Clipping

Projection
(to Screen Space)


Scan Conversion
(Rasterization)

Visibility / Display



- The objects are projected to the 2D imaging plane (screen space)



NDC



Screen Space



The Graphics Pipeline

Modelling
Transformations

Illumination
(Shading)

Viewing Transformation
(Perspective / Orthographic)


Clipping

Projection
(to Screen Space)

Scan Conversion
(Rasterization)

Visibility / Display

- Rasterizes objects into pixels
- Interpolate values inside objects (color, depth, etc.)



The Graphics Pipeline

Modelling
Transformations

Illumination
(Shading)

Viewing Transformation
(Perspective / Orthographic)


Clipping

Projection
(to Screen Space)

Scan Conversion
(Rasterization)

Visibility / Display

- Handles occlusions and transparency blending
- Determines which objects are closest and therefore visible
- Depth buffer



What do we want to do?

- Computer-generated imagery (CGI) in **real-time**
- Very computationally demanding:
 - full HD at 60hz:
 $1920 \times 1080 \times 60\text{hz} = 124 \text{ Mpx/s}$
 - and that's just the output data


→ use specialized hardware for
immediate mode (real-time) graphics

Solution

Most of real-time graphics is based on

- rasterization of graphic *primitives*

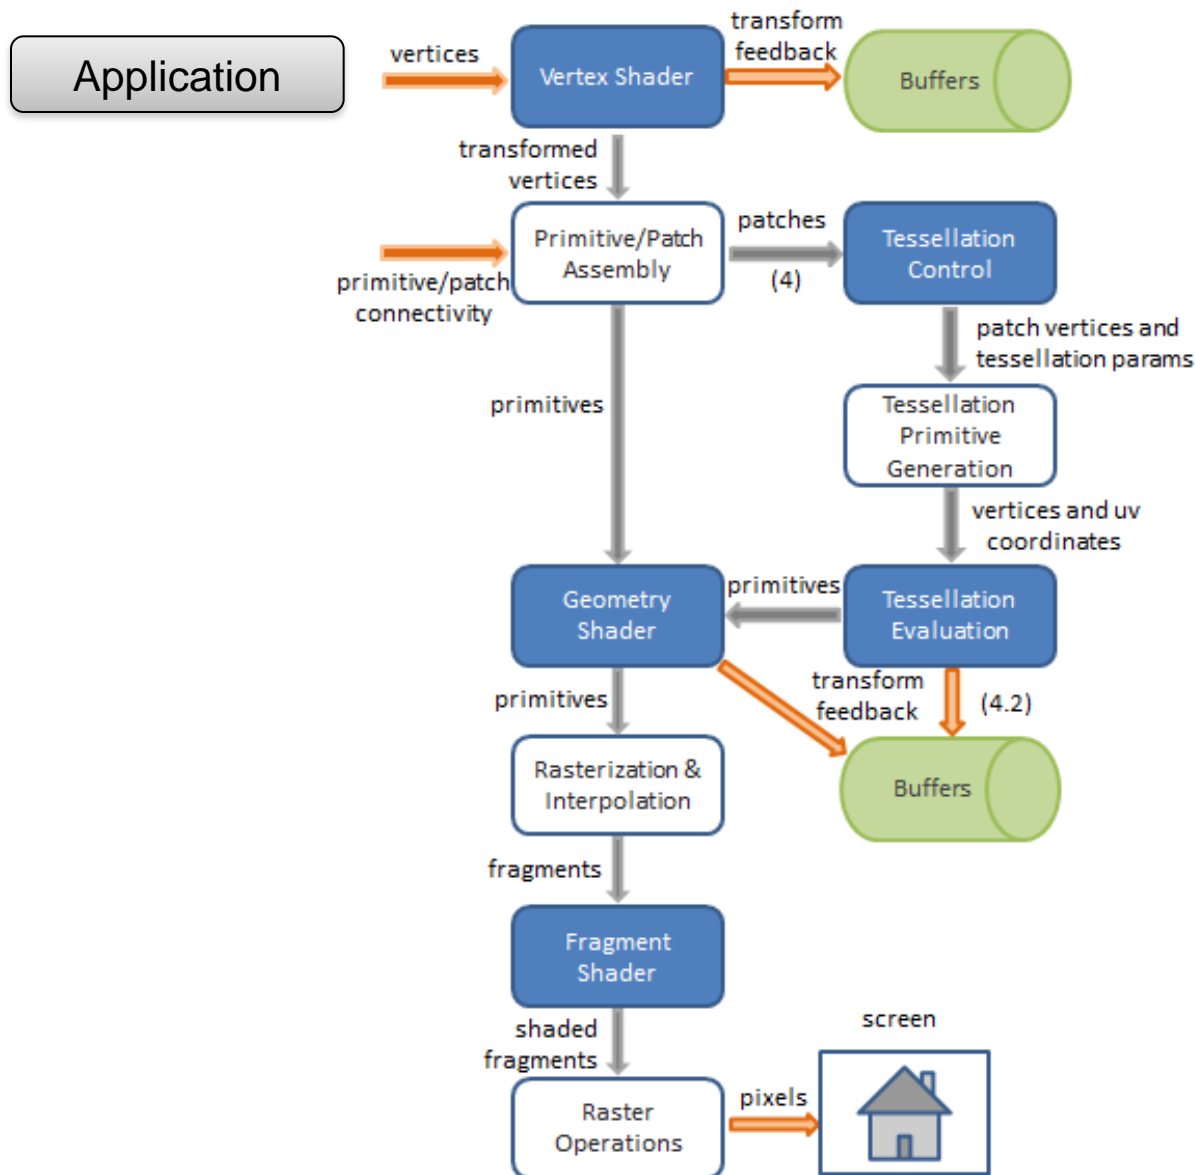
- points
- lines
- triangles
- ...



- Implemented in hardware

- *graphics processing unit* (GPU)
- controlled through an API such as OpenGL
- certain parts of graphics pipeline are programmable, e.g. using GLSL
 - ➔ shaders

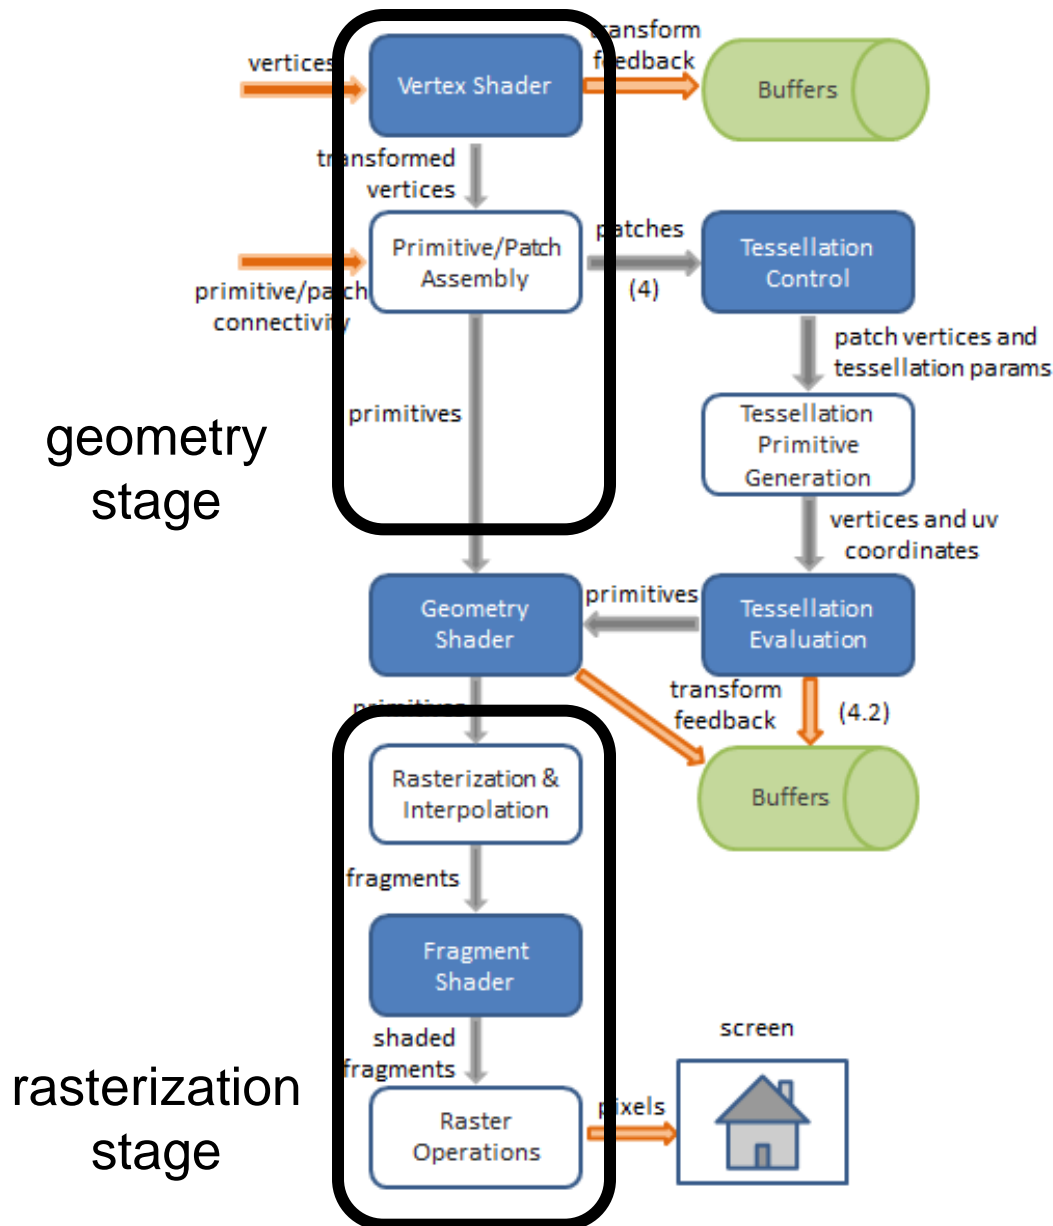
The Graphics Pipeline: OpenGL 3.2 and later



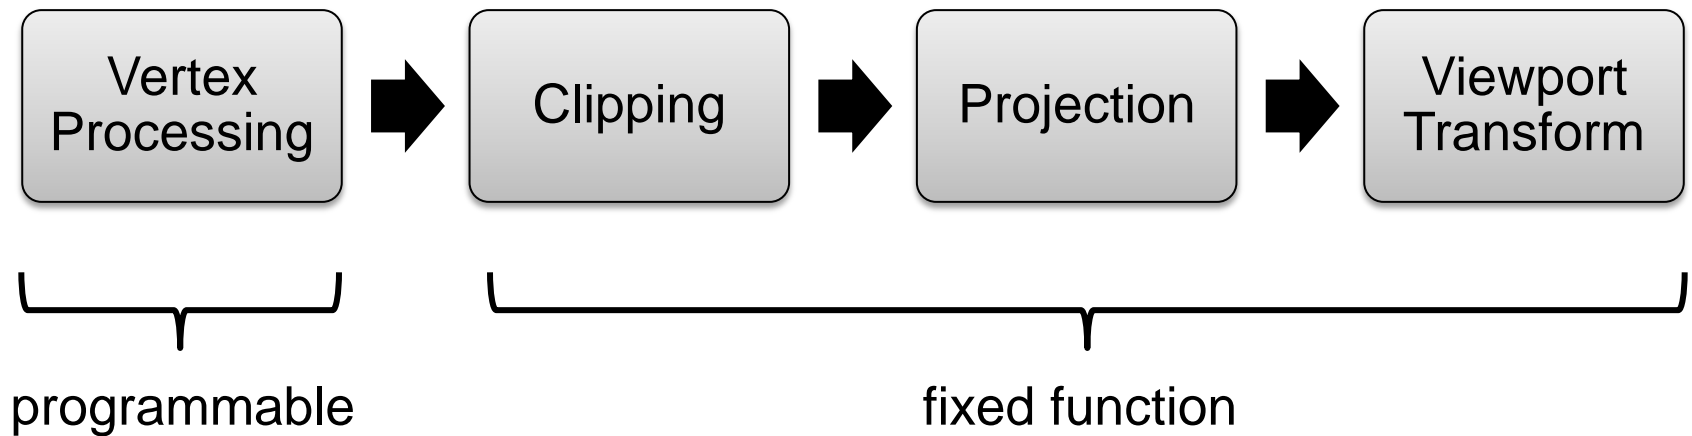
Programmable

Fixed function

The Graphics Pipeline: OpenGL 3.2 and later




Geometry Stage



Geometry Stage: Vertex Processing

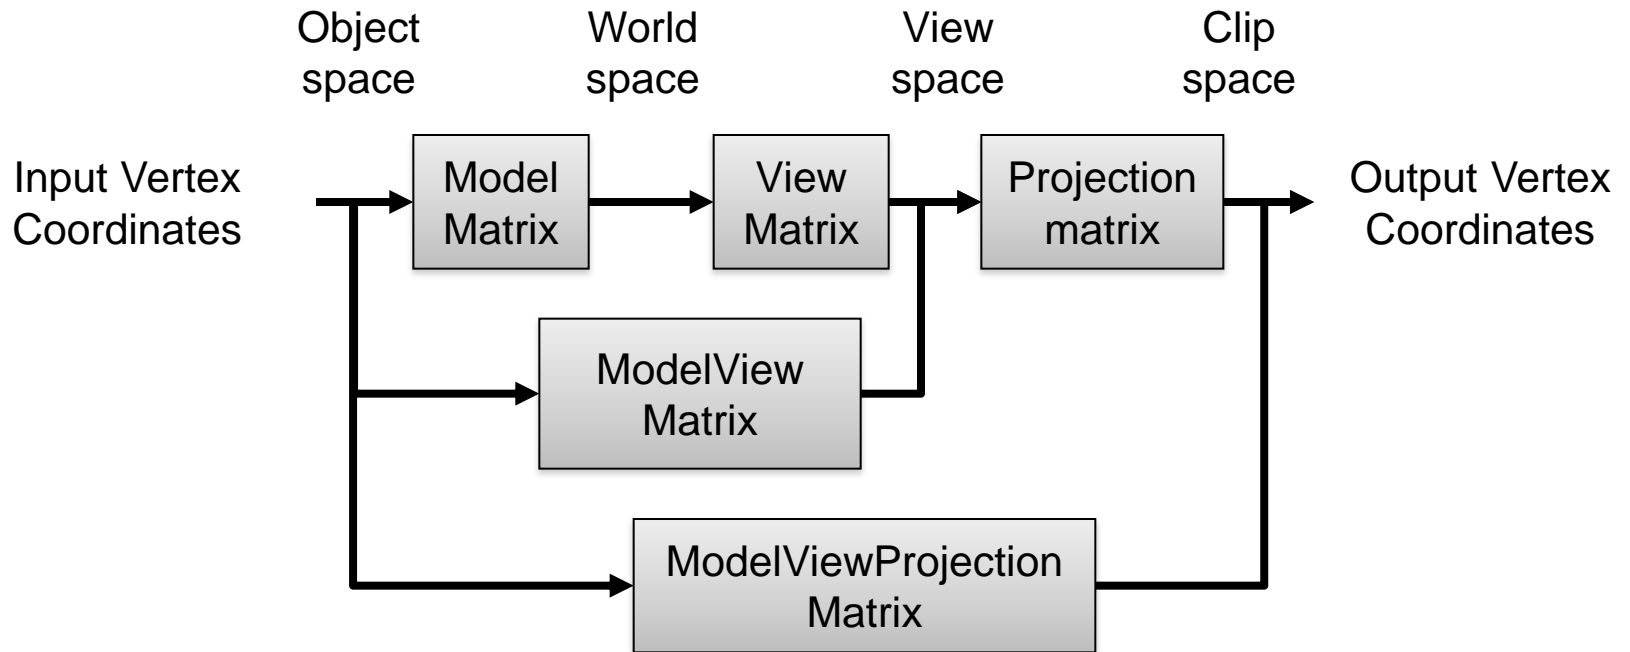
- The input vertex stream
 - composed of arbitrary vertex attributes (position, color, ...).
- is transformed into stream of vertices mapped onto the screen
 - composed of their clip space coordinates and additional user-defined attributes (color, texture coordinates, ...).
 - clip space: homogeneous coordinates
- by the **vertex shader**
 - GPU program that implements this mapping.



- Historically, “Shaders” were small programs performing lighting calculations, hence the name.




Geometry Stage: Vertex Post-Processing

- Uses a common transformation model in rasterization-based 3D graphics:





Geometry Stage: Vertex Post-Processing

- Clipping
 - Primitives not entirely in view are clipped to avoid projection errors
- Projection
 - Projects clip space coordinates to the image plane
 - Primitives in normalized device coordinates
- Viewport Transform:
 - Maps resolution-independent normalized device coordinates to a rectangular window in the frame buffer, the viewport.
 - Primitives in window (pixel) coordinates

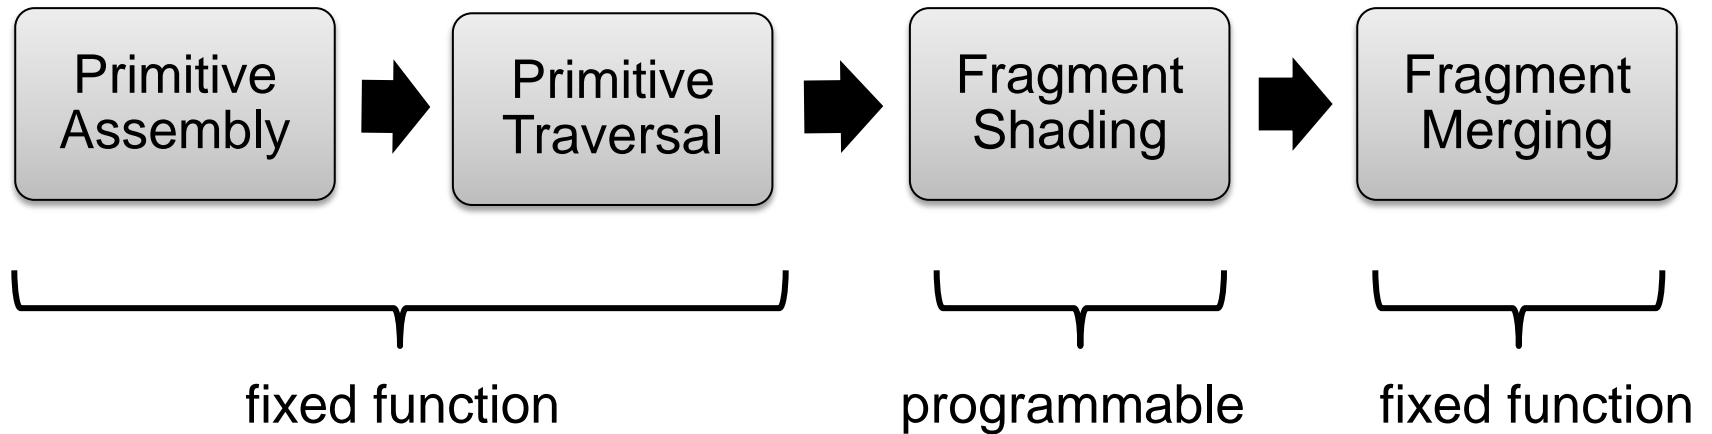


Geometry Shader

- Optional stage between vertex and fragment shader
- In contrast to the vertex shader, the geometry shader has full knowledge of the primitive it is working on
 - For each input primitive, the geometry shader has access to all the vertices that make up the primitive, including adjacency information.
- Can generate primitives dynamically
 - Procedural geometry, e.g. growing plants




Rasterization Stage




Rasterization Stage

- Primitive assembly
 - Backface culling
 - Setup primitive for traversal
- Primitive traversal (“scan conversion”)
 - Sampling (triangle \rightarrow fragments)
 - Interpolation of vertex attributes (depth, color, ...)
- Fragment shading
 - Compute fragment colors
- Fragment merging
 - Compute pixel colors from fragments
 - Depth test, blending, ...




Rasterization – Coordinates




Rasterization – Rules

- Different rules for each primitive type
 - “fill convention”
- Non-ambiguous!
 - artifacts...
- Polygons:
 - Pixel center contained in polygon
 - Pixels on edge: only one is rasterized




Fragment Shading

- “Fragment”:
 - Sample produced during rasterization
 - Multiple fragments are *merged* into pixels.
- Given the interpolated vertex attributes,
 - output by the Vertex Shader
- the *Fragment Shader* computes color values for each fragment.
 - Apply textures
 - Lighting calculations
 - ...

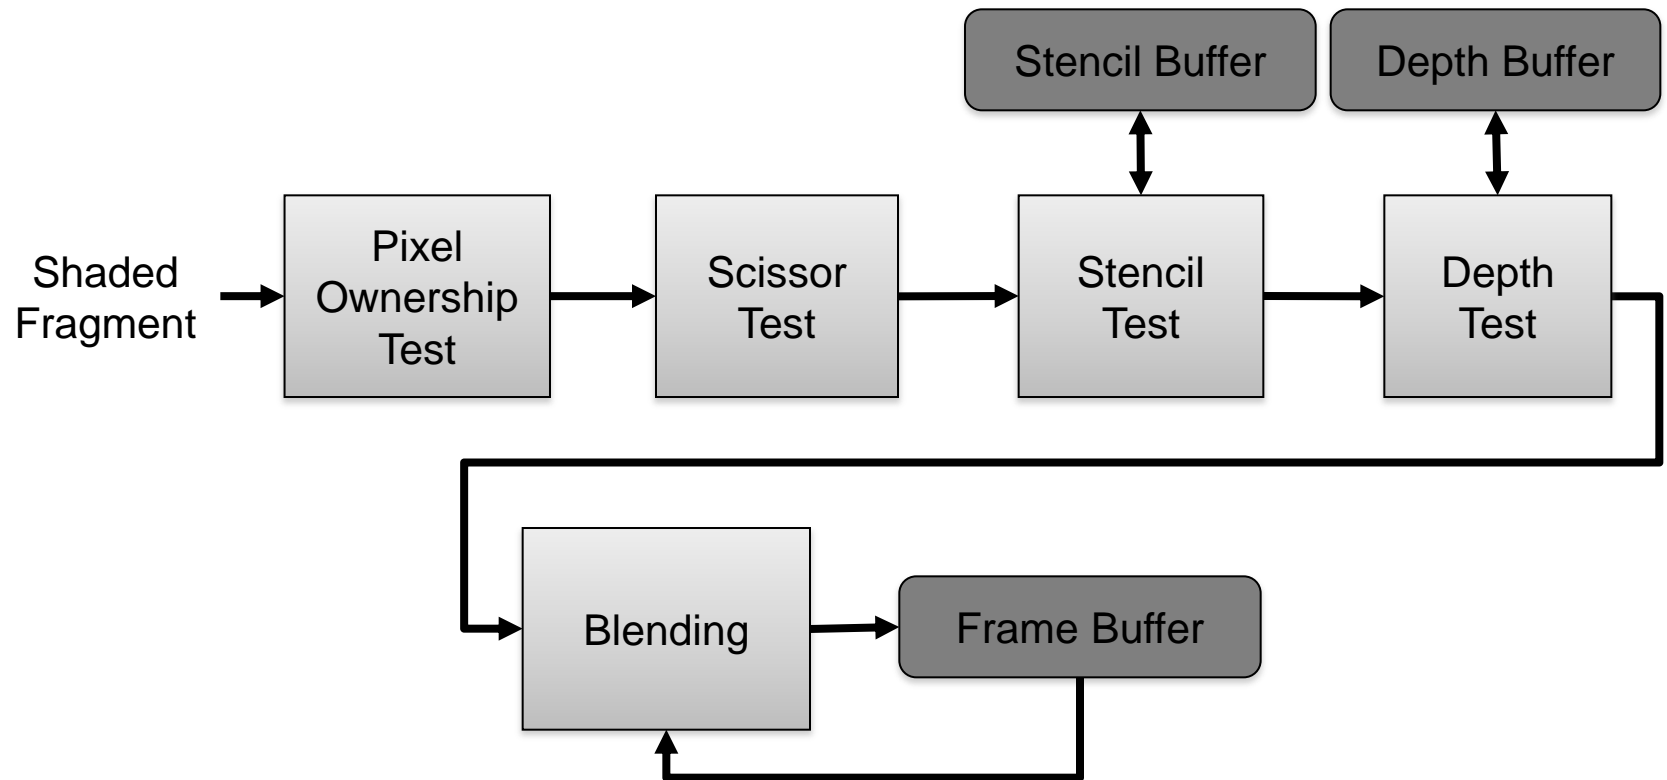


Fragment Merging

- Multiple primitives can cover the same pixel.
- Their Fragments need to be composed to form the final pixel values.
 - Blending
 - Resolve Visibility
 - Depth buffering




Fragment Merging



Display Stage

- Gamma correction
- Historically: Digital to Analog conversion
- Today: Digital scan-out, HDMI encryption, etc.



Display Format

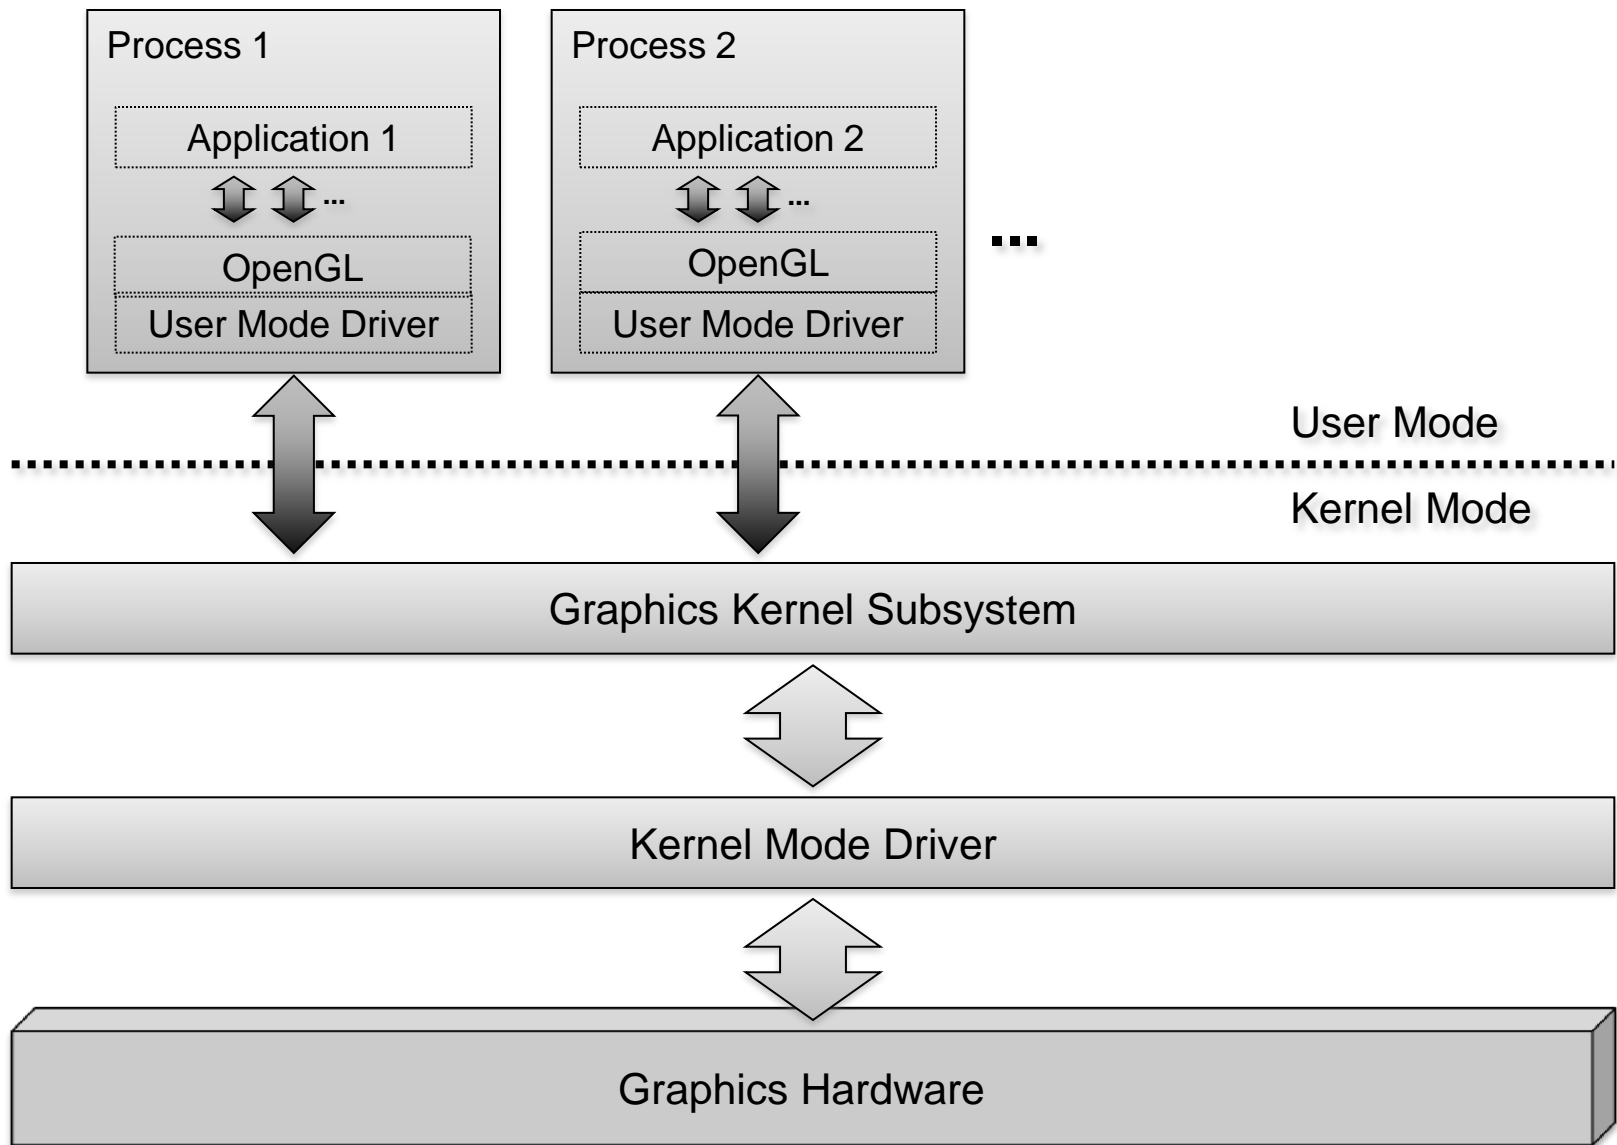
- Frame buffer pixel format:
RGBA vs. index (obsolete)
- Bits: 16, 32, 64, 128 bit floating point, ...
- Double buffer vs. single buffer
- Quad-buffered stereo
- Overlays (extra bitplanes)
- Auxilliary buffers: alpha, stencil

Functionality vs. Frequency

- Geometry processing = per-vertex
 - Transformation and Lighting (T&L)
 - Historically floating point, complex operations
 - Millions of vertices per second
 - Today: Vertex Shader
- Fragment processing = per-fragment
 - Blending, texture combination
 - Historically fixed point and limited operations
 - Billions of fragments (“Gigapixel” per second)
 - Today: Fragment Shader

Architectural Overview

- Graphics Hardware is a shared resource
- User Mode Driver (UMD)
 - Prepares Command Buffers for the hardware
- Graphics Kernel Subsystem
 - Schedules access to the hardware
- Kernel Mode Driver (KMD)
 - Submits Command Buffers to the hardware



What is OpenGL?

- a low-level graphics API specification
 - not a library!
 - The interface is platform independent,
 - but the implementation is platform dependent.
 - Defines
 - an abstract rendering device.
 - a set of functions to operate the device.
 - “immediate mode” API
 - drawing commands
 - no concept of permanent objects

What is OpenGL?

- Platform provides OpenGL *implementation*.
 - Part of the graphics driver, or
 - runtime library built on top of the driver
- Initialization through platform specific API
 - WGL (Windows)
 - GLX (Unix/Linux)
 - EGL (mobile devices)
 - ...
- State machine for high efficiency!

Basic Concepts

- Context
- Resources
- Object Model
 - Objects
 - Object Names
 - Bind Points (Targets)

Context

- Represents an instance of OpenGL
- A process can have multiple contexts
 - These can share resources
- A context can be *current* for a given thread
 - one to one mapping
 - only one current context per thread
 - context only current in one thread at the same time
 - OpenGL operations work on the current context

Resources

- Act as
 - sources of input
 - sinks for output
- Examples:
 - buffers
 - images
 - state objects
 - ...

Resources

- Buffer objects
 - linear chunks of memory
- Texture images
 - 1D, 2D, or 3D arrays of *texels*
 - Can be used as input for *texture sampling*

Object Model

- OpenGL is object oriented
 - but in its own, strange way
- Object instances are identified by a *name*
 - basically just an unsigned integer handle
- Commands work on *targets*
 - Each target has an object currently *bound* to the target
 - That's the one commands will work with
- Object oriented, you said?
 - target \Leftrightarrow type
 - commands \Leftrightarrow methods

Object Model

- By binding a name to a target
 - the object it identifies becomes current for that target
 - “latched state”
 - might change soon (`EXT_direct_state_access`)
 - An object is created when a name is first bound.
- Notable exceptions: Shader Objects, Program Objects
 - Some commands work directly on object names.

Example: Buffer Object

```
GLuint my_buffer;

// request an unused buffer object name
glGenBuffers(1, &my_buffer);

// bind name as GL_ARRAY_BUFFER
// bound for the first time ⇒ creates
glBindBuffer(GL_ARRAY_BUFFER, my_buffer);

// put some data into my_buffer
glBufferStorage(GL_ARRAY_BUFFER, ...);


// “unbind” buffer
glBindBuffer(GL_ARRAY_BUFFER, 0);

// probably do something else...
glBindBuffer(GL_ARRAY_BUFFER, my_buffer);
// use my_buffer...

glDrawArrays(GL_TRIANGLES, 0, 33);
// draw content example (type, startIdx, numer of elements)

// delete buffer object, free resources, release buffer object name
glDeleteBuffers(1, &my_buffer);
```


Primitive types



Draw Call

- After pipeline is configured:
 - issue *draw call* to actually draw something

e. g.:

```
glBegin(GL_TRIANGLE_STRIP);  
glColor3f(0.0, 1.0, 0.0);  
glVertex3f(1.0, 0.0, 0.0);  
...  
glEnd();
```

primitive type

Color "state"

vertex index

Shaders

- Shader Objects
 - parts of a pipeline (Vertex Shader, Fragment Shader, etc.)
 - compiled during runtime from GLSL code
 - OpenGL Shading Language
 - C-like syntax
- Program Object
 - a whole pipeline
 - Shader objects linked together during runtime

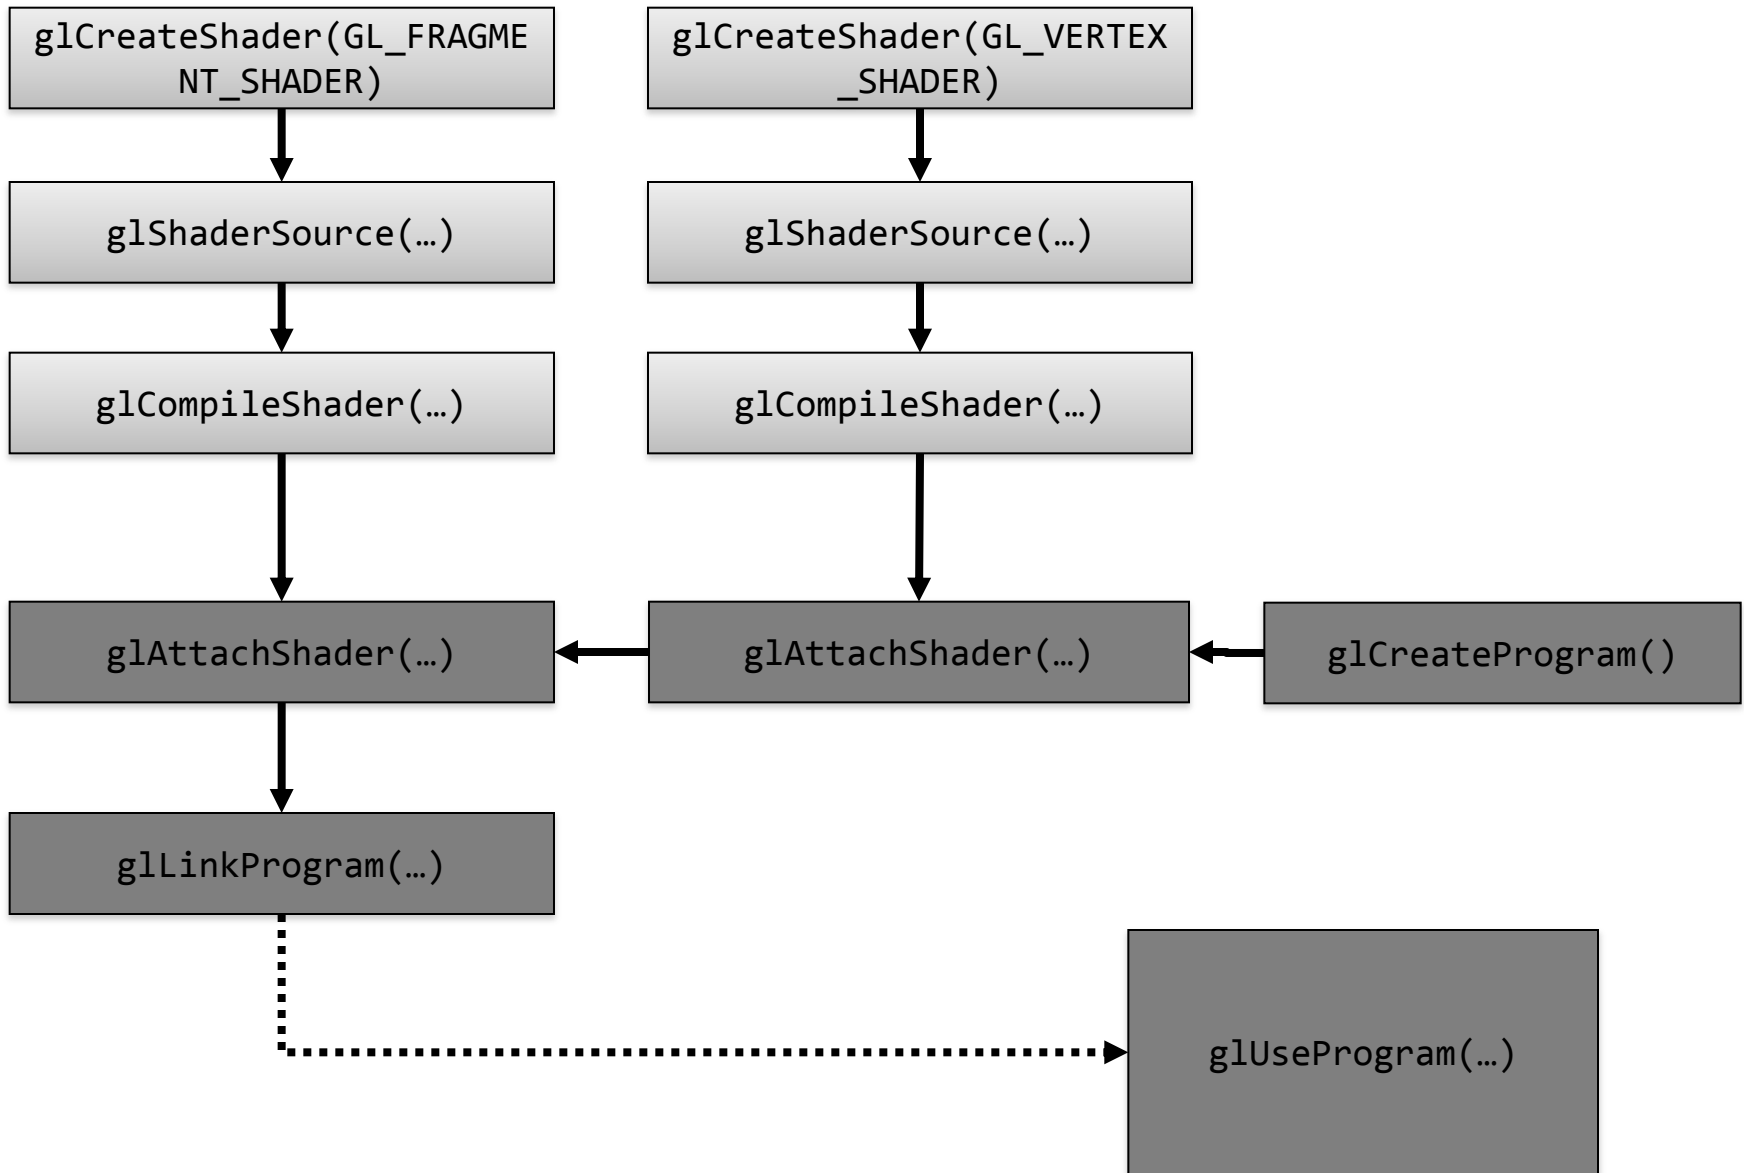
Anatomy of a GLSL Shader

```
1  #version 330
2
3  uniform vec4 some_uniform;
4
5  layout(location = 0) in vec3 some_input;
6  layout(location = 1) in vec4 another_input;
7
8  out vec4 some_output;
9  void main()
10 {
11
12 }
```

Set by application
(configuration values, e.g.
ModelViewProjection Matrix)

Optional flexible
register
configuration
between shaders

Output definition for
next shader stage



Vertex Shader

- Processes each vertex
- Input: vertex attributes
- Output: vertex attributes
 - mandatory: `gl_Position`

Rasterizer

- Fixed-function
- Rasterizes primitives
- Input: primitives
 - vertex attributes
- Output: fragments
 - interpolated vertex attributes

Fragment Shader

- Processes each fragment
- Input: interpolated vertex attributes
- Output: fragment color

Example: Vertex Shader

```
1  #version 150 compatibility
2  layout(location = 0) in vec3 vertex_position;
3  layout(location = 1) in vec4 vertex_color;
4
5  out vertexData
6  {
7      vec3 pos;
8      vec3 normal;
9      vec4 color;
10 }vertex;
11
12 void main()
13 {
14     vertex.pos = vec3(gl_ModelViewMatrix * gl_Vertex);
15     vertex.normal = normalize(gl_NormalMatrix * gl_Normal);
16     gl_Position = gl_ModelViewProjectionMatrix * gl_Vertex;
17     vertex.color = vec4(1.0,0.0,0.0,1.0);
18 }
```

Example: Fragment Shader

```
1  #version 150 compatibility
2  in vec4 color;
3
4  in fragmentData
5  {
6      vec3 pos;
7      vec3 normal;
8      vec4 color;
9  }frag;
10
11 void main()
12 {
13     vec4 outcol = frag.color;
14     gl_FragColor = outcol;
15 }
```